

Self-Driving Networks

Breaking new ground in network automation



Laurent Vanbever
nsg.ee.ethz.ch

ETH Zurich
May 29 2019

1962



Paul Baran (1926-2011)

American electrical engineer

created the notion of a distributed network
which could maintain communications
in the face of a thermonuclear attack

inventor of packet switching
(with Donald Davies)

On distributed communication networks

Paul Baran—Sept. 1962

Introduces the concept of

- switching
- distributed routing
- reliable transport

among many others...

ON DISTRIBUTED COMMUNICATIONS NETWORKS
Paul Baran*
The RAND Corporation, Santa Monica, California

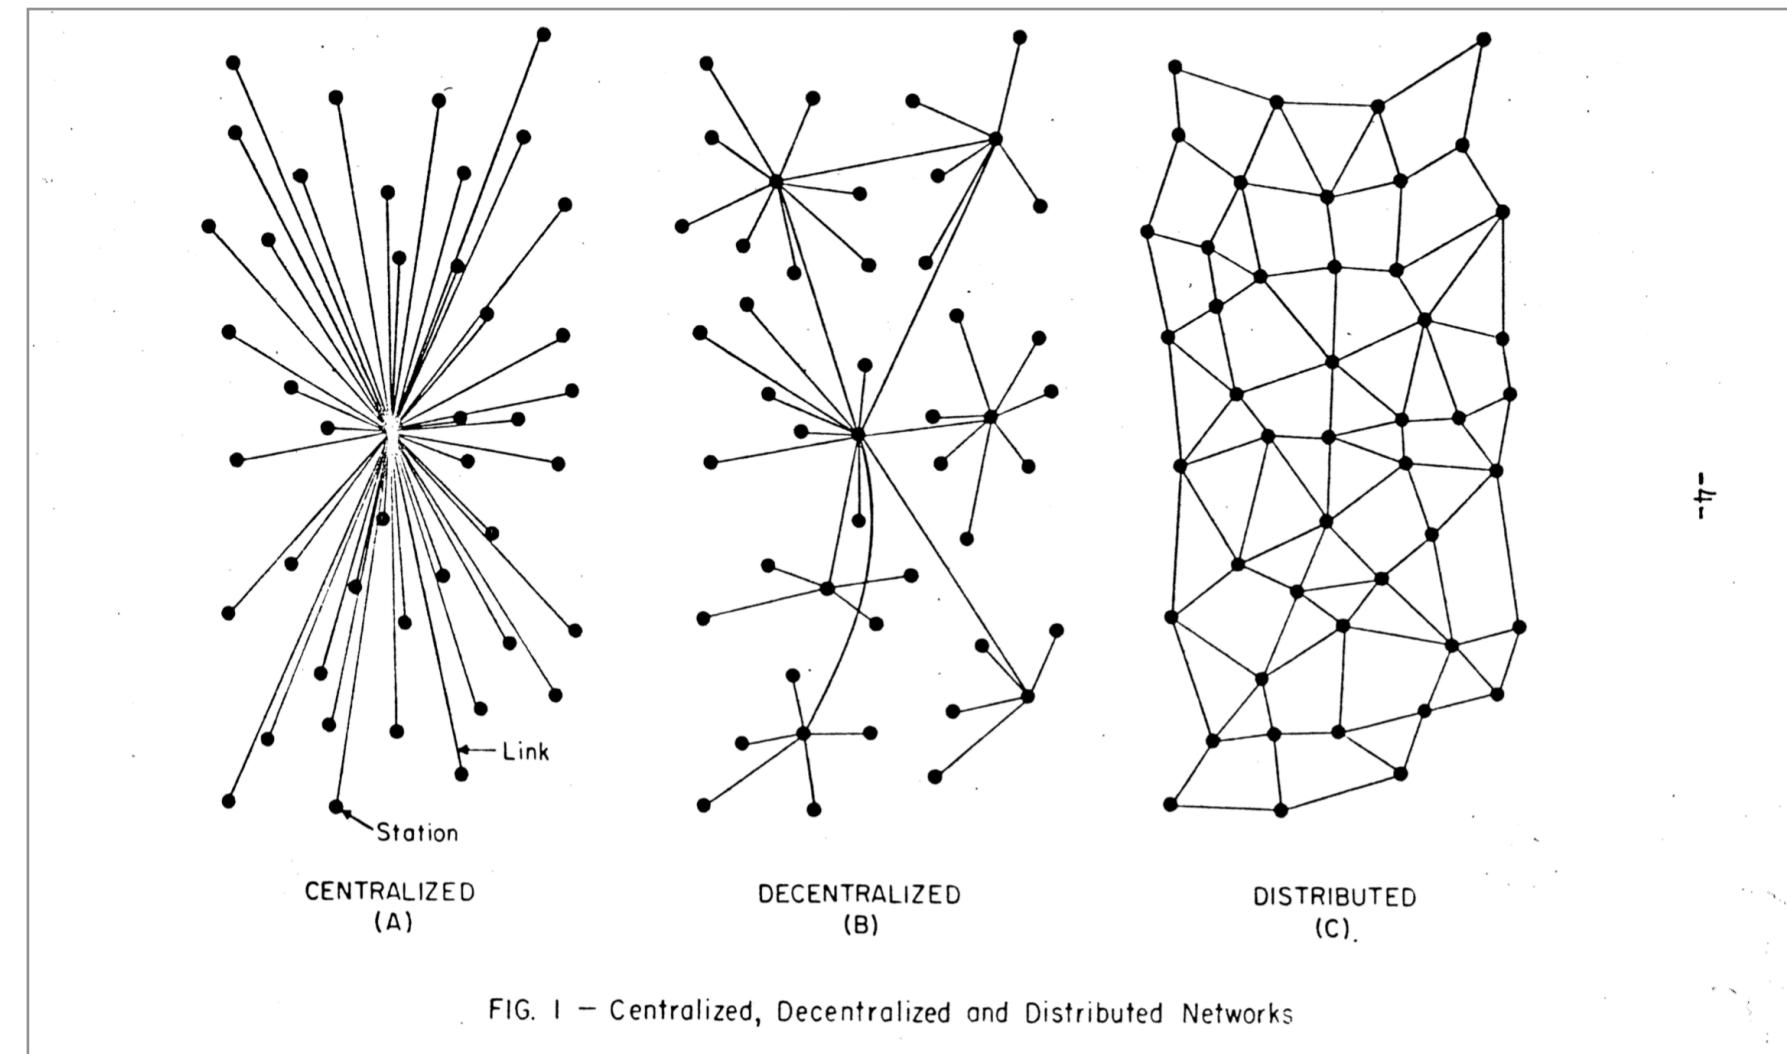
INTRODUCTION
The previous paper** described how redundancy of coding can be used to build efficient digital data links out of transmission links of variable and often less than presently useful quality. An arbitrarily low over-all error rate can be purchased with a modest redundancy of coding and clever terminal equipment. But even links with low error rates can have less than perfect reliability.
We should like to extend the remarks of the previous paper and address ourselves to the problem of building

*Any views expressed in this paper are those of the author. They should not be interpreted as reflecting the views of The RAND Corporation or the official opinion or policy of any government, organization or private research sponsor. Papers are reproduced by The RAND Corporation as a courtesy to members of its staff.
This paper was prepared for presentation at the First Congress on Information Systems Sciences, sponsored by The MITRE Corporation and the USAF Electronic Systems Division, November, 1962.
The writer is indebted to John Bower for his suggestions that switching and store-and-forward system can be described as a model of a postman sitting at a switchboard. Programming assistance provided by Sharla Boehm, John Derr, and Joseph Smith is gratefully acknowledged.

Assuming symmetrical bi-directional links, the postman can infer the "best" paths to transmit mail to any station merely by looking at the cancellation time or the equivalent handover number tag. If the postman sitting in the center of the United States received letters from San Francisco, he would find that letters from San Francisco arriving from channels to the west would come in with later cancellation dates than if such letters had

HOT-POTATO HEURISTIC ROUTING DOCTRINE
To achieve real-time operation it is desirable to respond to change in network status as quickly as possible so we shall seek to derive the network status information directly into each message block.
Each standardized message block contains a "to" address, a "from" address, a handover number tag, and error detecting bits together with other housekeeping data. The message block is analogous to a letter. The "from" address is equivalent to the return address of the letter.
The handover number is a tag in each message block set to zero upon initial transmission of the message block into the network. Every time the message block is passed on, the handover number is incremented. The handover number tag on each message block indicates the length of time in the network or path length. This tag is somewhat analogous to the cancellation date of a conventional letter.

INDUCTIVE DETERMINATION OF BEST PATH
Assuming symmetrical bi-directional links, the postman can infer the "best" paths to transmit mail to any station merely by looking at the cancellation time or the equivalent handover number tag. If the postman sitting in the center of the United States received letters from San Francisco, he would find that letters from San Francisco arriving from channels to the west would come in with later cancellation dates than if such letters had



Thanks to Paul Barlan et al.

the Internet infrastructure is extremely **resilient**

ability to route packets around failures
(not necessarily fast though)

Thanks to Paul Barlan et al.
the Internet infrastructure is **extremely resilient**

... at least when it comes to **thermonuclear wars**...

Thanks to Paul Barlan et al.
the Internet infrastructure is **extremely resilient**

... at least when it comes to thermonuclear wars...

but how does it fare against, say...

Thanks to Paul Barlan et al.
the Internet infrastructure is **extremely resilient**

... at least when it comes to thermonuclear wars...

but how does it fare against, say... **us humans?**



Login

Gear Gaming Entertainment Tomorrow The Buyer's Guide Video Reviews US Edition



Google accidentally broke the internet throughout Japan

A mistake led to internet outages for about half of the country.



Mallory Locklear, @mallorylocklear
08.28.17 in Internet

22
Comments

1815
Shares



Sponsored Links



27 August 2017

Someone in Google made a configuration mistake
which caused Google Chicago to wrongly advertise
160k IP prefixes to its neighbors.

Someone in Google made a configuration mistake which caused Google Chicago to wrongly advertise 160k IP prefixes to its neighbors.

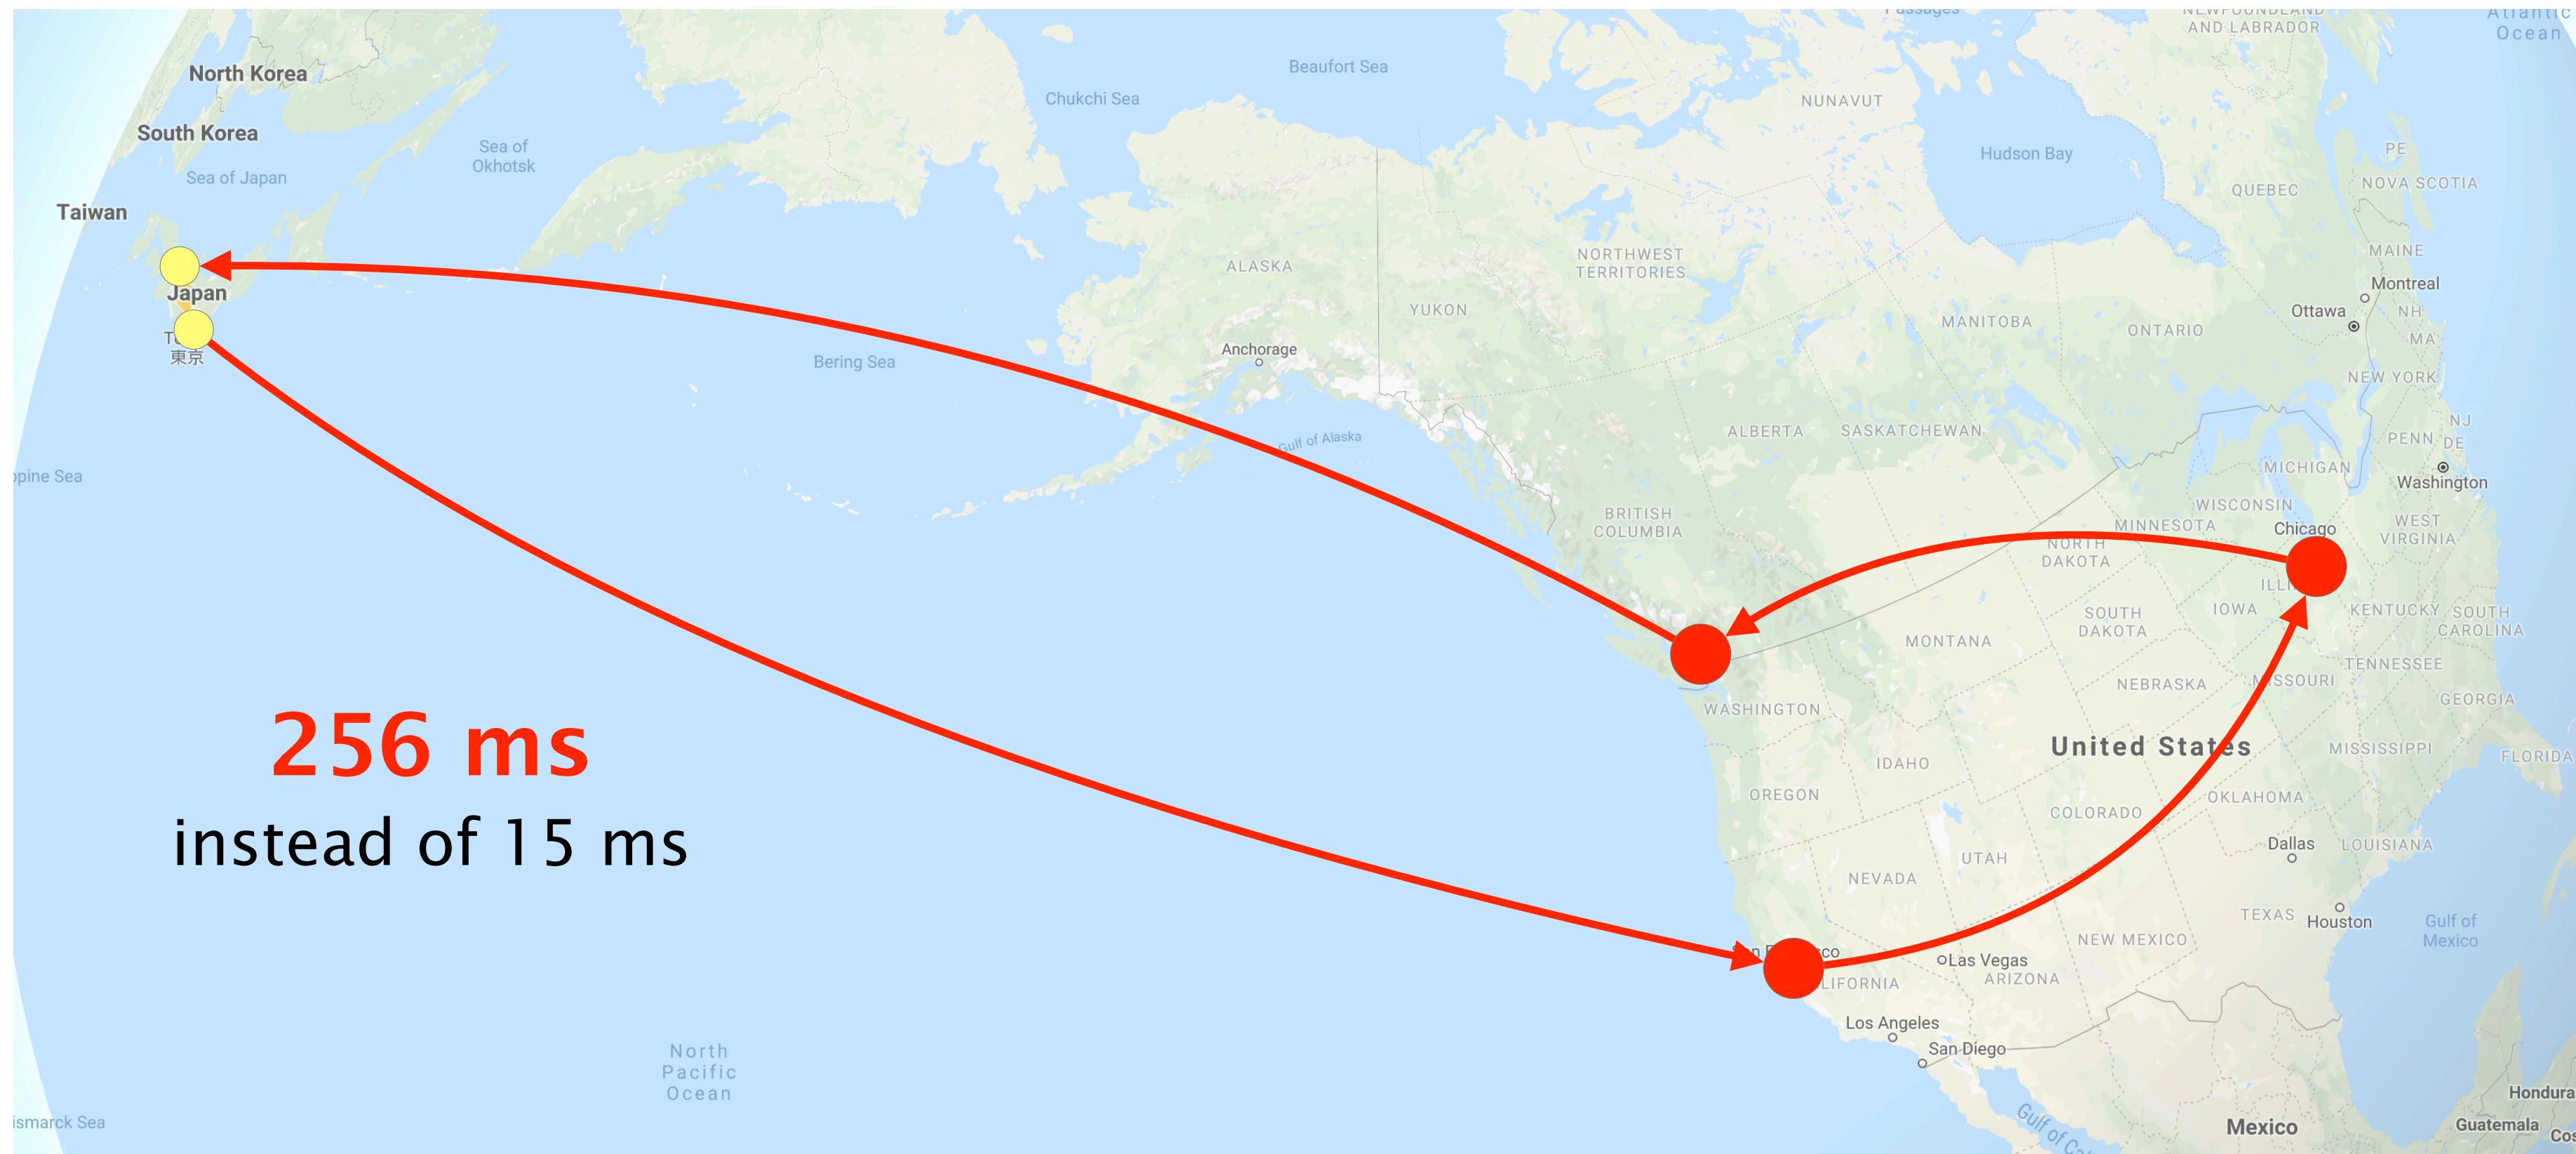
These advertisements propagated in the Internet and got picked up by Japanese giants (IIJ and KDDI) which started to direct local Japanese traffic to Google.

Tokyo to Nagoya (normally)



Tokyo to Nagoya via... Chicago?!

A 17x increased latency!



Someone in Google made a configuration mistake which caused Google Chicago to wrongly advertise 160k IP prefixes to its neighbors.

These advertisements propagated in the Internet and got picked up by Japanese giants (NTT and KDDI) which started to direct local, Japanese traffic to Google.

The outage in Japan *only* lasted a couple of hours but was so severe that the country's ministries want carriers to report on what went wrong.

This is **far** from being an isolated event...

JUL 8, 2015 @ 03:36 PM

11,261 VIEWS

United Airlines Blames Router for Grounded Flights

'Configuration Error' Blamed for AWS Outage

By David Ramel ■ 08/12/2015

Amazon's massive AWS outage was caused by human error

One incorrect command and the whole internet suffers.

By Jason Del Rey | @DelRey | Mar 2, 2017, 2:20pm EST



The summer of network misconfigurations

CONNECTIVITY MANAGEMENT FIREWALL CHANGE MANAGEMENT

SECURITY RISK MANAGEMENT AND VULNERABILITIES

ICY MANAGEMENT

dfrey | Aug 11, 2016



Data Centre ▶ Networks

Level3 switch config blunder for US-wide VoIP blackout

Widespread internet outages affected Comcast, Spectrum, Verizon and AT&T customers

BY: CNN
POSTED: 11:17 PM Nov 6, 2017

Data Centre ▶ Networks

CloudFlare apologizes for Telia screwing you over

Unhappy about

By Kieren McCarthy in San Francisco

Facebook struggles to deal with epic outage



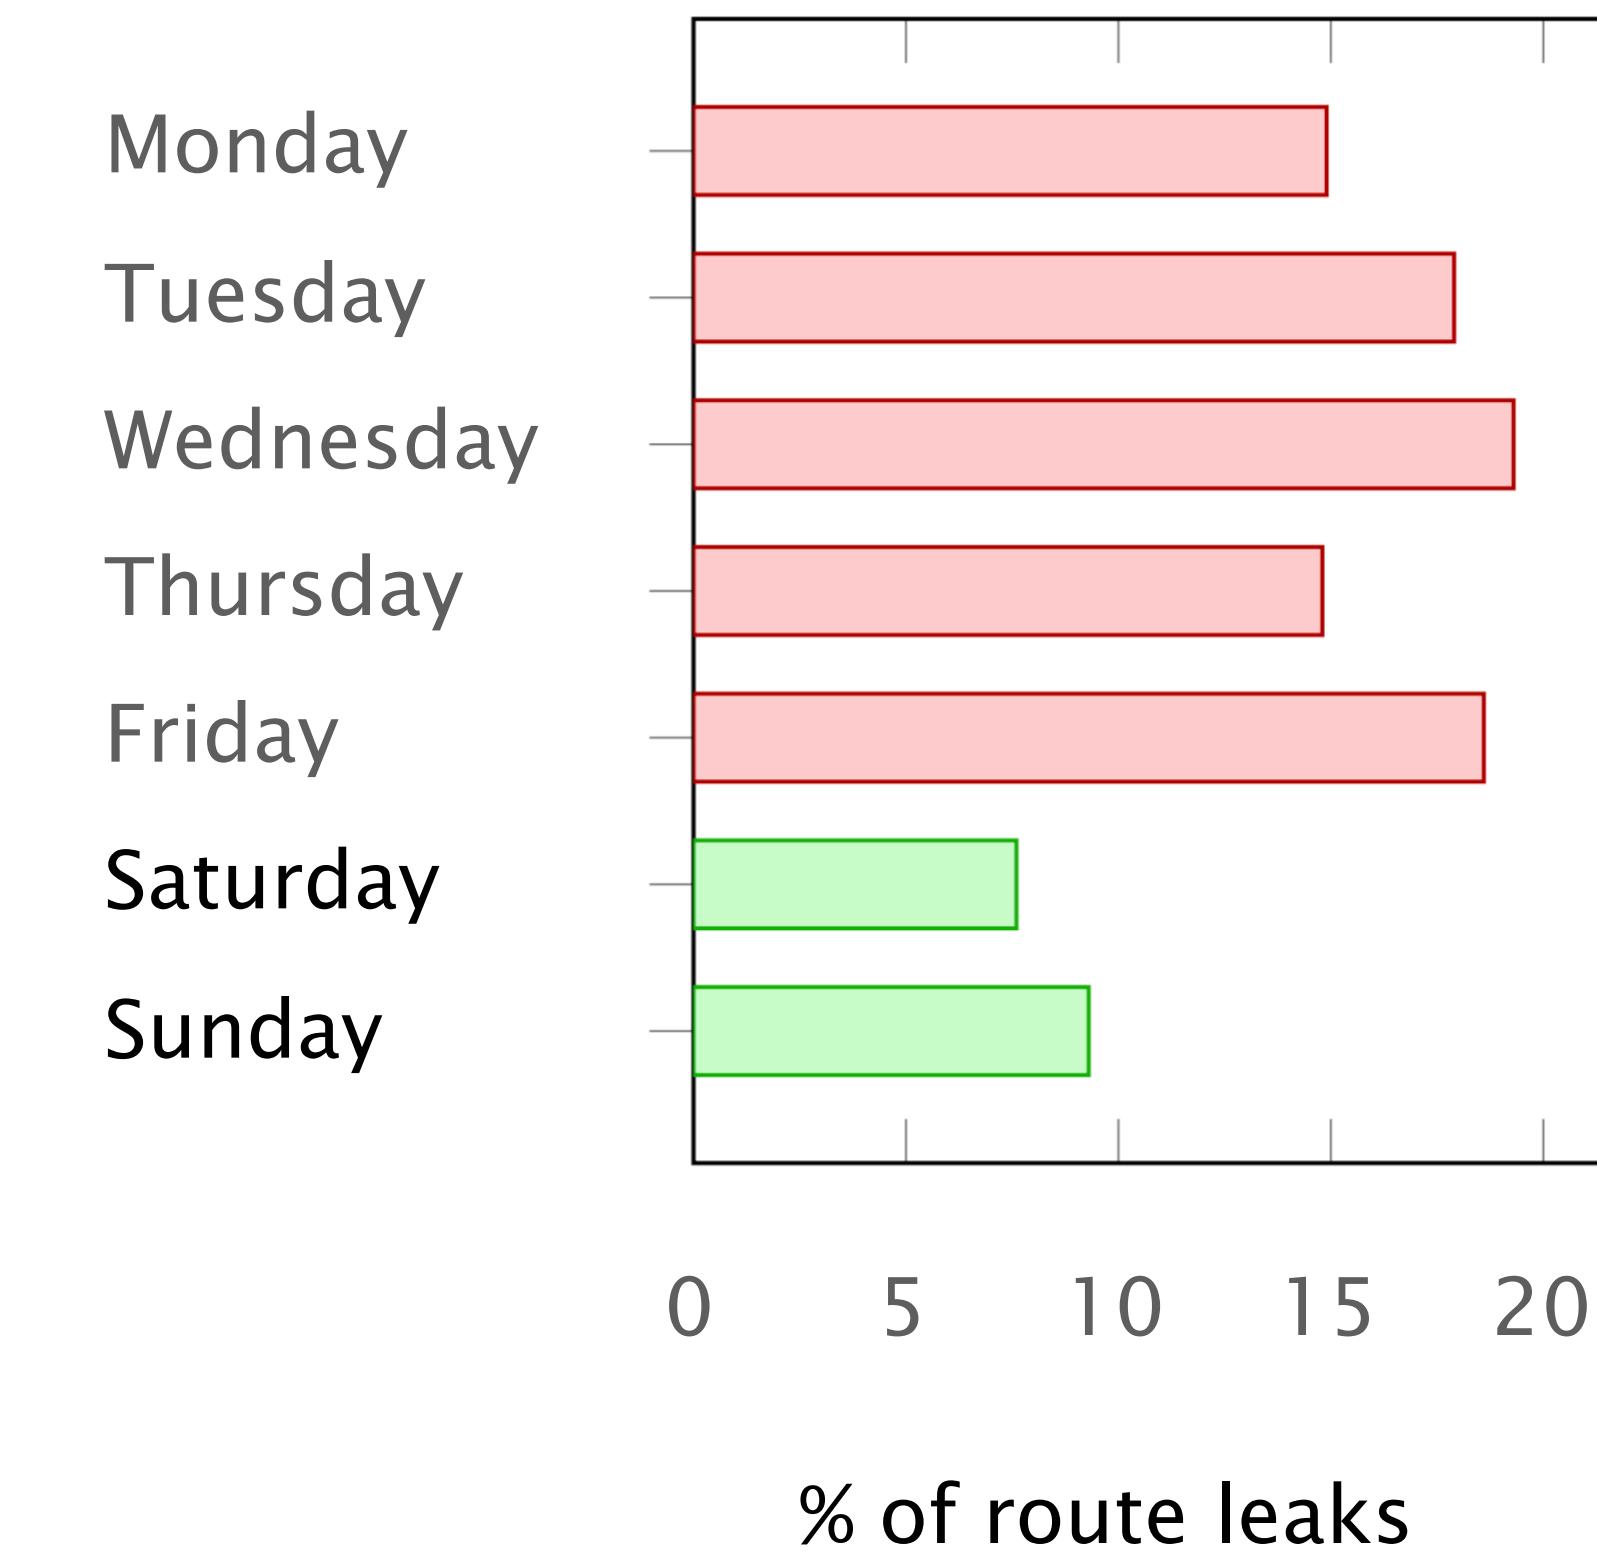
By Donie O'Sullivan and Heather Kelly, CNN Business

Updated 0654 GMT (1454 HKT) March 14, 2019

“Human factors are responsible
for 50% to 80% of network outages”

Juniper Networks, *What's Behind Network Downtime?*, 2008

A perhaps ironic consequence is that
the Internet works better during the week-ends



source: Job Snijders, 2008-2016

Need more proof?

Ask our students!

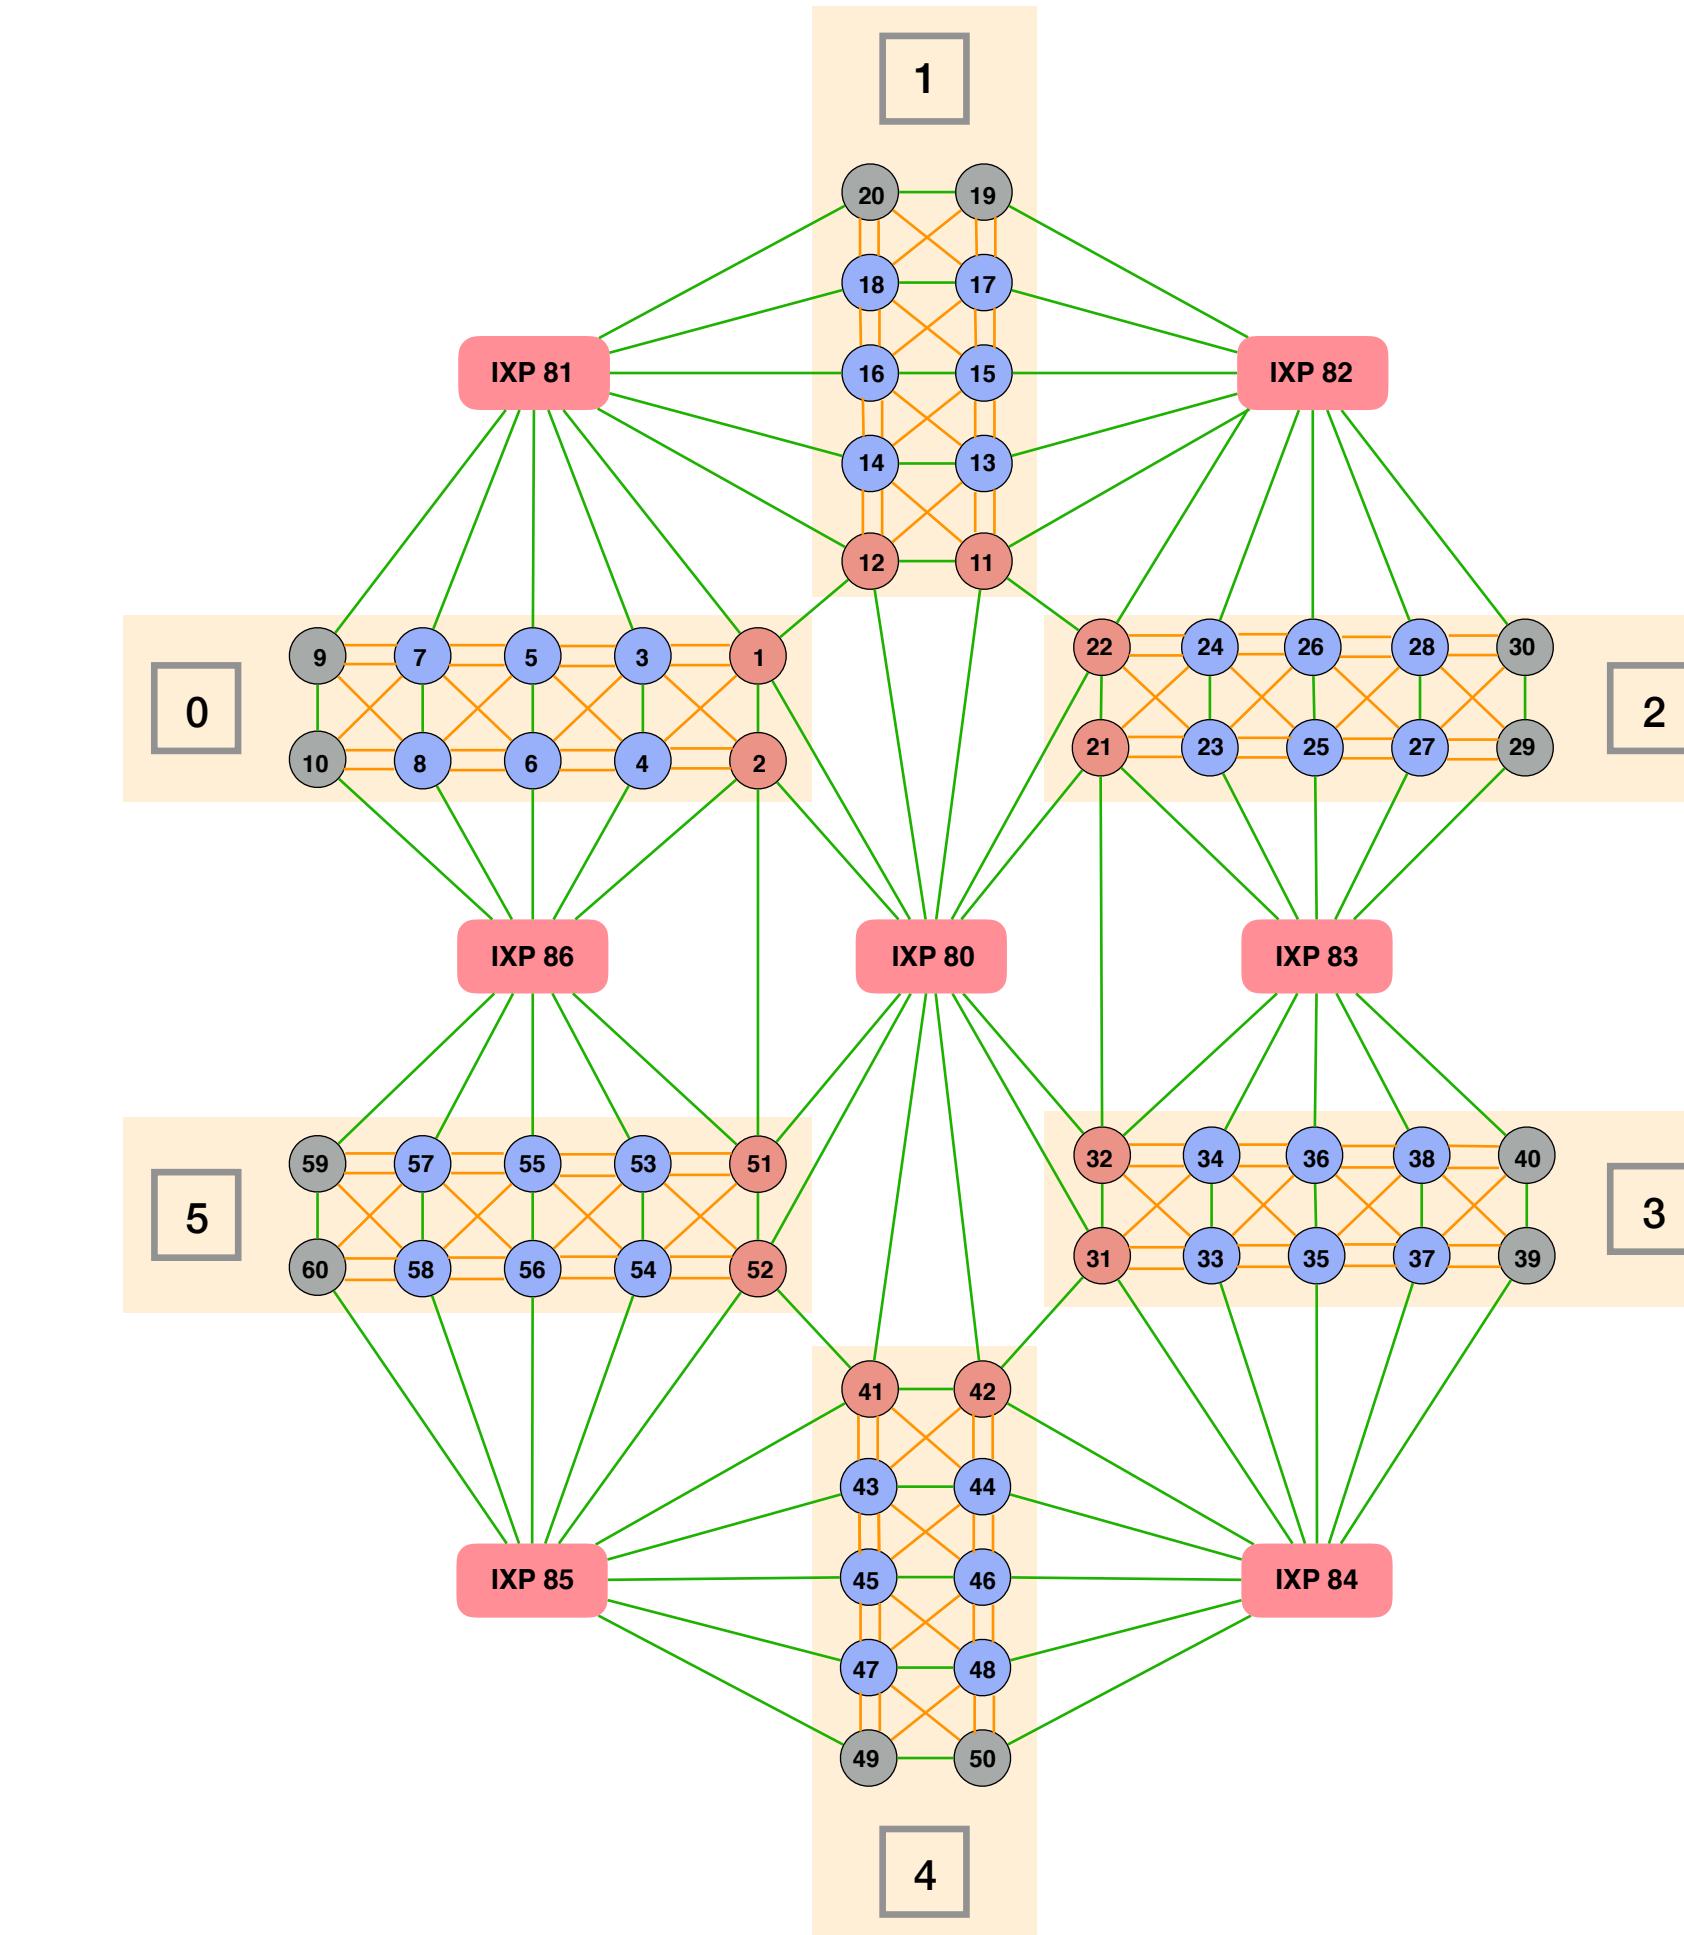
Yearly Communication Networks "Hackathon"



Each year, D-ITET students build and operate their very own Internet infrastructure

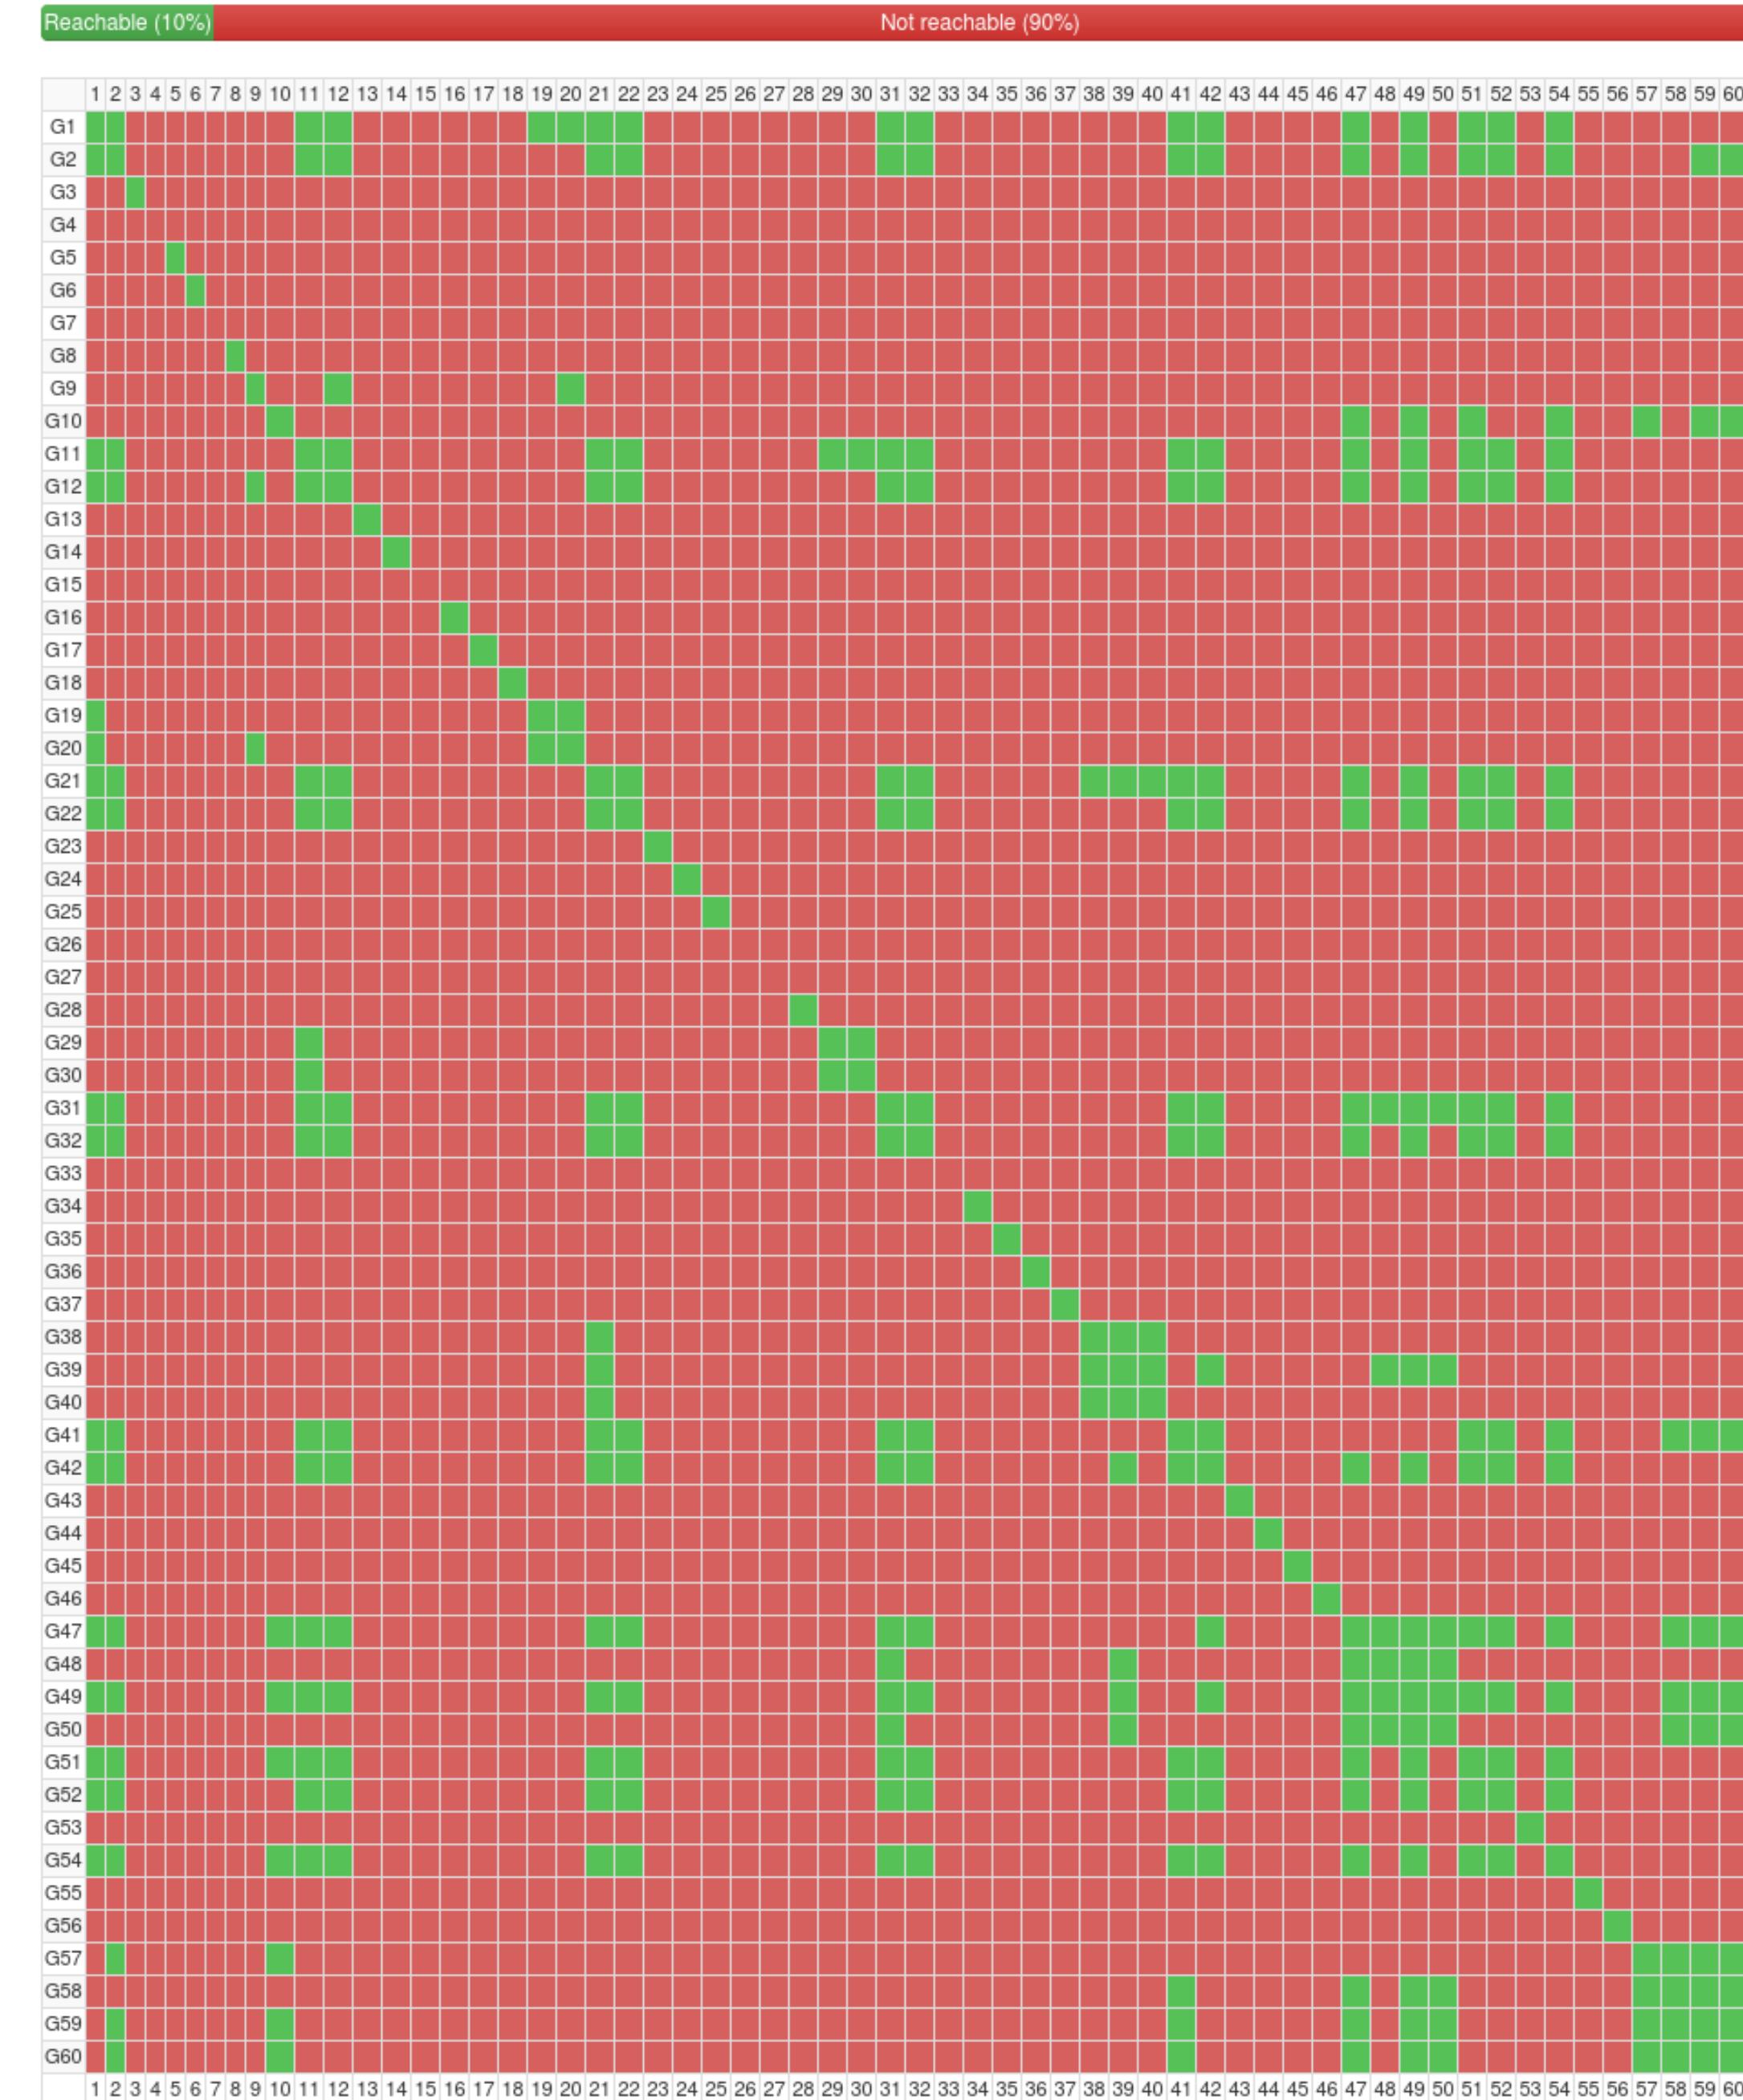
Internet topology 2019

~30 groups, 100 students



Internet Project: Connectivity Matrix

This connectivity matrix indicates the networks that each group can (green) or cannot reach (red). Matrix updated at Thu Apr 4 08:32:56 2019.



ITET's Internet connectivity

initial ~10%

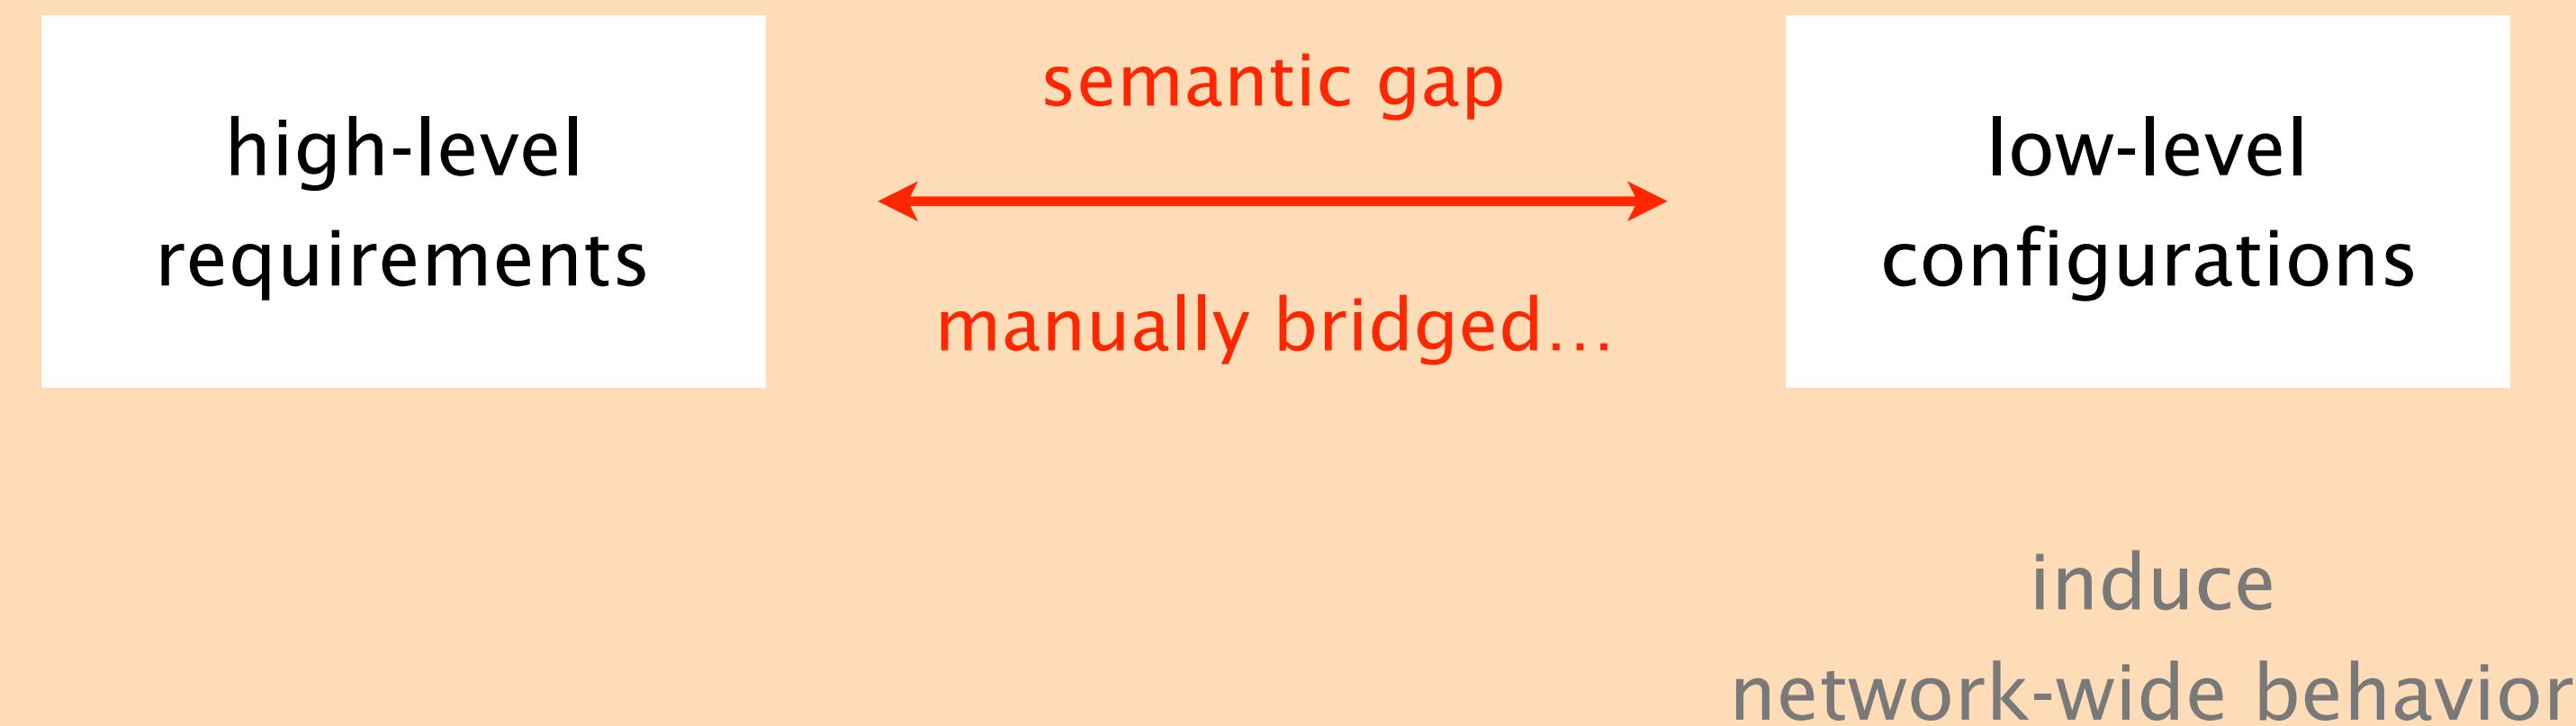
final ~97%

highest since 2016! 😊

Why ?!?

Configuring networks today is hard
because of a **fundamental semantic gap**

Configuring networks today is hard
because of a **fundamental semantic gap**



High-level requirements encode operational objectives

connectivity

guarantee

reachability

business/
performance

minimize

transit costs

congestion

convergence time

maximize

load-balancing

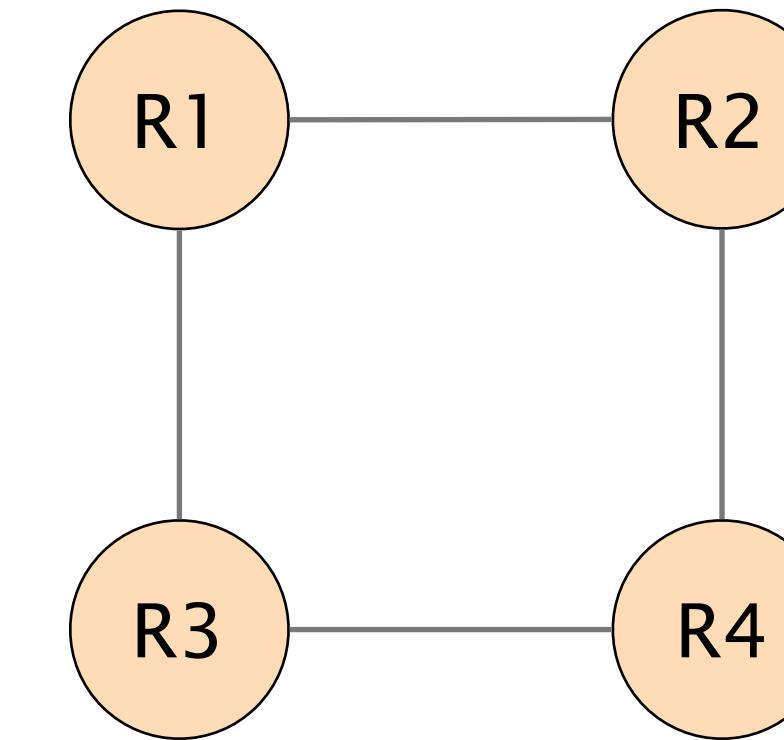
prevent

X from reaching Y

traffic crosses a firewall

Operators implement these requirements by
manually **configuring** their network devices

setting parameter values



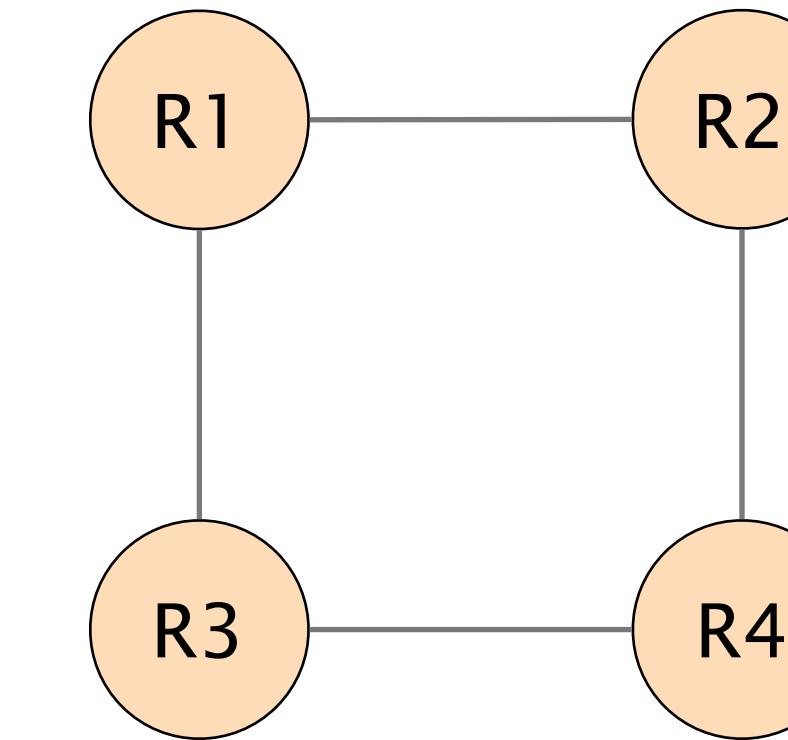
Operators implement these requirements by manually configuring their network devices

R1's configuration

```
id      BRUX
lo0    10.0.0.1
...
paramX  OSPF
weight (1,2) 10
weight (1,3) 20
```

R2's configuration

```
id      PARIS
lo0   10.0.0.2
...
paramX  OSPF
weight (2,1) 10
weight (2,4) 20
```



R3's configuration

```
id      LON
lo0   10.0.0.3
...
paramX  OSPF
weight (3,1) 10
weight (3,4) 20
```

R4's configuration

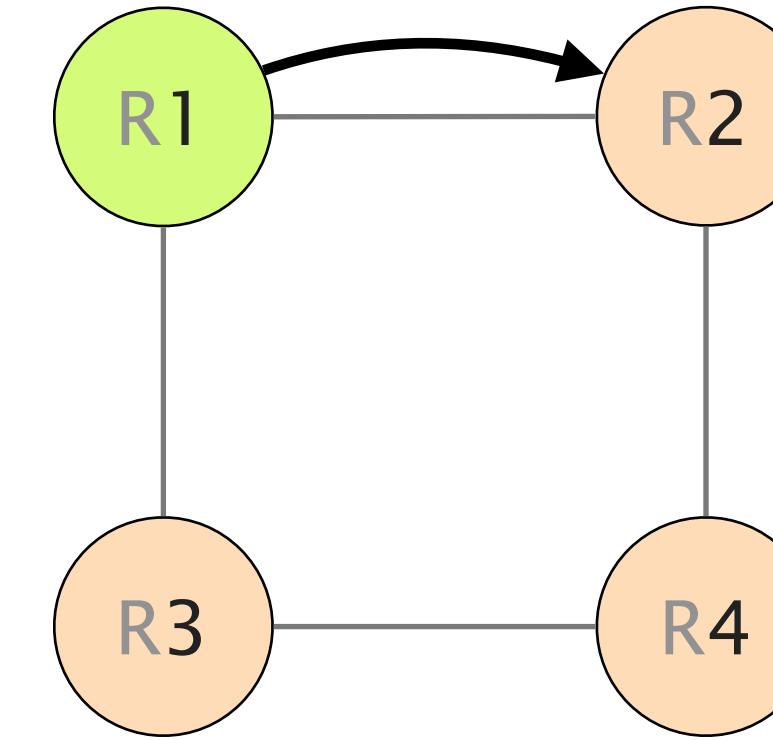
```
id      ZRH
lo0   10.0.0.4
...
paramX  OSPF
weight (4,1) 10
weight (4,2) 20
```

These parameters constraint the behavior of
the distributed routing protocols each node runs

R1's configuration

id	BRUX
lo0	10.0.0.1
...	
paramX	OSPF
weight (1,2)	10
weight (1,3)	20

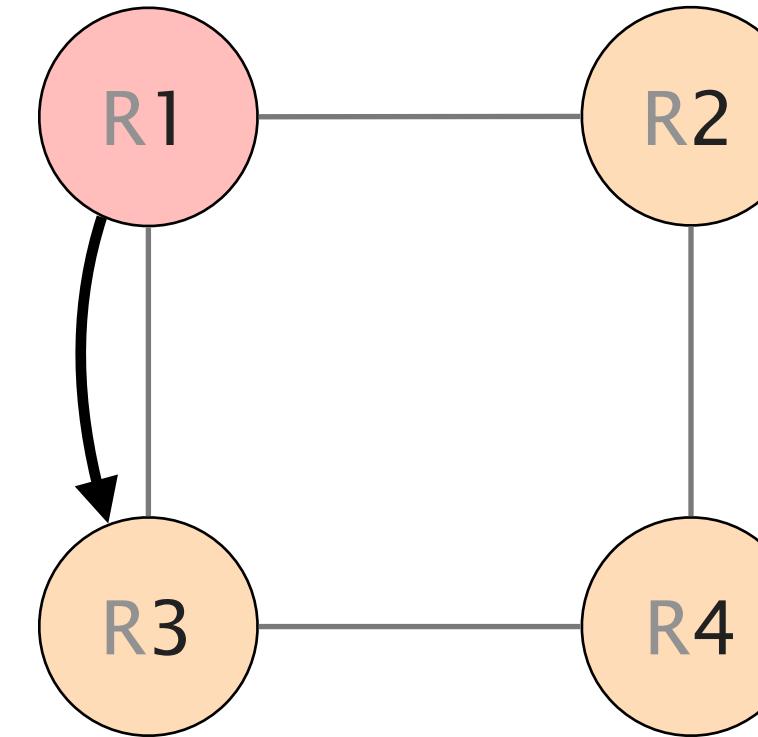
Eurovision
video traffic



R1's configuration

```
id      BRUX
lo0    10.0.0.1
...
paramX OSPF
weight (1,2) 100
weight (1,3) 20
```

Eurovision
video traffic



Defining the routing behavior of a large network
requires writing **millions of lines** of configuration

Configuration length

avg.	14k	LoC per device
max.	97k	

200+ devices, large European network

“Human factors are responsible
for 50% to 80% of network outages”

Juniper Networks, *What's Behind Network Downtime?*, 2008

“Human factors are responsible
for [REDACTED] of [REDACTED]”

“Human factors are responsible
for >90% of car accidents ”

NHTSA, National Motor Vehicle Crash Causation Survey, February 2015

Enters...
Self-Driving Car

Enters...

Self-Driving Car

Wikipedia

A vehicle that is capable of
sensing its environment and
moving with little or no human input

Enters... Self-Driving Car

Wikipedia

A vehicle that is capable of
sensing its environment and
moving with little or no human input

Promise

A drastic reduction in the number of (fatal) accidents

Enters...
Self-Driving Car

Enters...

Self-Driving

Enters...

Self-Driving Network

Enters...

Self-Driving Network

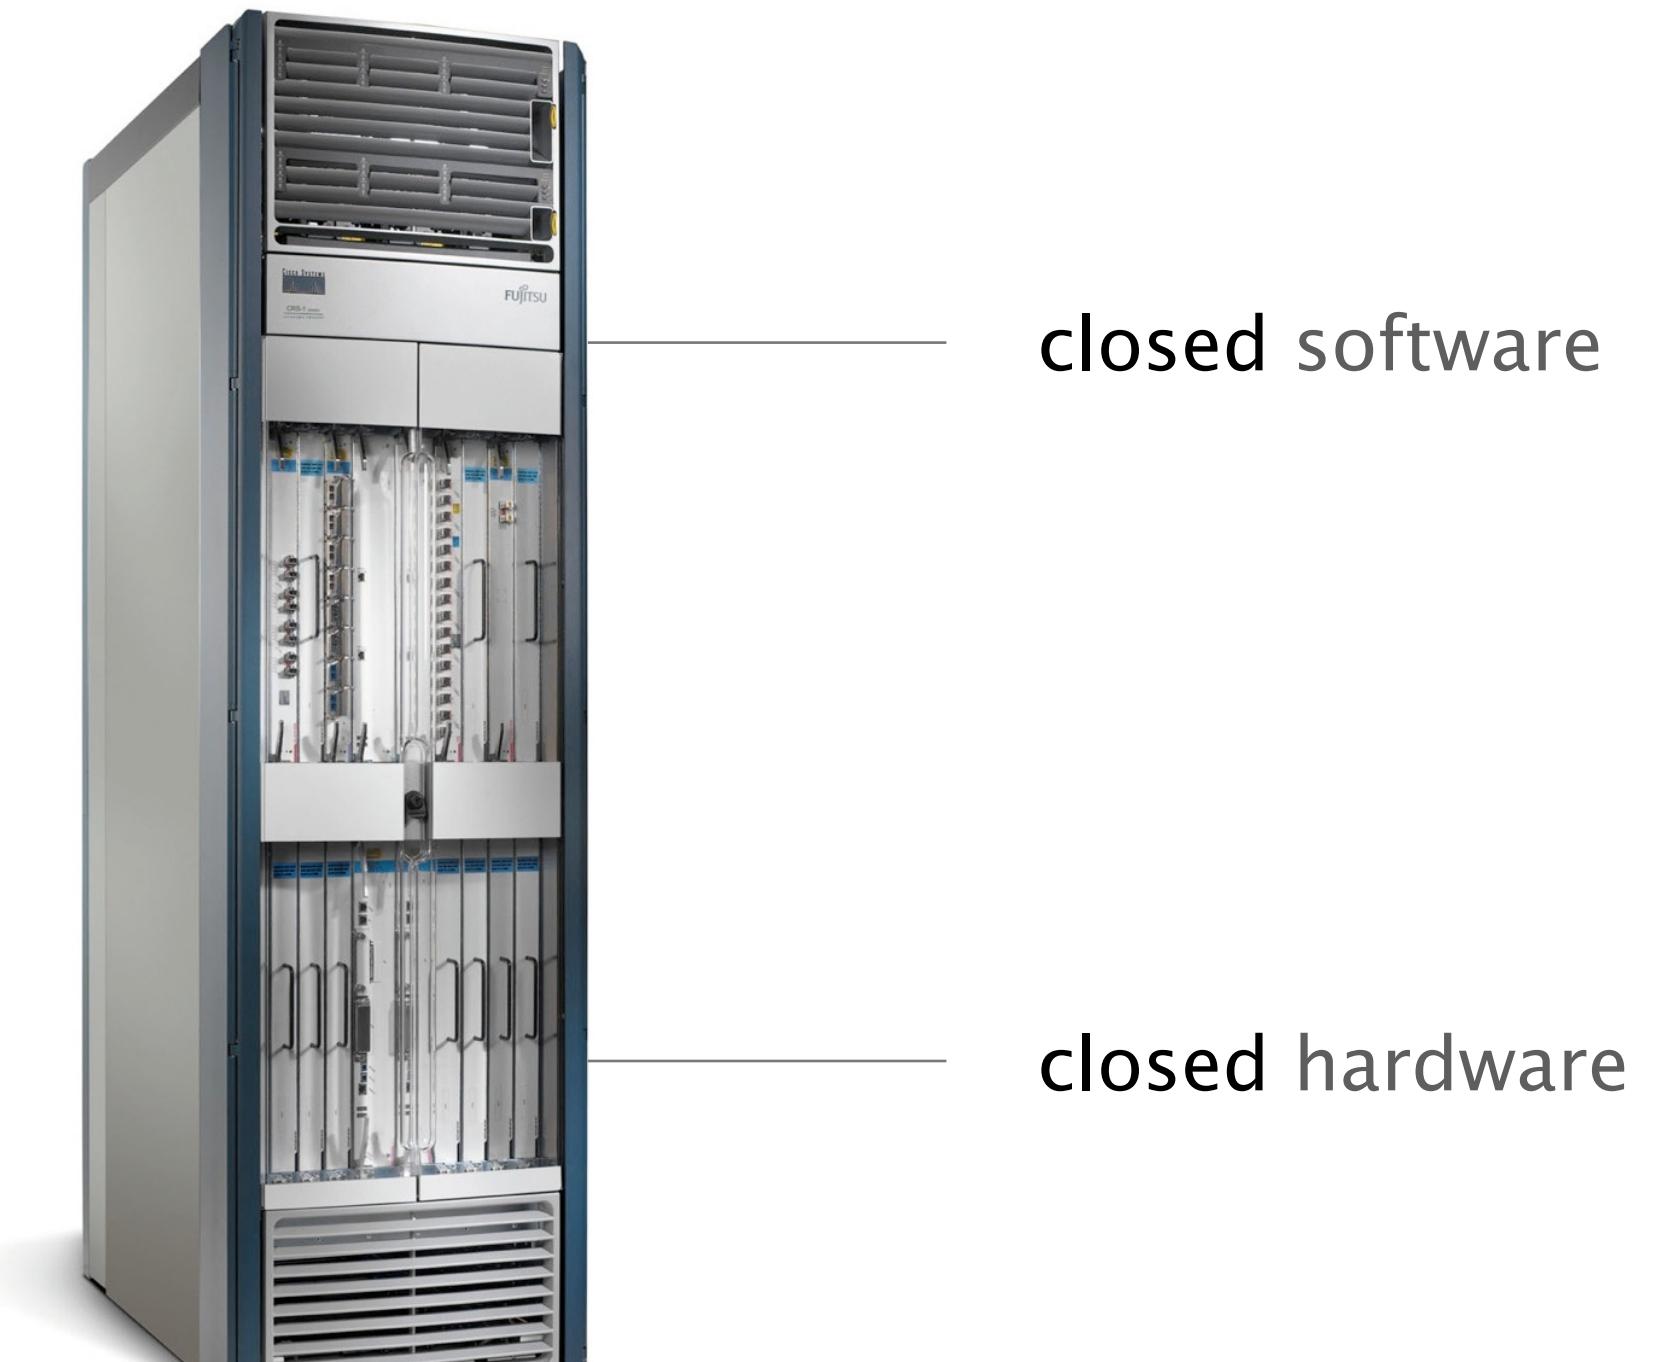
Definition A network that is capable of
 sensing its environment and
 adapting its behavior accordingly
 with little or no human input

Enters... Self-Driving Network

Definition A network that is capable of
 sensing its environment and
 adapting its behavior accordingly
 with little or no human input

Questions **How do we build and deploy such networks?**

Until recently, innovating in networks was hard because devices were completely locked down



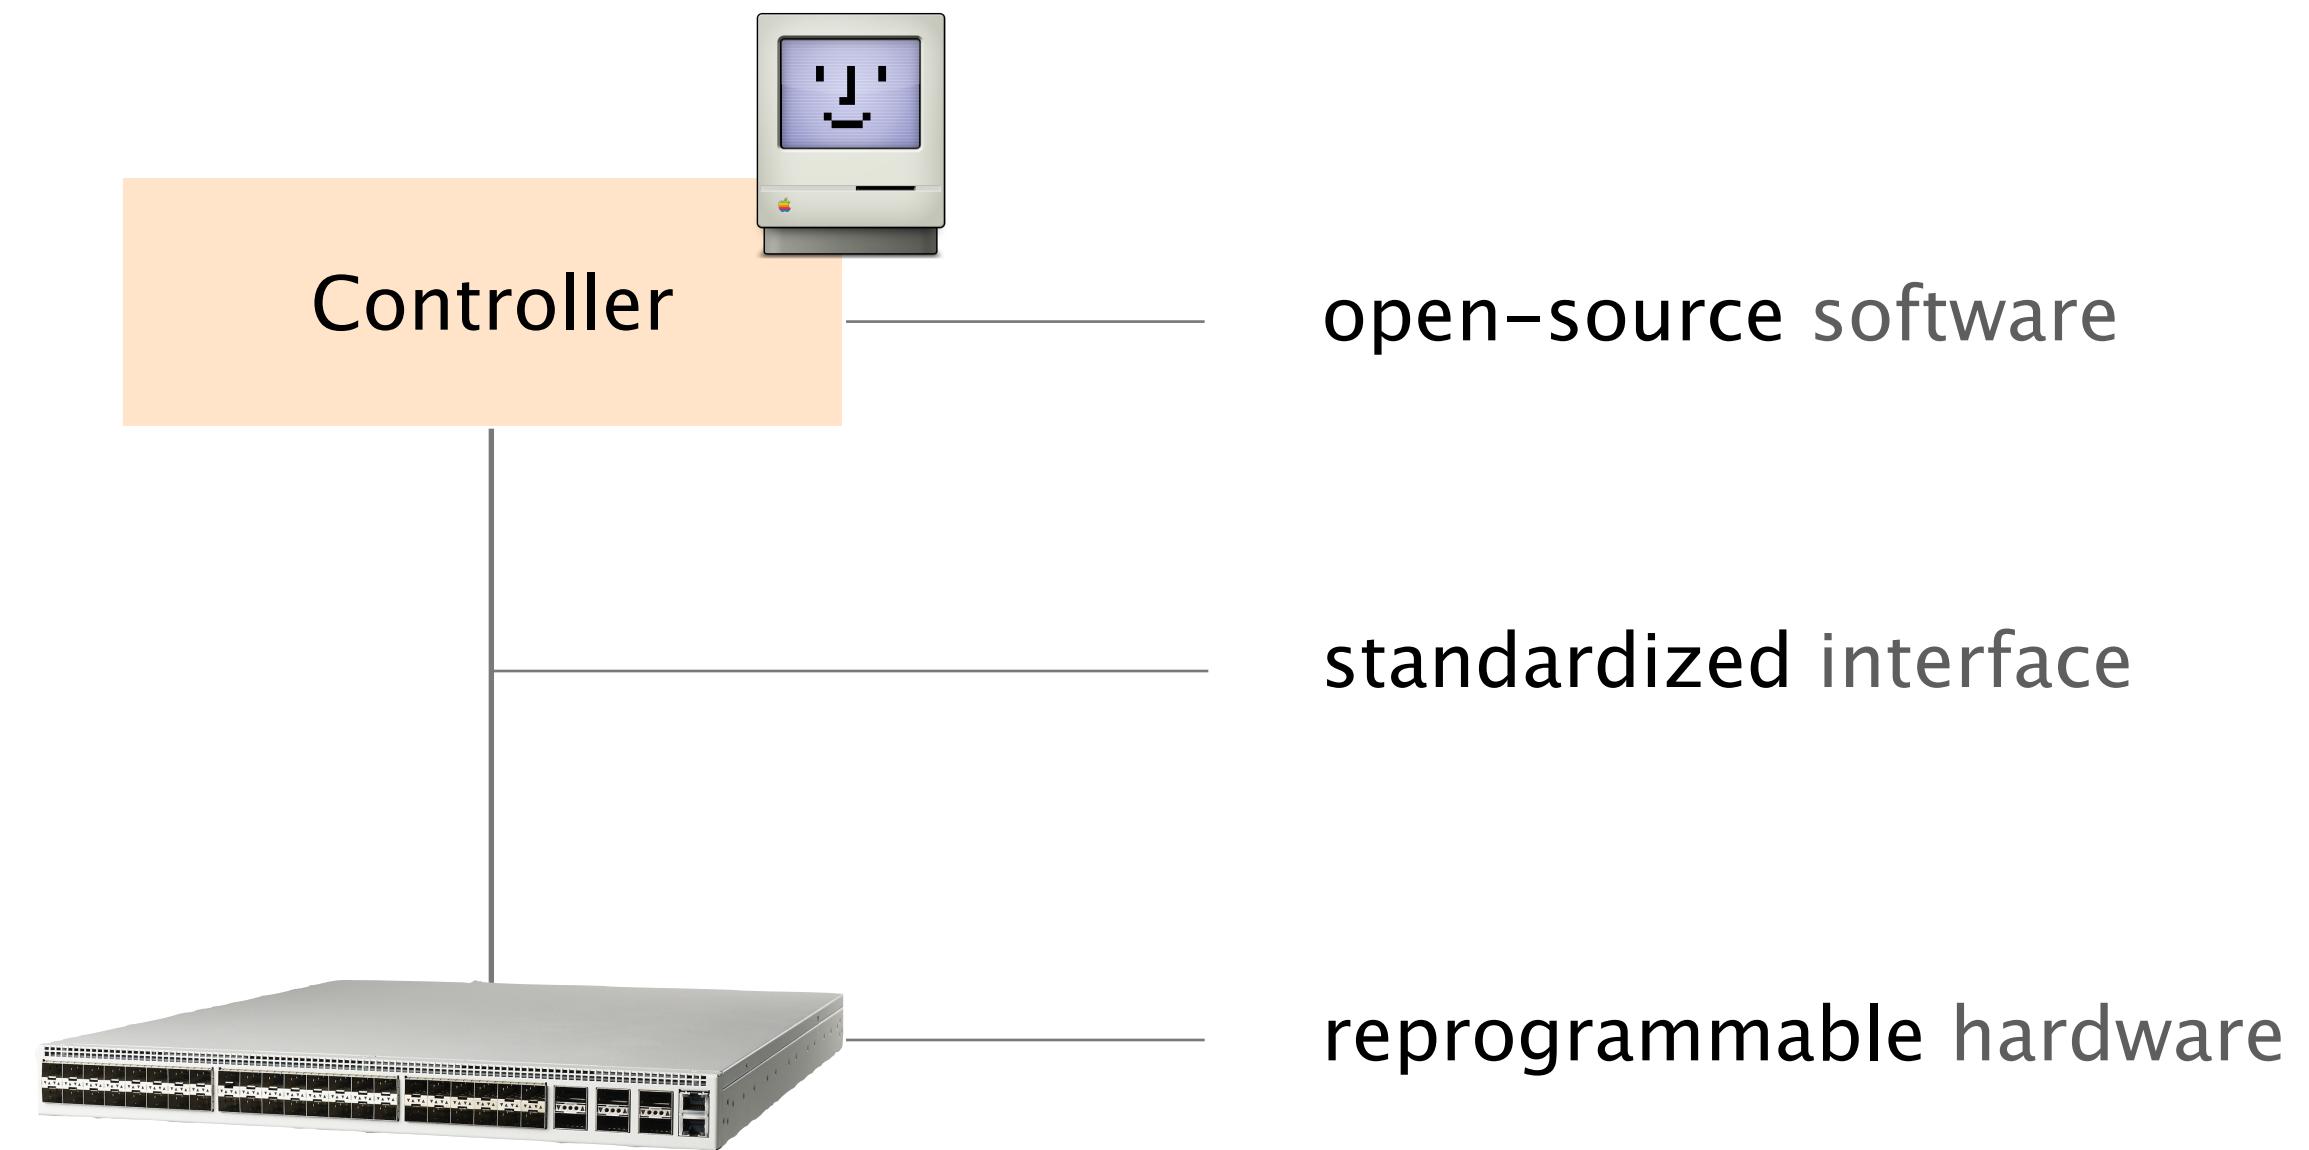
Cisco™ device

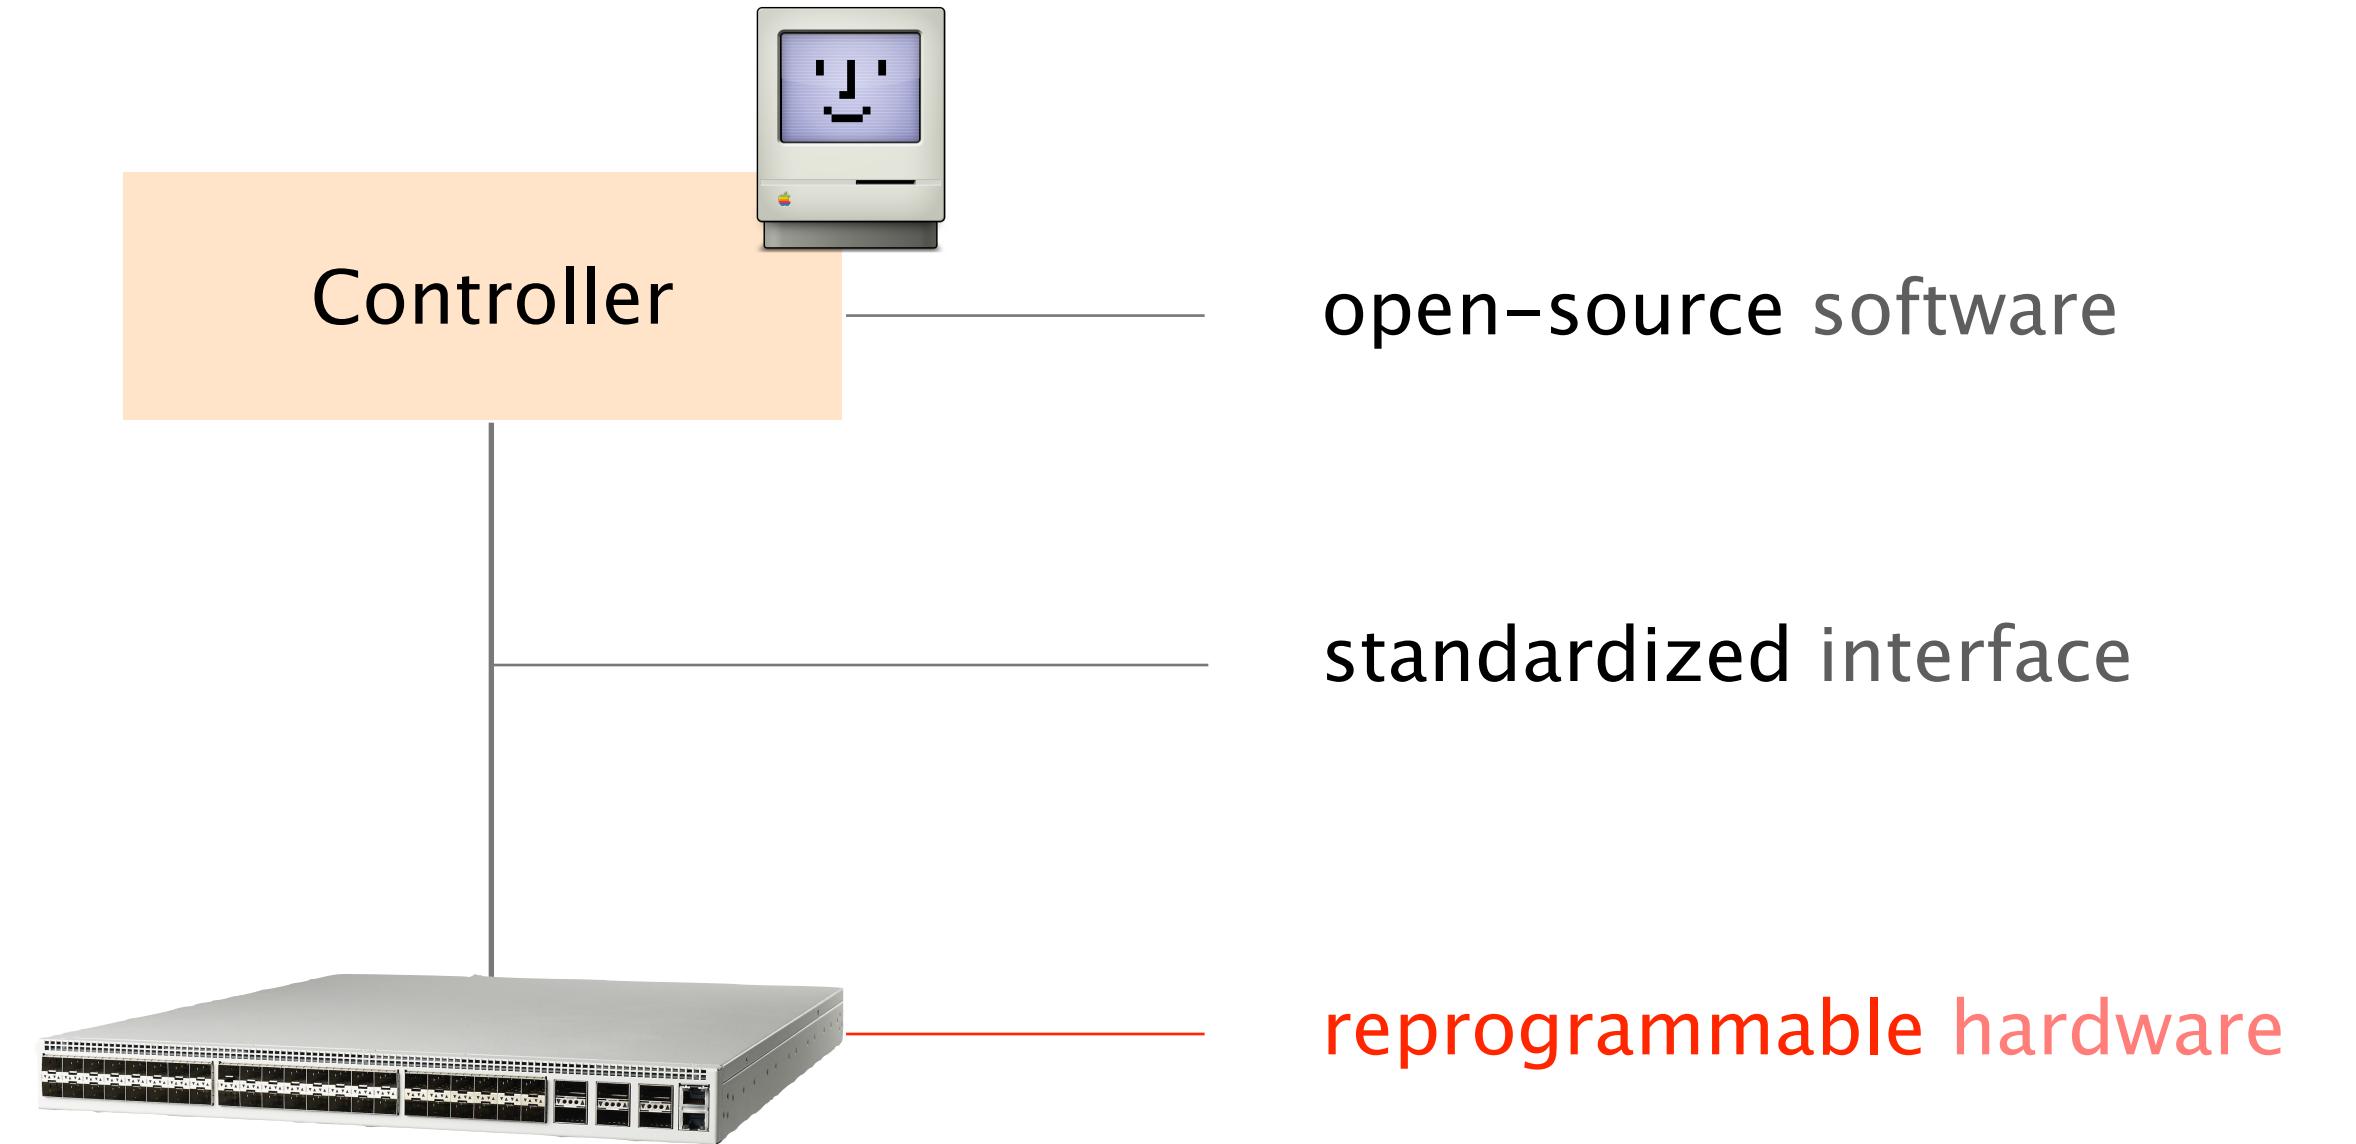
closed software

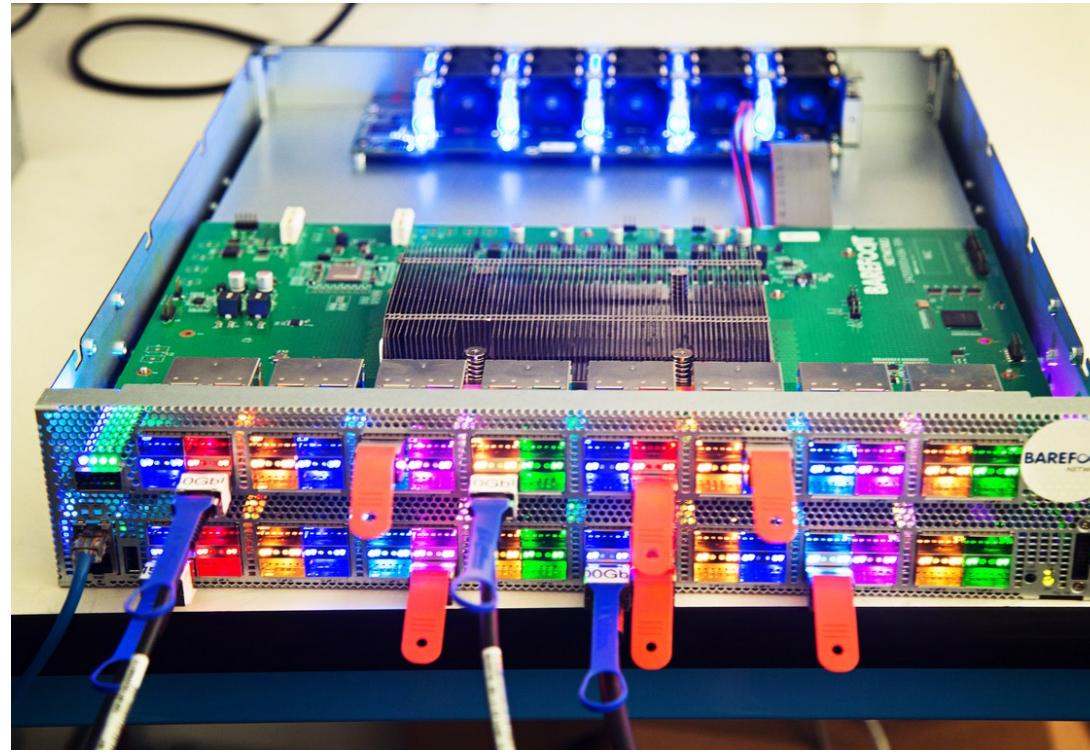
closed hardware

Things are changing though!

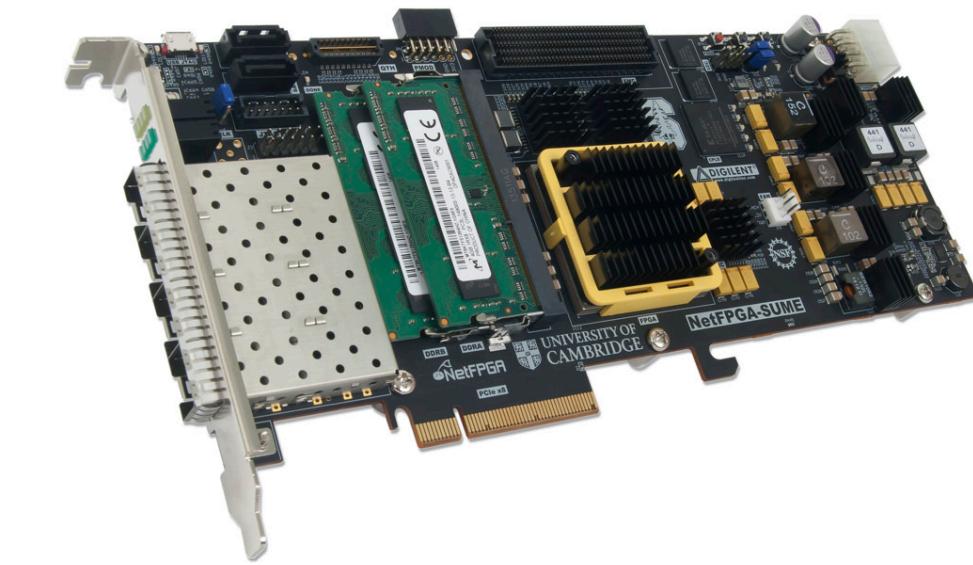
Networks are on the verge of a paradigm shift towards **deep programmability**



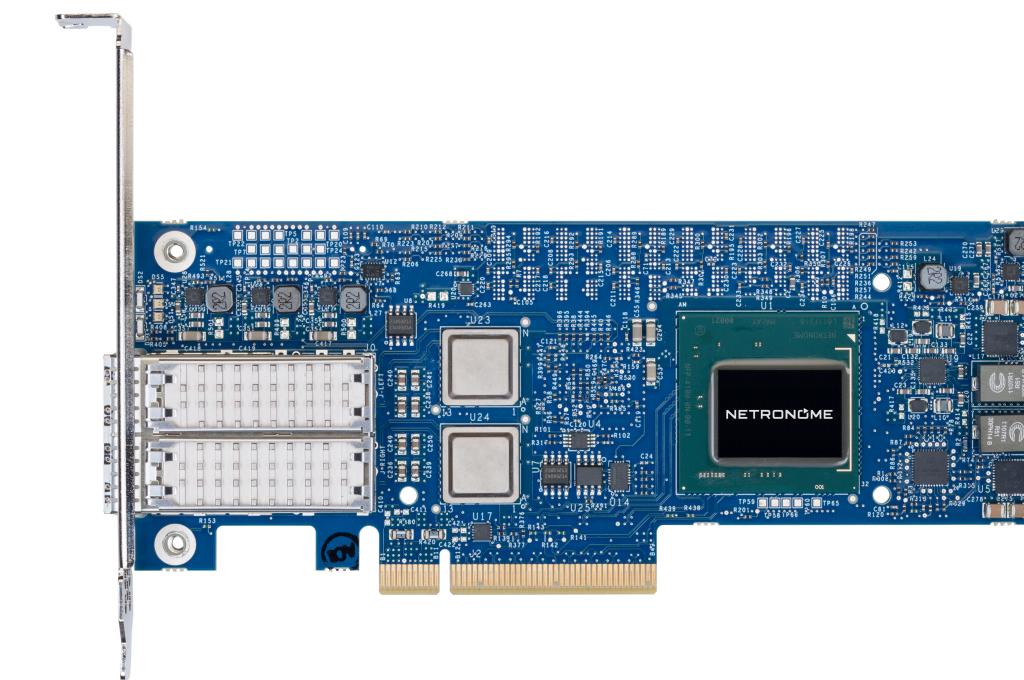




Barefoot Tofino 12.8 Tbps
programmable network switch

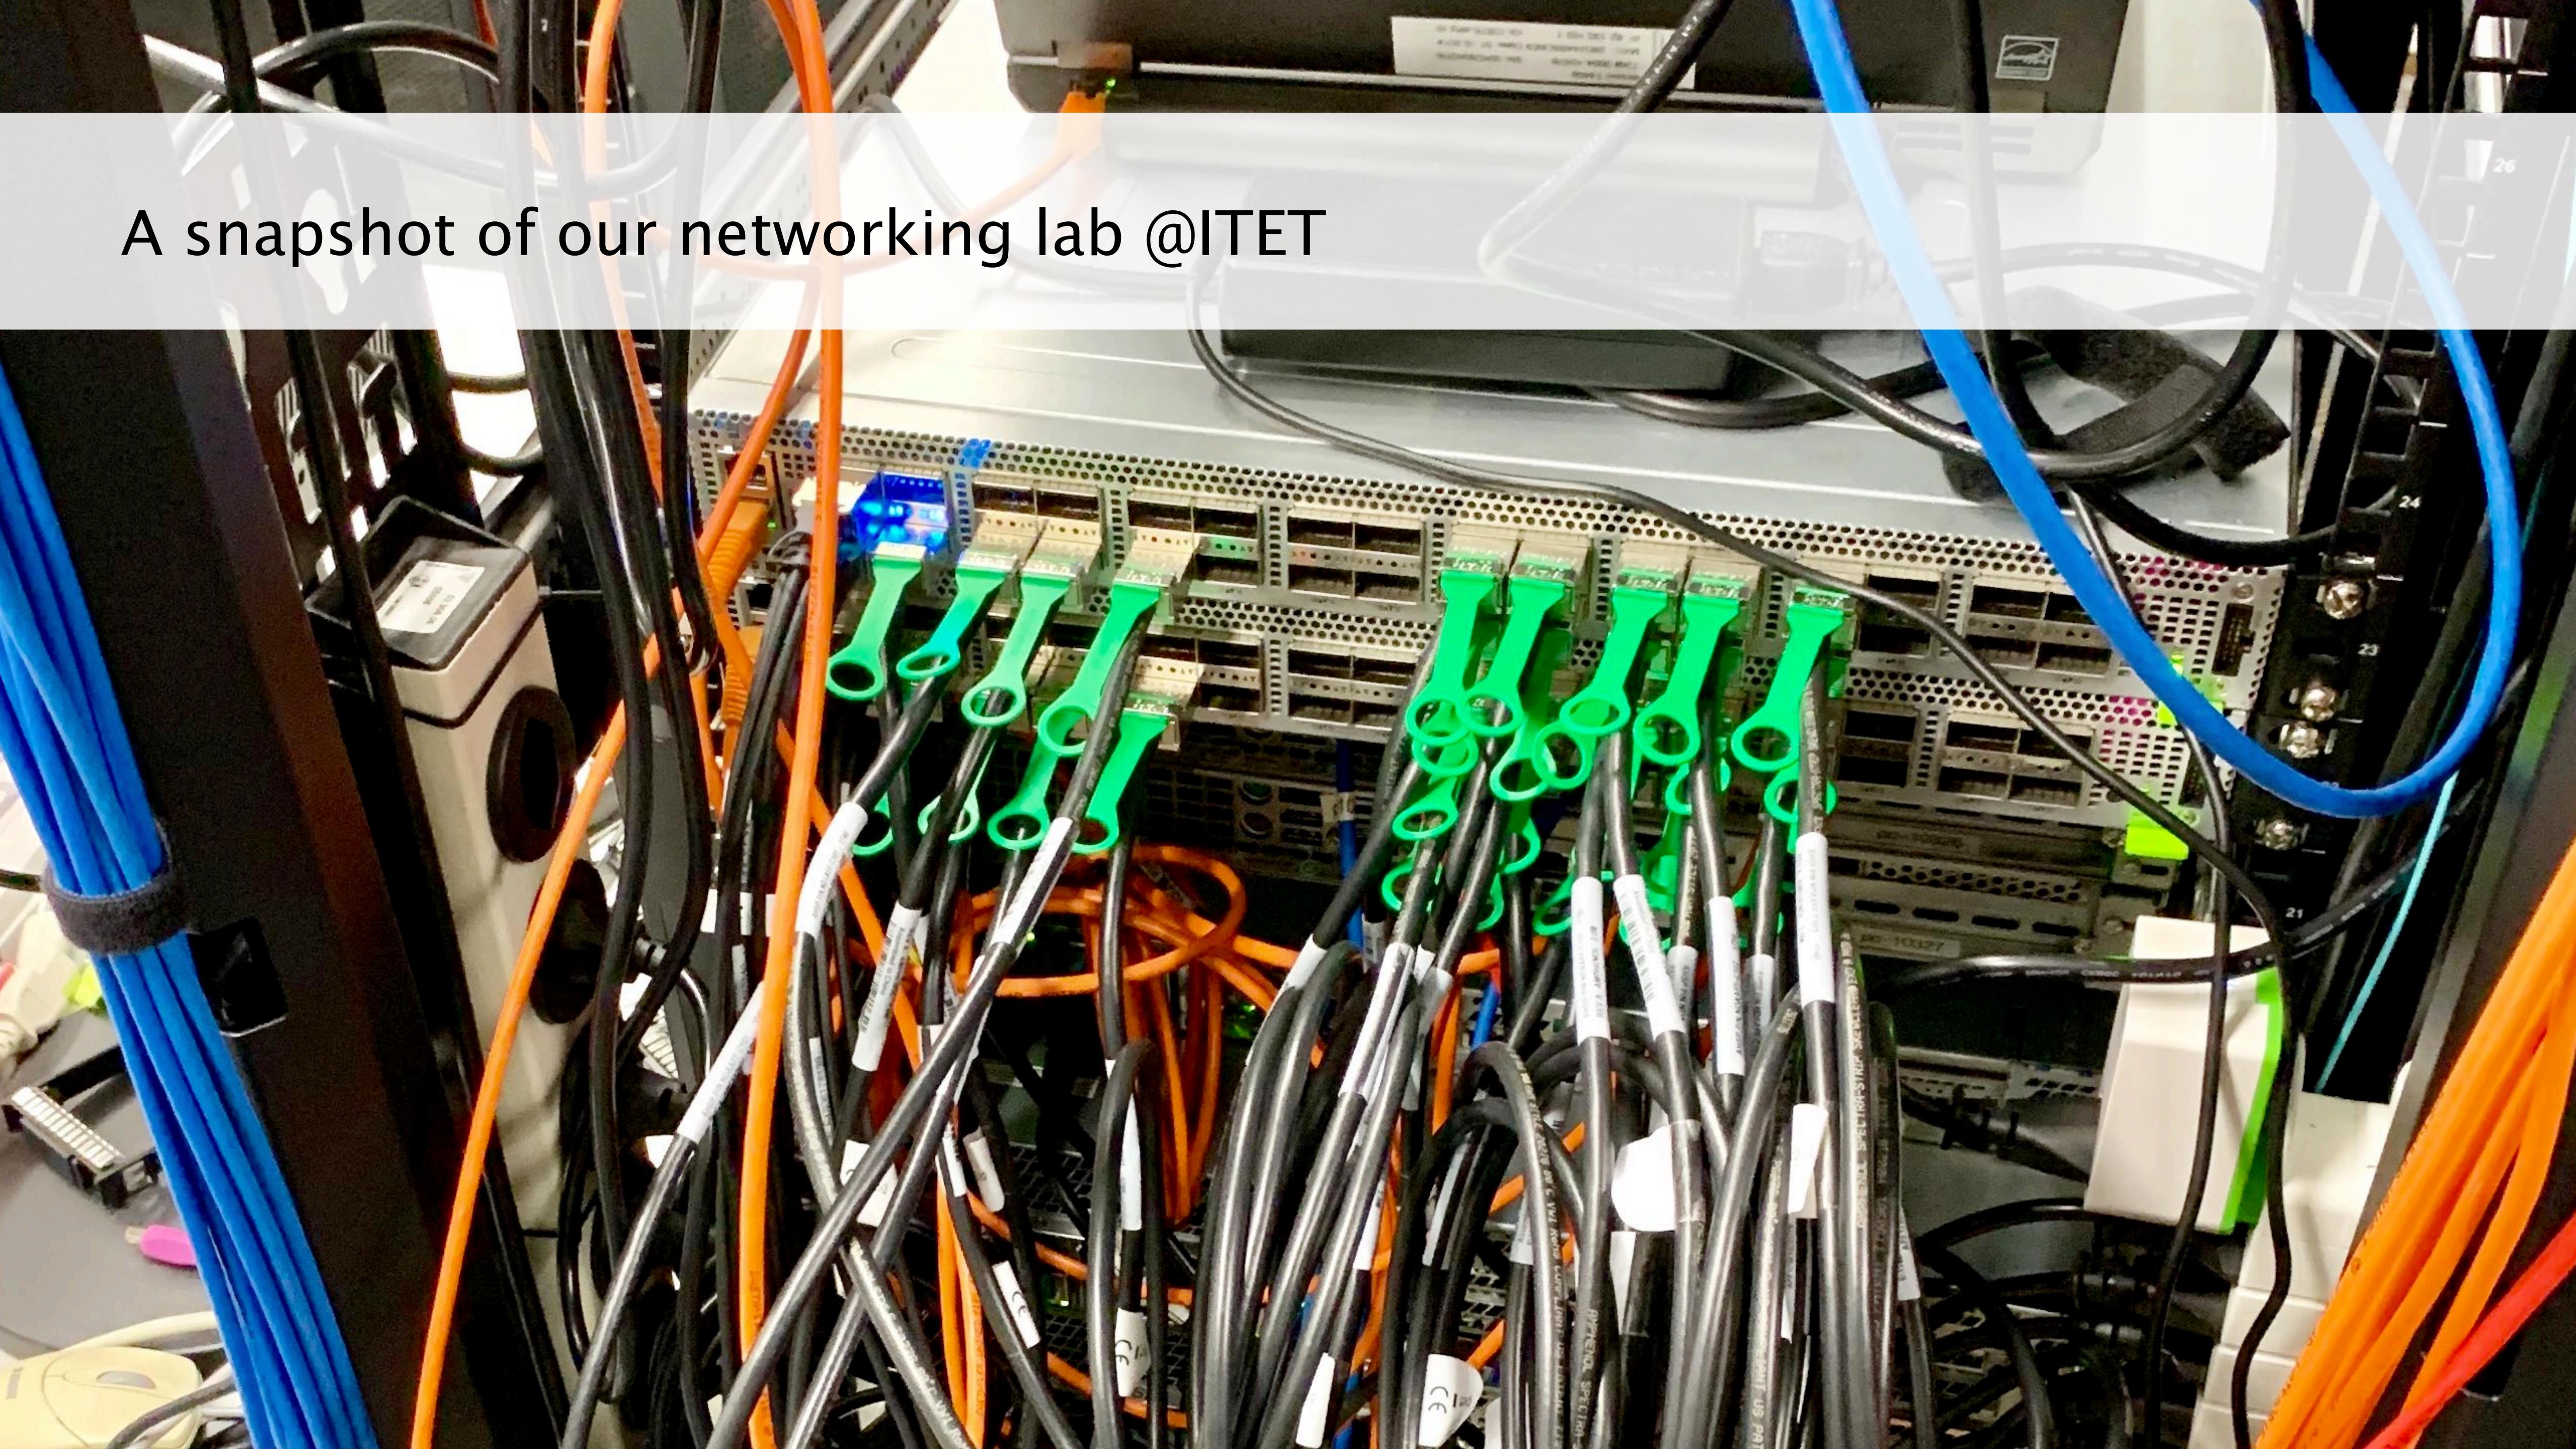


NetFPGA SUME
100 Gbps programmable network card



Netronome Agilio CX

A snapshot of our networking lab @ITET



Programmable networks can be made self-driving

Programmable devices can

- measure
- perform statistical inference
- adapt their forwarding decisions

at line rate, on a per-packet basis

Enters...

Self-Driving Network

Definition

A network that is capable of sensing its environment and adapting its behavior accordingly with little or no human input

Questions

How do we build and **deploy** such networks?

Levels of autonomy in self-driving cars

level 0	no automation	
1	driver assistance	human monitors the environment
2	partial automation	
3	conditional automation	system monitors human as fallback
4	high automation	
5	full automation	no more human

Self-Driving Networks

Breaking new ground in network automation



- 1 operator assistance
- 2 partial automation

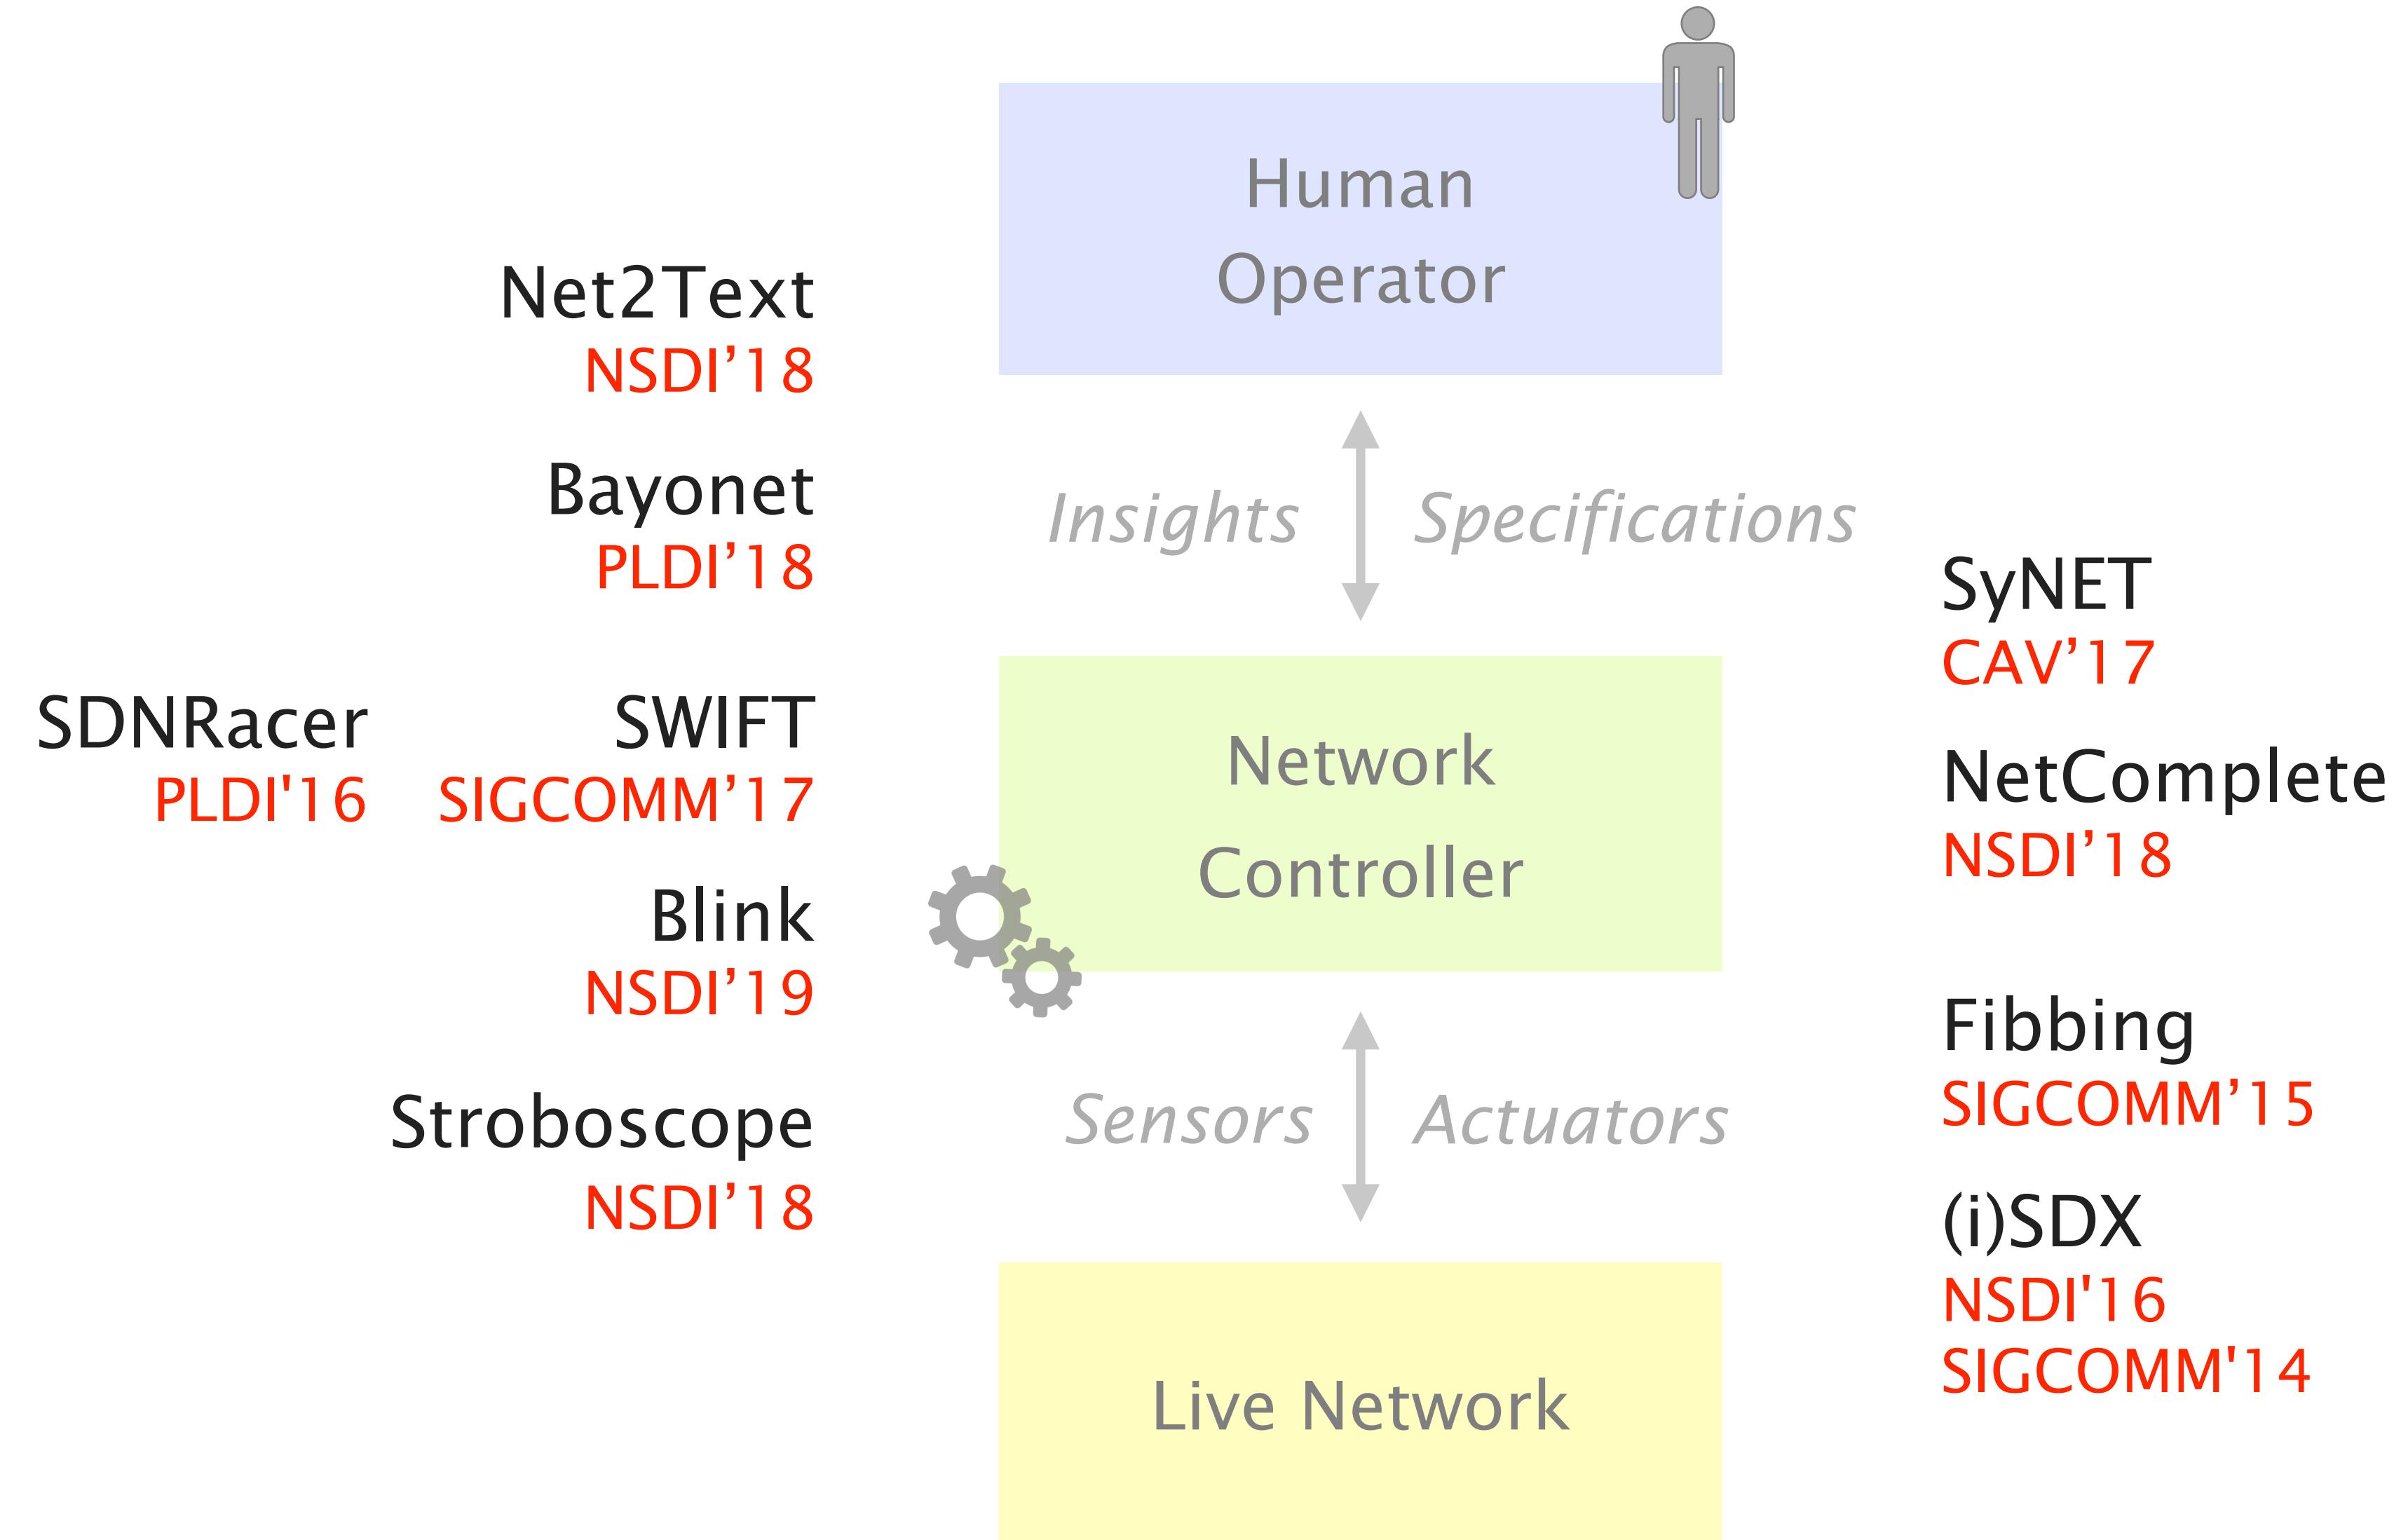
Part I
assisting operators

- 3 conditional automation
- 4 high automation

Part II
taking control

- 5 full automation

too futuristic? 😊



Self-Driving Networks

Breaking new ground in network automation



- 1 operator assistance
- 2 partial automation

Part I
assisting operators

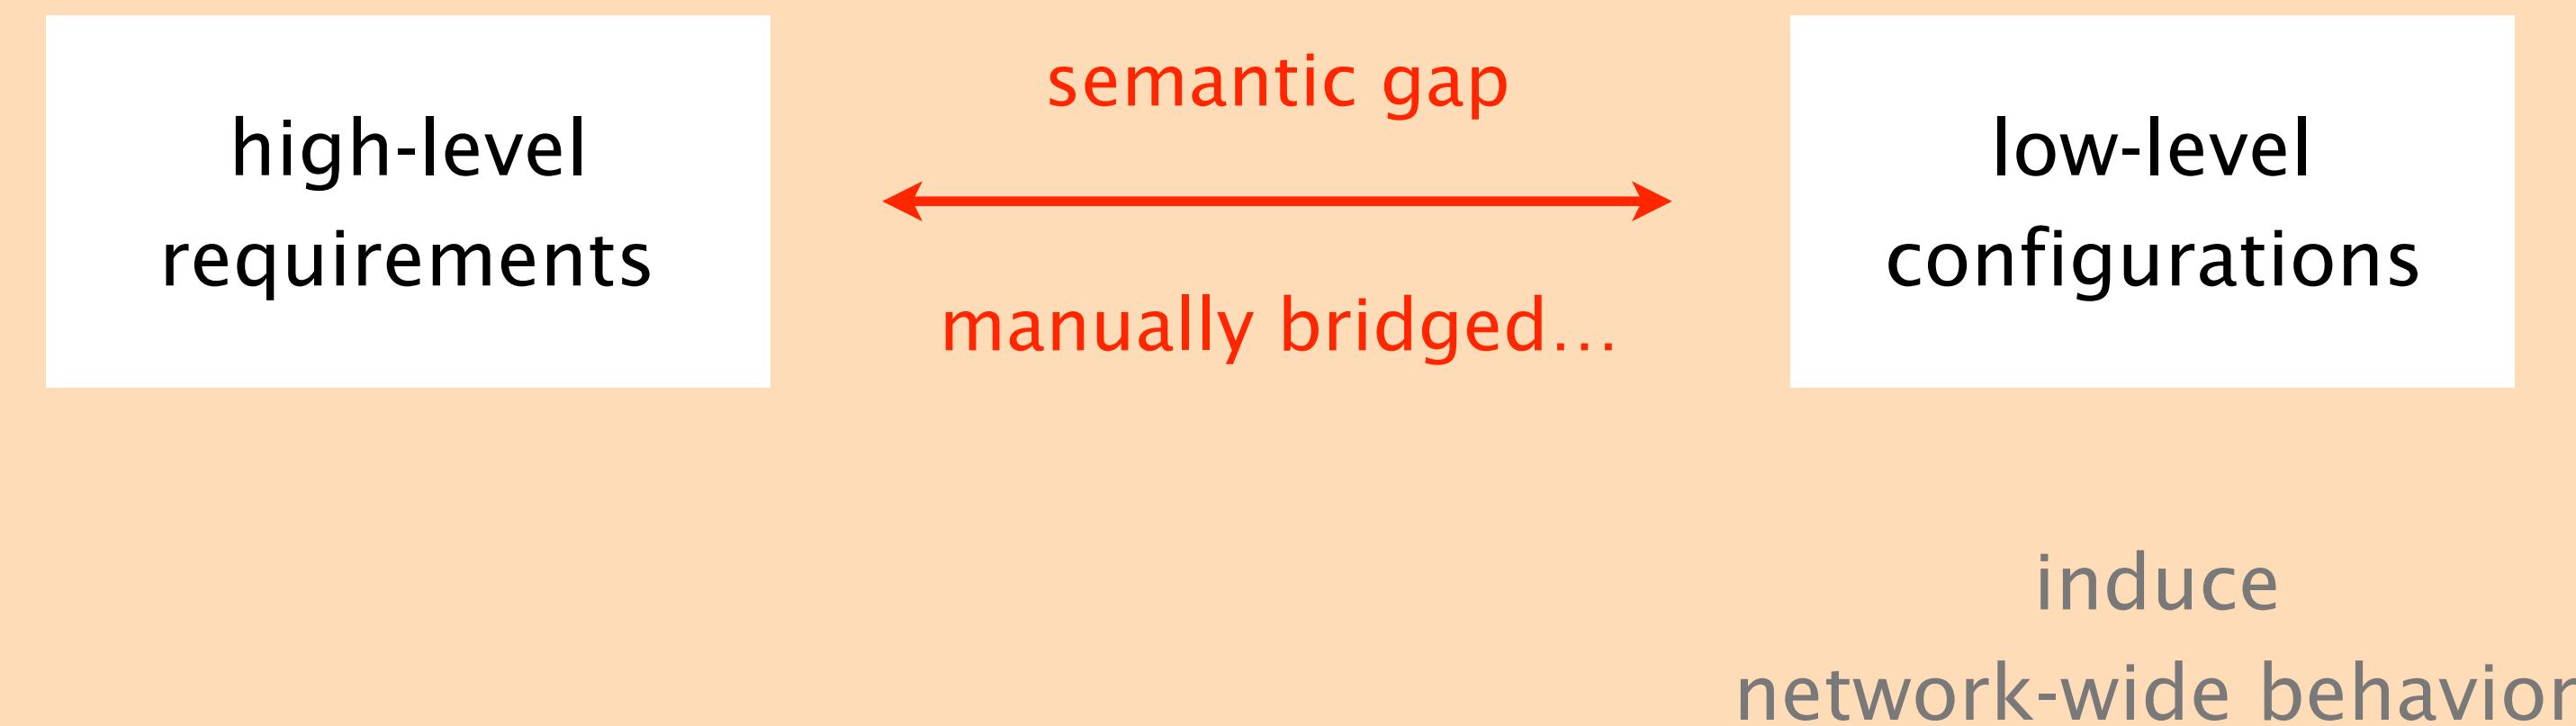
- 3 conditional automation
- 4 high automation

Part II
taking control

- 5 full automation

too futuristic? 😊

Part I: Assisting operators in configuring their networks



NetComplete: Practical Network-Wide Configuration Synthesis with Autocompletion



Ahmed El-Hassany



Petar Tsankov



Martin Vechev



Laurent Vanbever

USENIX Symposium on Networked Systems Design and Implementation. April 2018.

NetComplete assists network operators in bridging the semantic gap between requirements and low-level configurations

NetComplete takes as inputs configuration sketches
together with a set of high-level requirements

A configuration with “holes”

```
interface TenGigabitEthernet1/1/1
```

```
  ip address ? ?
```

```
  ip ospf cost 10 < ? < 100
```

```
router ospf 100
```

```
  ?
```

```
  ...
```

```
router bgp 6500
```

```
  ...
```

```
  neighbor AS200 import route-map imp-p1
```

```
  neighbor AS200 export route-map exp-p1
```

```
  ...
```

```
  ip community-list C1 permit ?
```

```
  ip community-list C2 permit ?
```

```
route-map imp-p1 permit 10
```

```
  ?
```

```
route-map exp-p1 ? 10
```

```
  match community C2
```

```
route-map exp-p1 ? 20
```

```
  match community C1
```

```
  ...
```

NetComplete “autocompletesthe holes such that
the output configuration complies with the requirements

```
interface TenGigabitEthernet1/1/1
```

```
  ip address ? ?
```

```
  ip ospf cost 10 < ? < 100
```

```
router ospf 100
```

```
  ?
```

```
  ...
```

```
router bgp 6500
```

```
  ...
```

```
  neighbor AS200 import route-map imp-p1
```

```
  neighbor AS200 export route-map exp-p1
```

```
  ...
```

```
  ip community-list C1 permit ?
```

```
  ip community-list C2 permit ?
```

```
route-map imp-p1 permit 10
```

```
  ?
```

```
route-map exp-p1 ? 10
```

```
  match community C2
```

```
route-map exp-p1 ? 20
```

```
  match community C1
```

```
  ...
```

```
interface TenGigabitEthernet1/1/1
    ip address 10.0.0.1 255.255.255.254
    ip ospf cost 15
```

```
router ospf 100
    network 10.0.0.1 0.0.0.1 area 0.0.0.0
```

```
router bgp 6500
```

```
...
```

```
neighbor AS200 import route-map imp-p1
neighbor AS200 export route-map exp-p1
```

```
...
```

```
ip community-list C1 permit 6500:1
ip community-list C2 permit 6500:2
```

```
route-map imp-p1 permit 10
    set community 6500:1
    set local-pref 50
route-map exp-p1 permit 10
    match community C2
route-map exp-p1 deny 20
    match community C1
...
```

NetComplete reduces the autocompletion problem
to a **constraint satisfaction problem**

First

- Encode the
- protocol semantics
 - high-level requirements as a logical formula (in SMT)
 - partial configurations

- First Encode the
 - protocol semantics
 - high-level requirements as a logical formula (in SMT)
 - partial configurations
- Then Use a solver (Z3) to find an assignment for the undefined configuration variables s.t. the formula evaluates to True

Main challenge:

Scalability

Insight #1

network-specific
heuristics

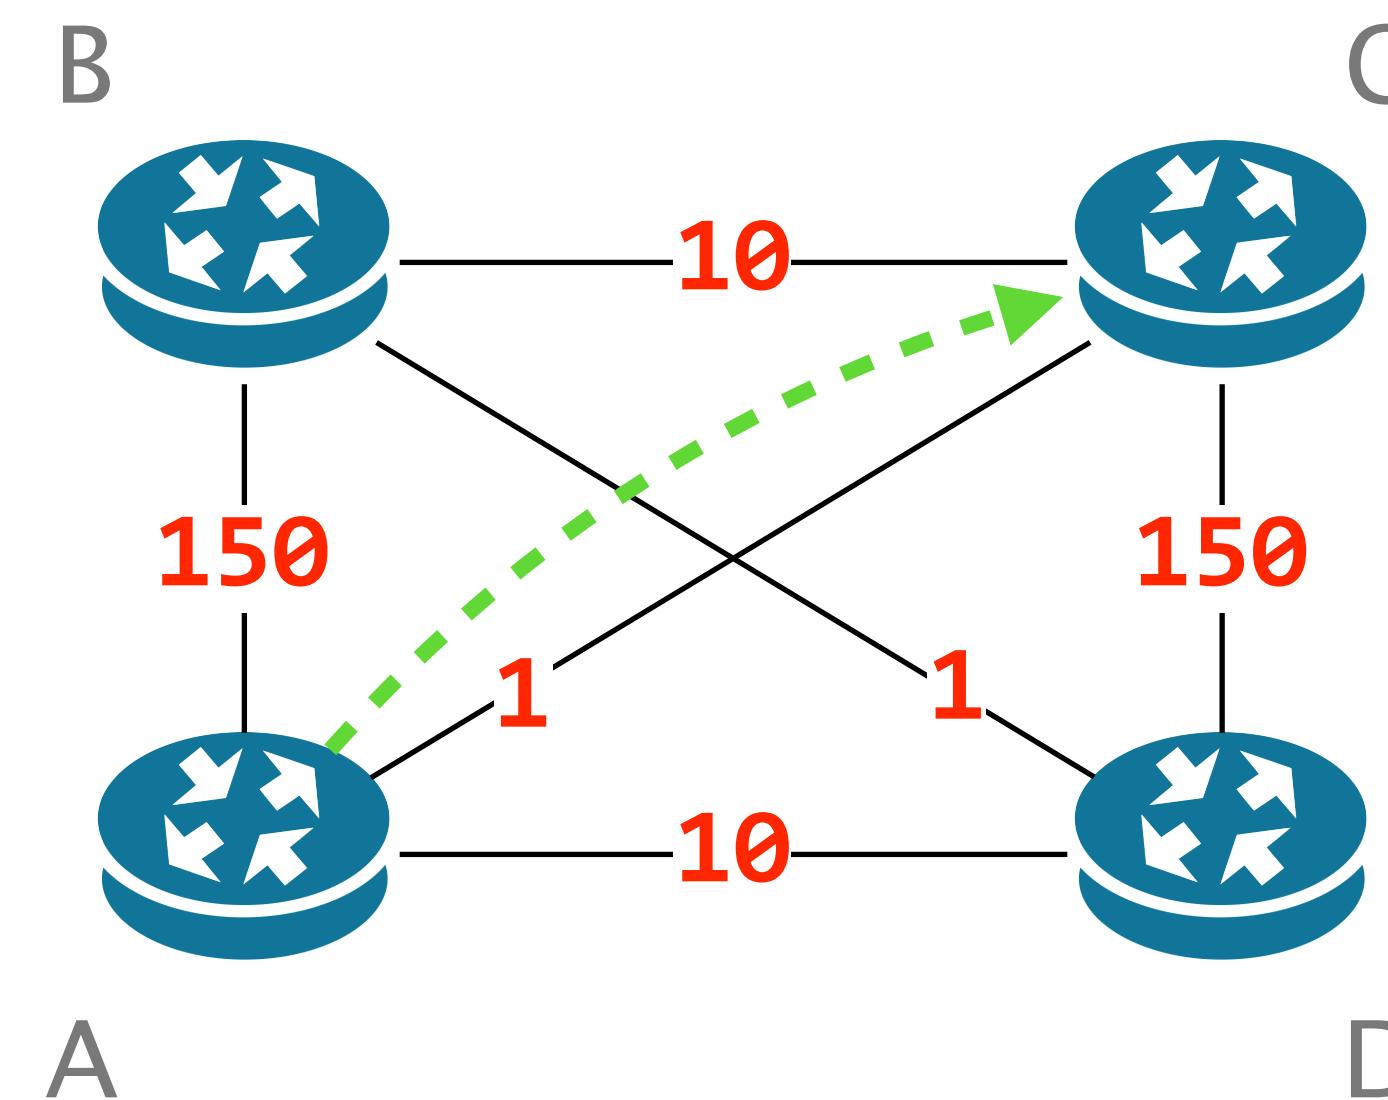
search space navigation

Insight #2

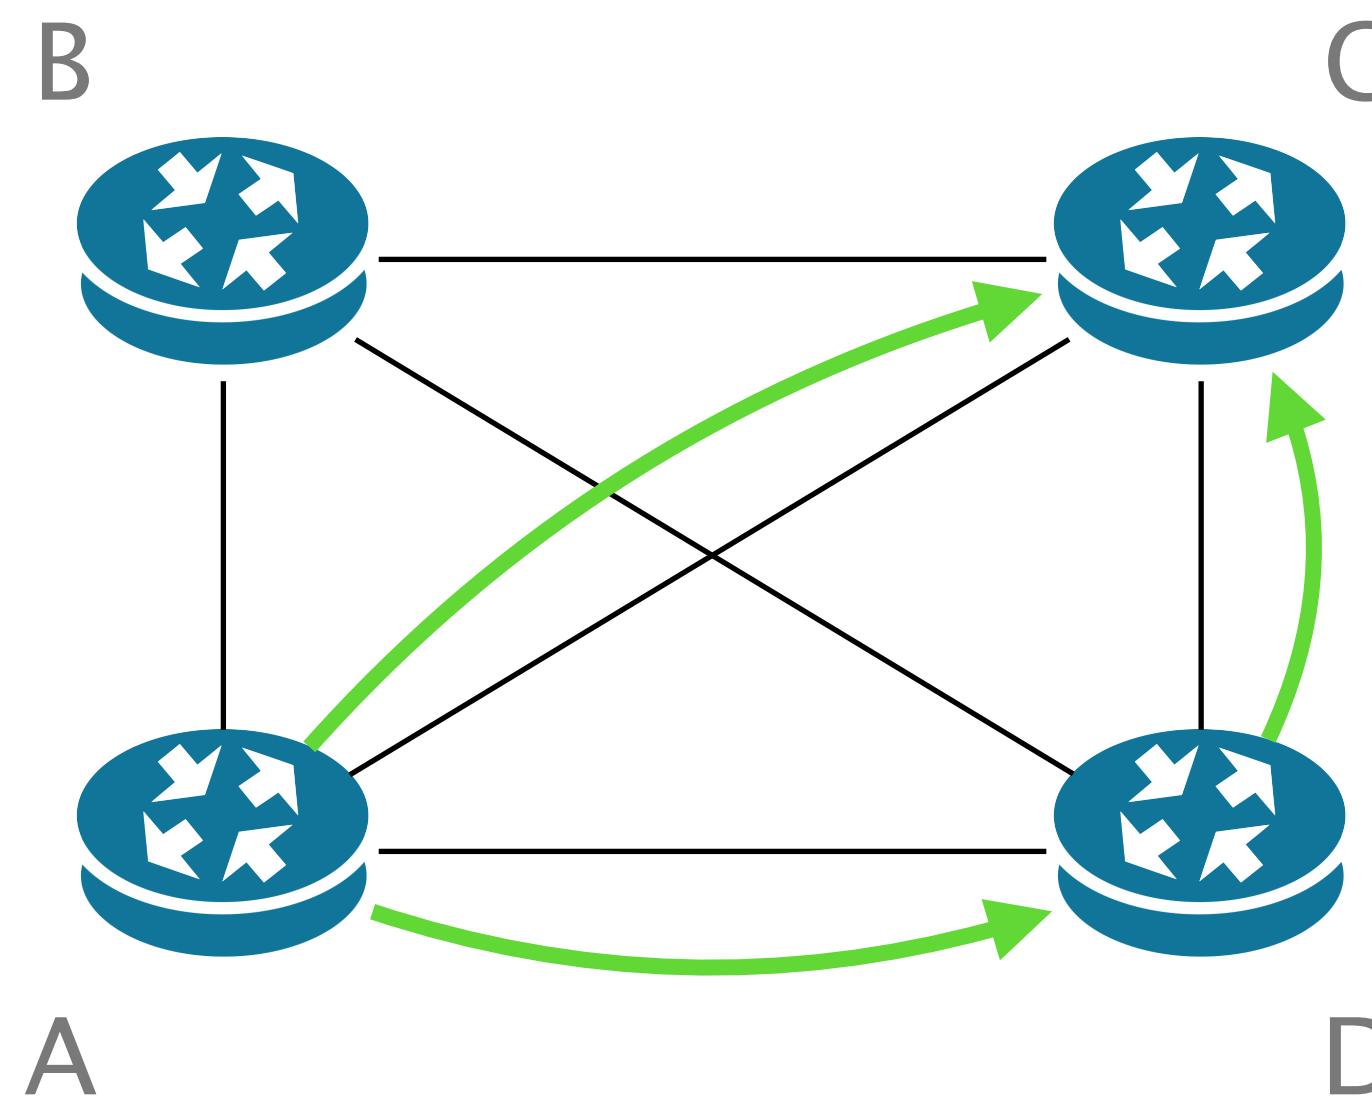
partial evaluation

search space reduction

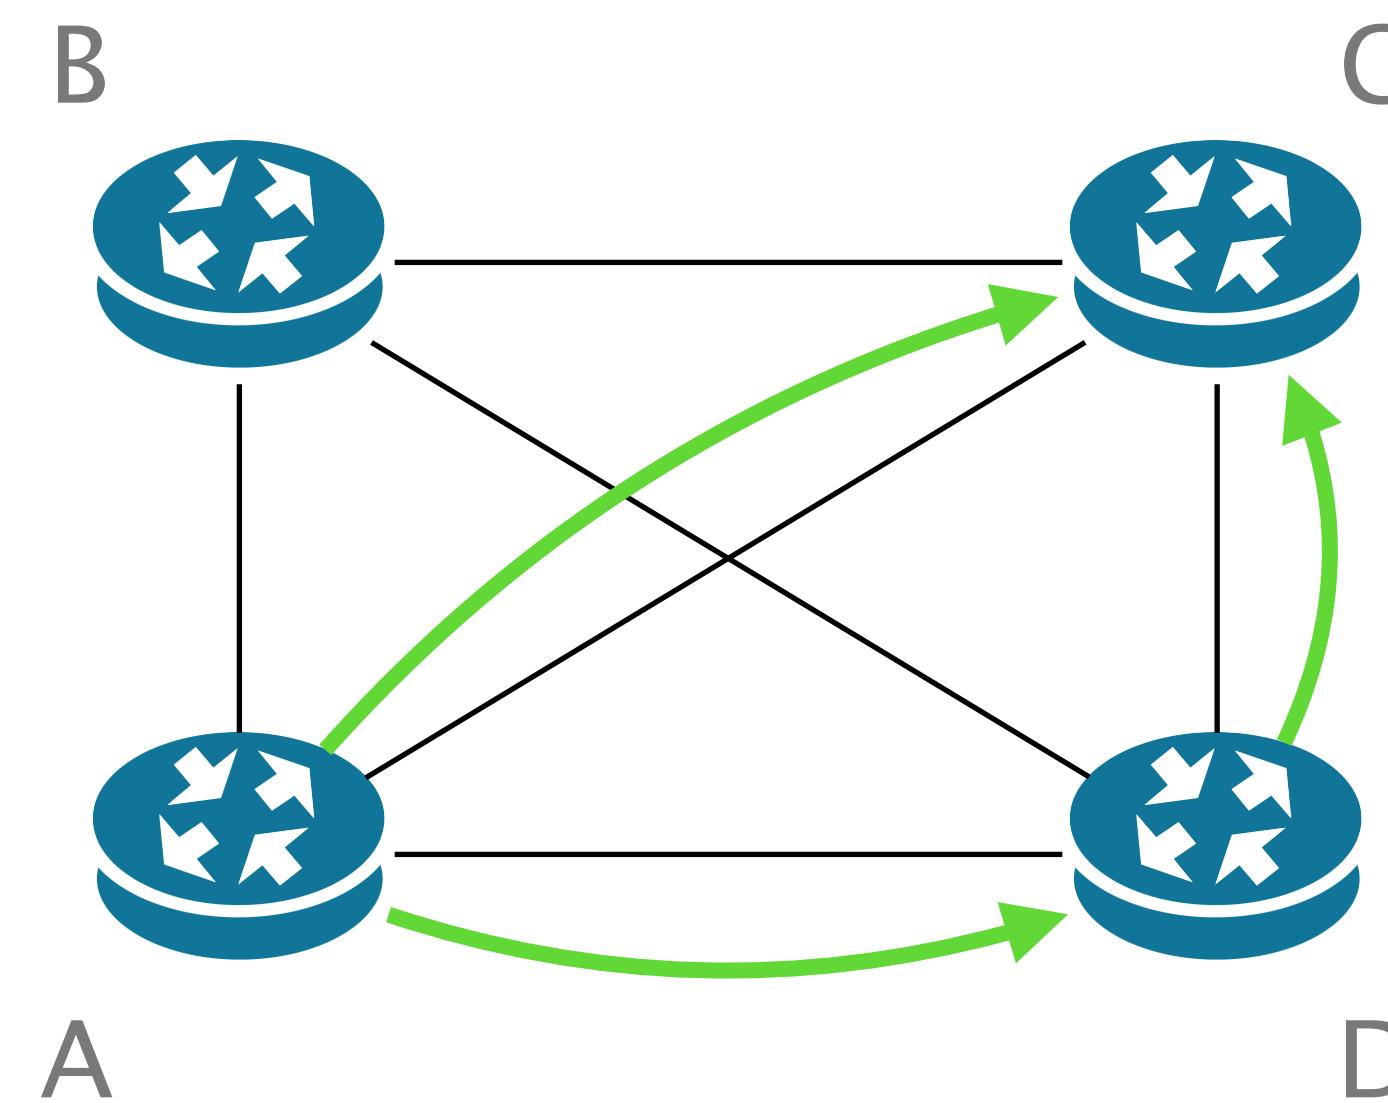
Consider this initial configuration in which
(A,C) traffic is forwarded along the direct link



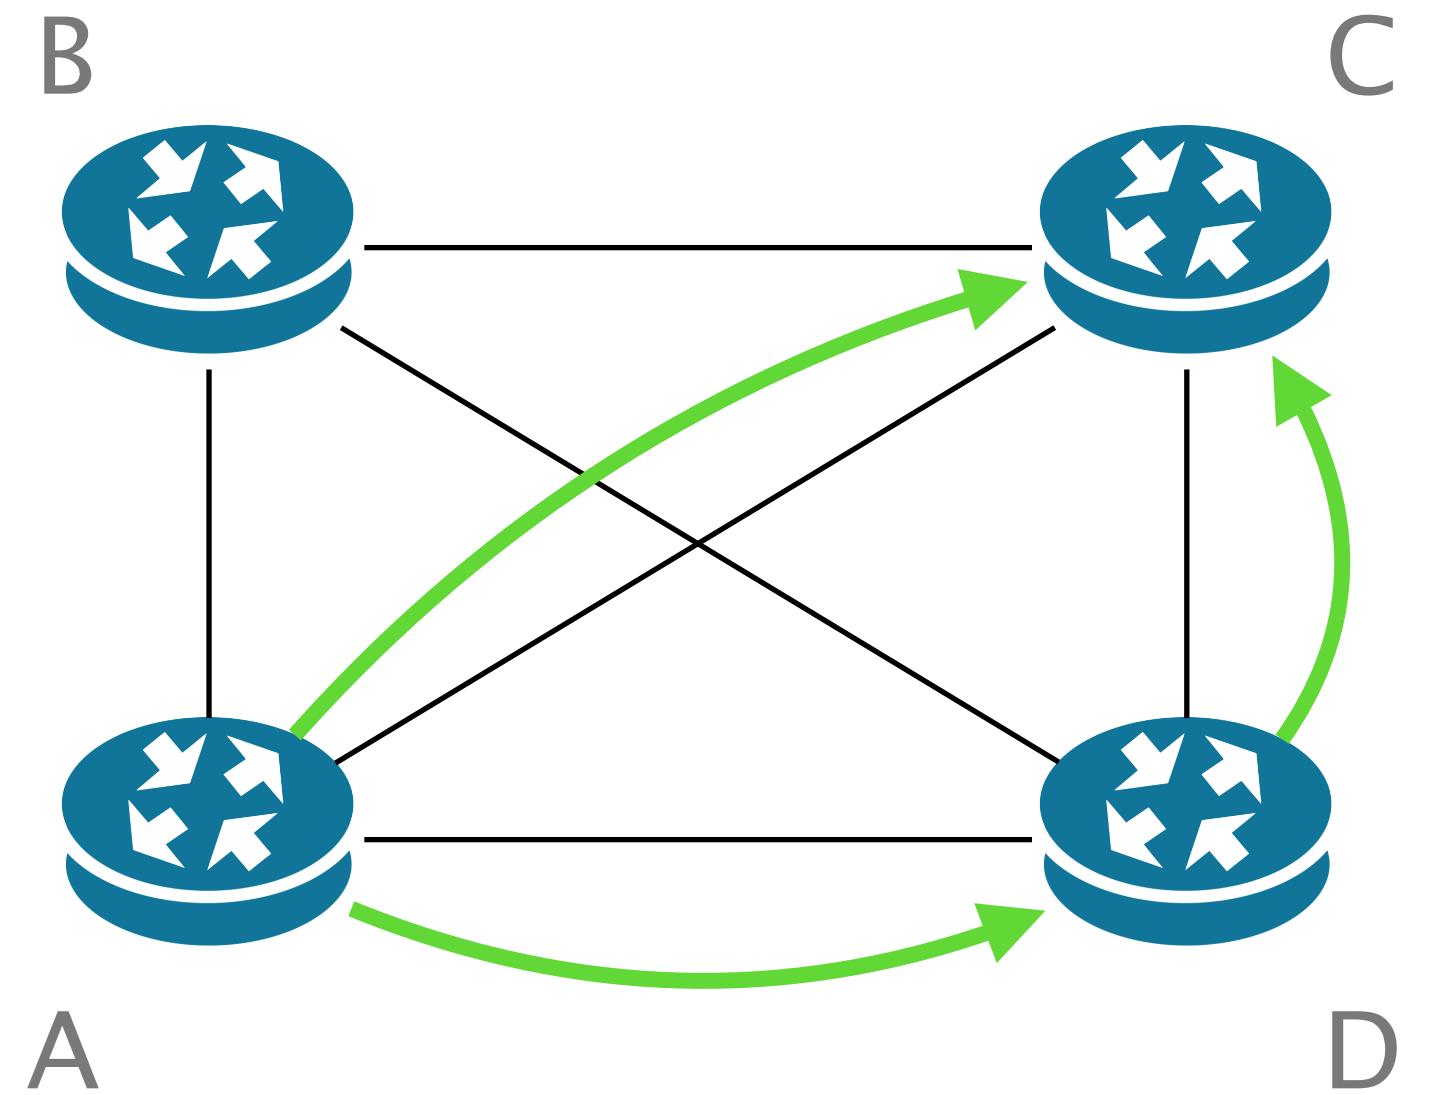
For performance reasons,
the operators want to enable load-balancing



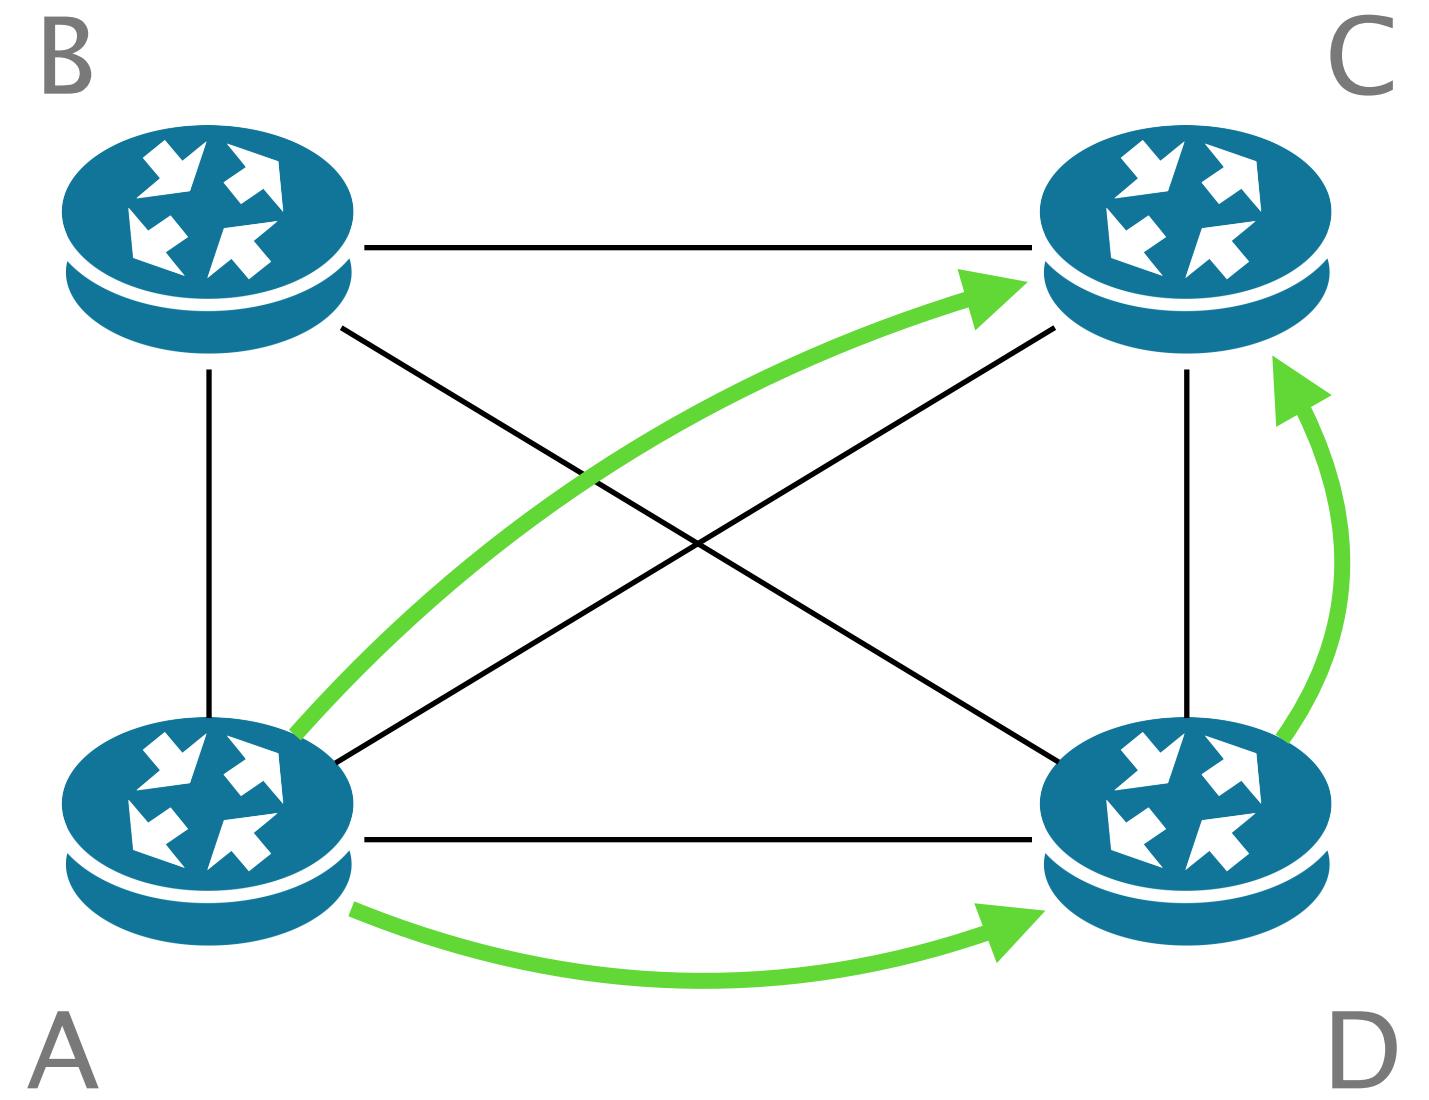
What should be the weights for this to happen?



input requirements

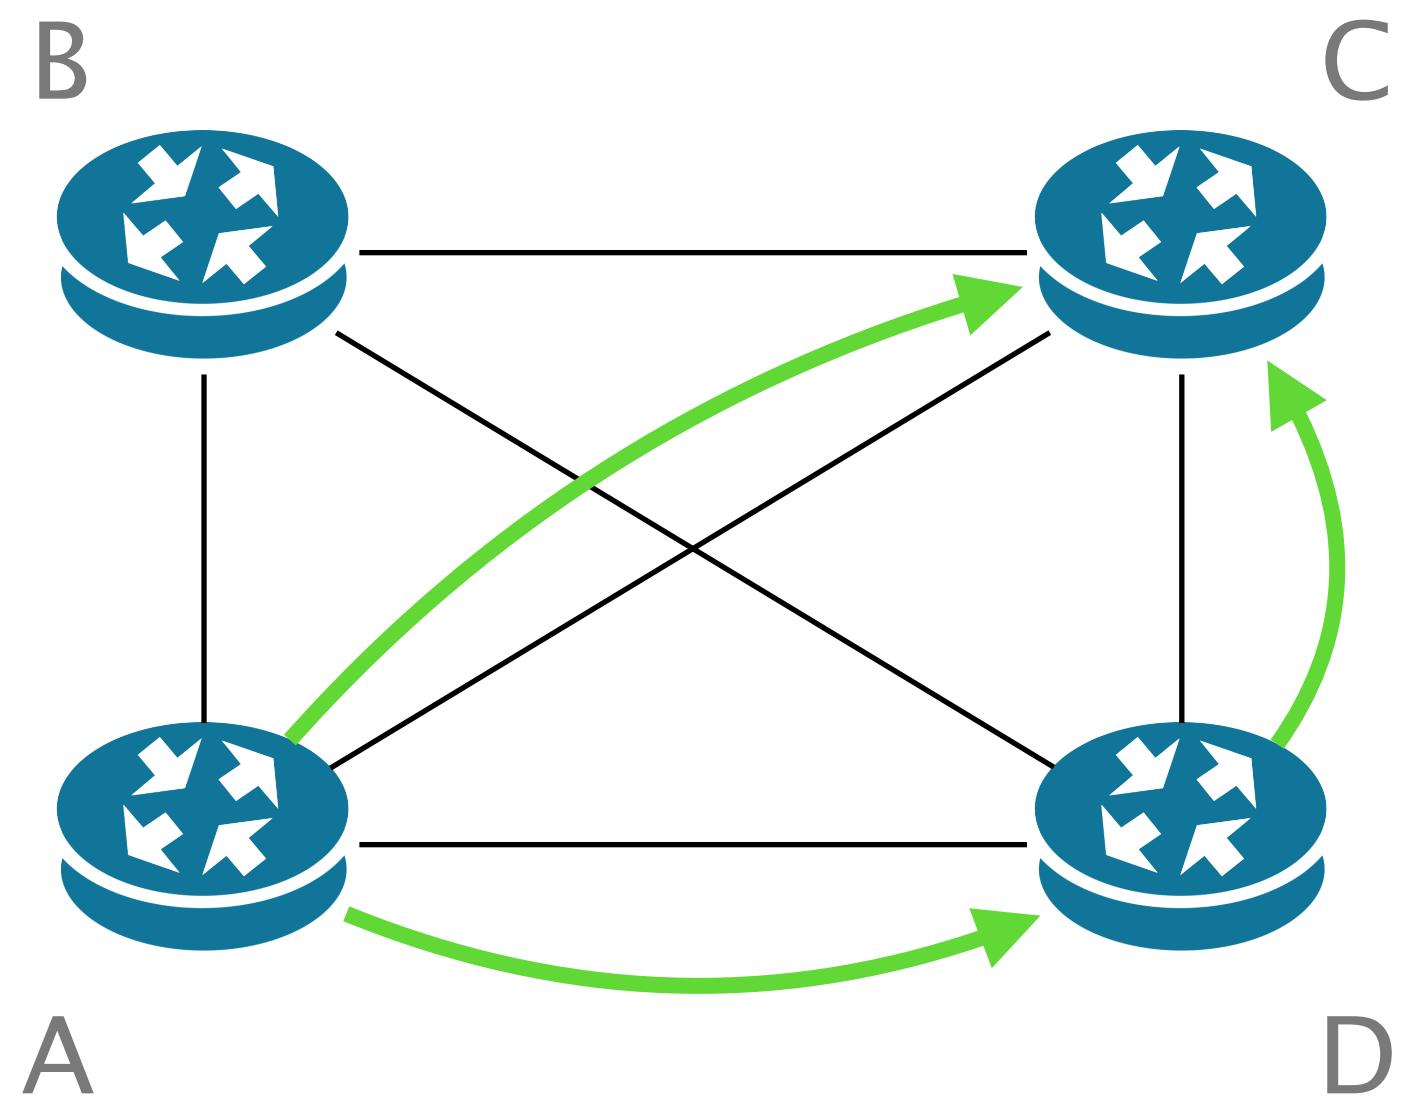


input requirements



synthesis procedure

input requirements

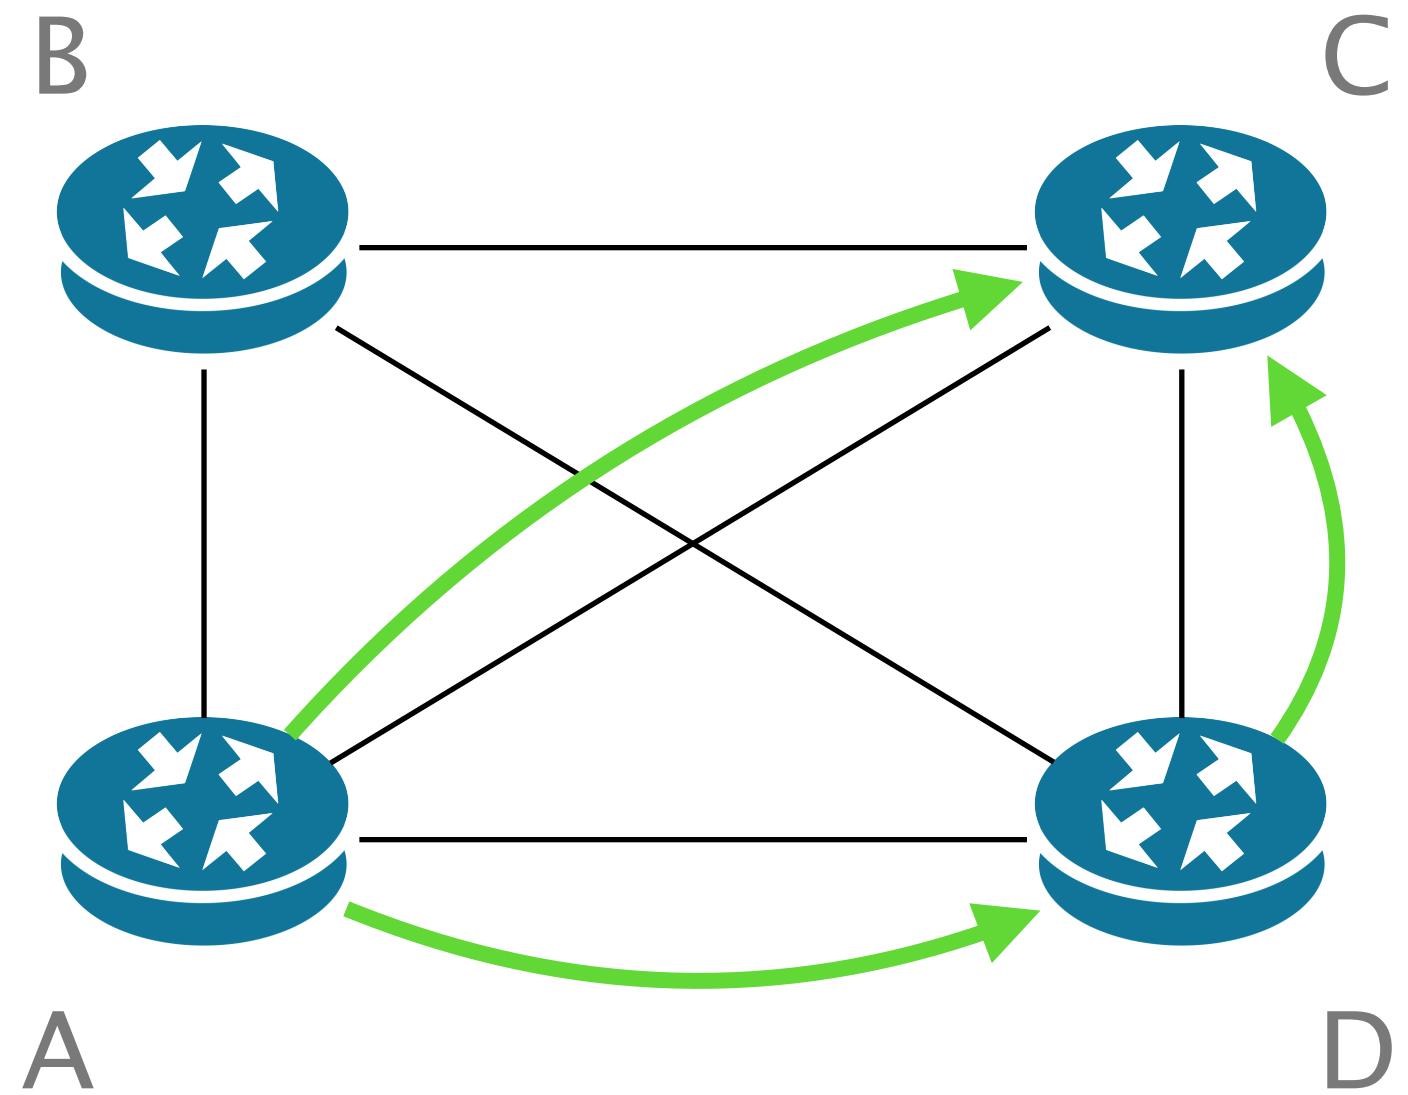


synthesis procedure

$$\forall X \in \text{Paths}(A,C) \setminus \text{Reqs}$$

$$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$$

input requirements



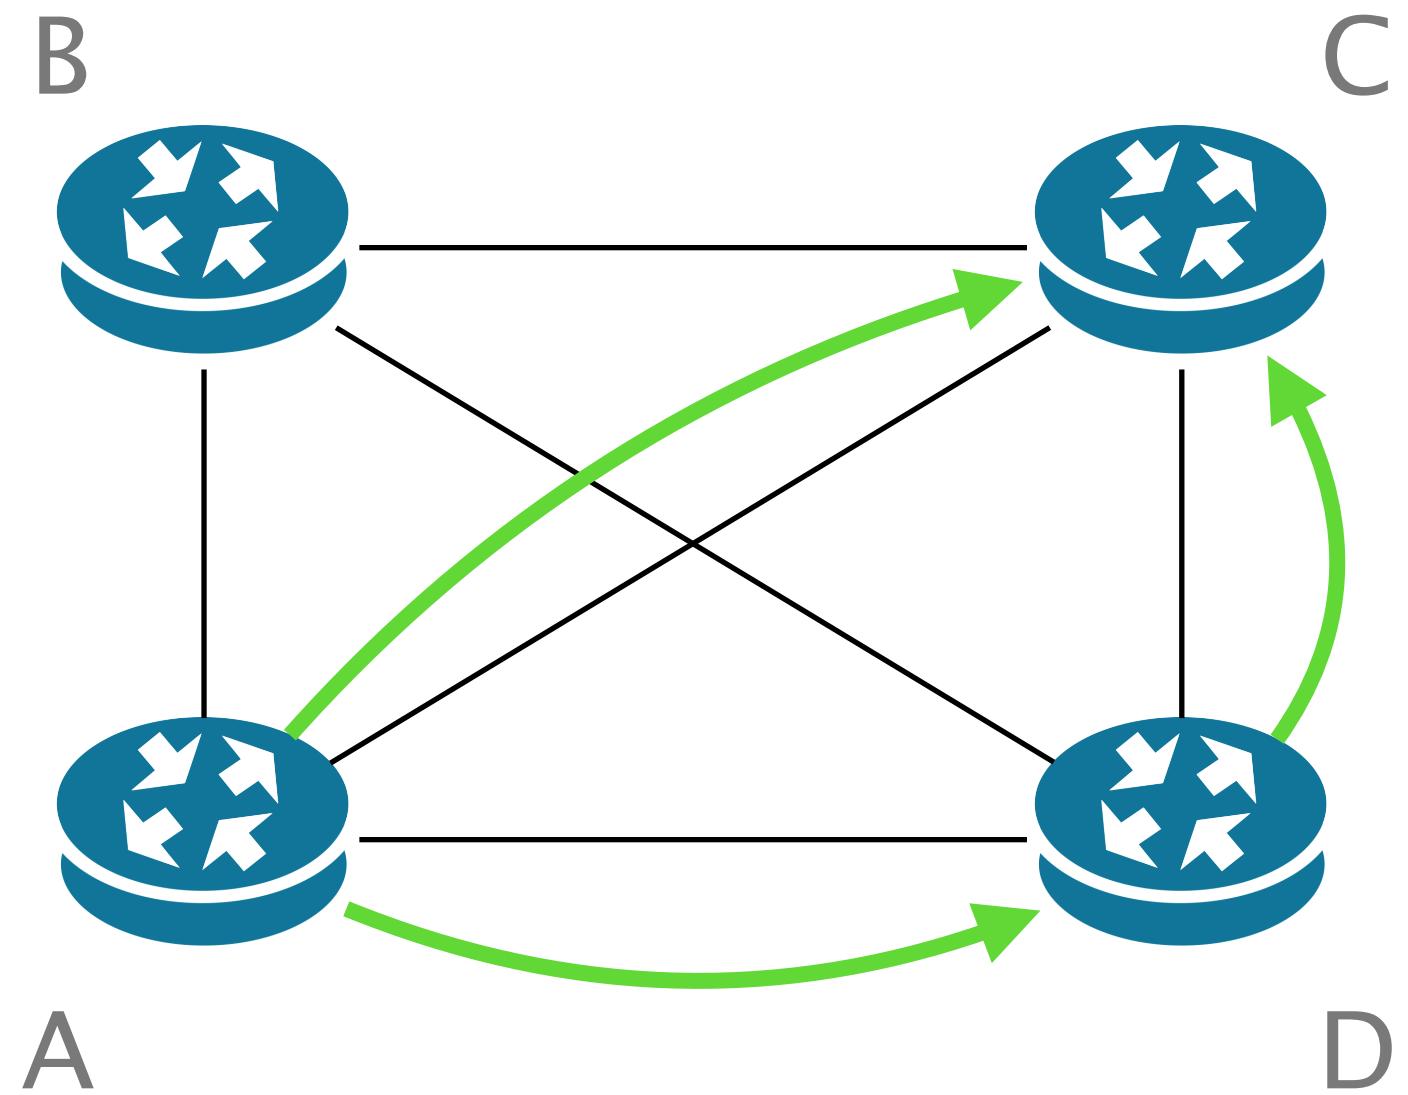
synthesis procedure

$\forall X \in \text{Paths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

input requirements



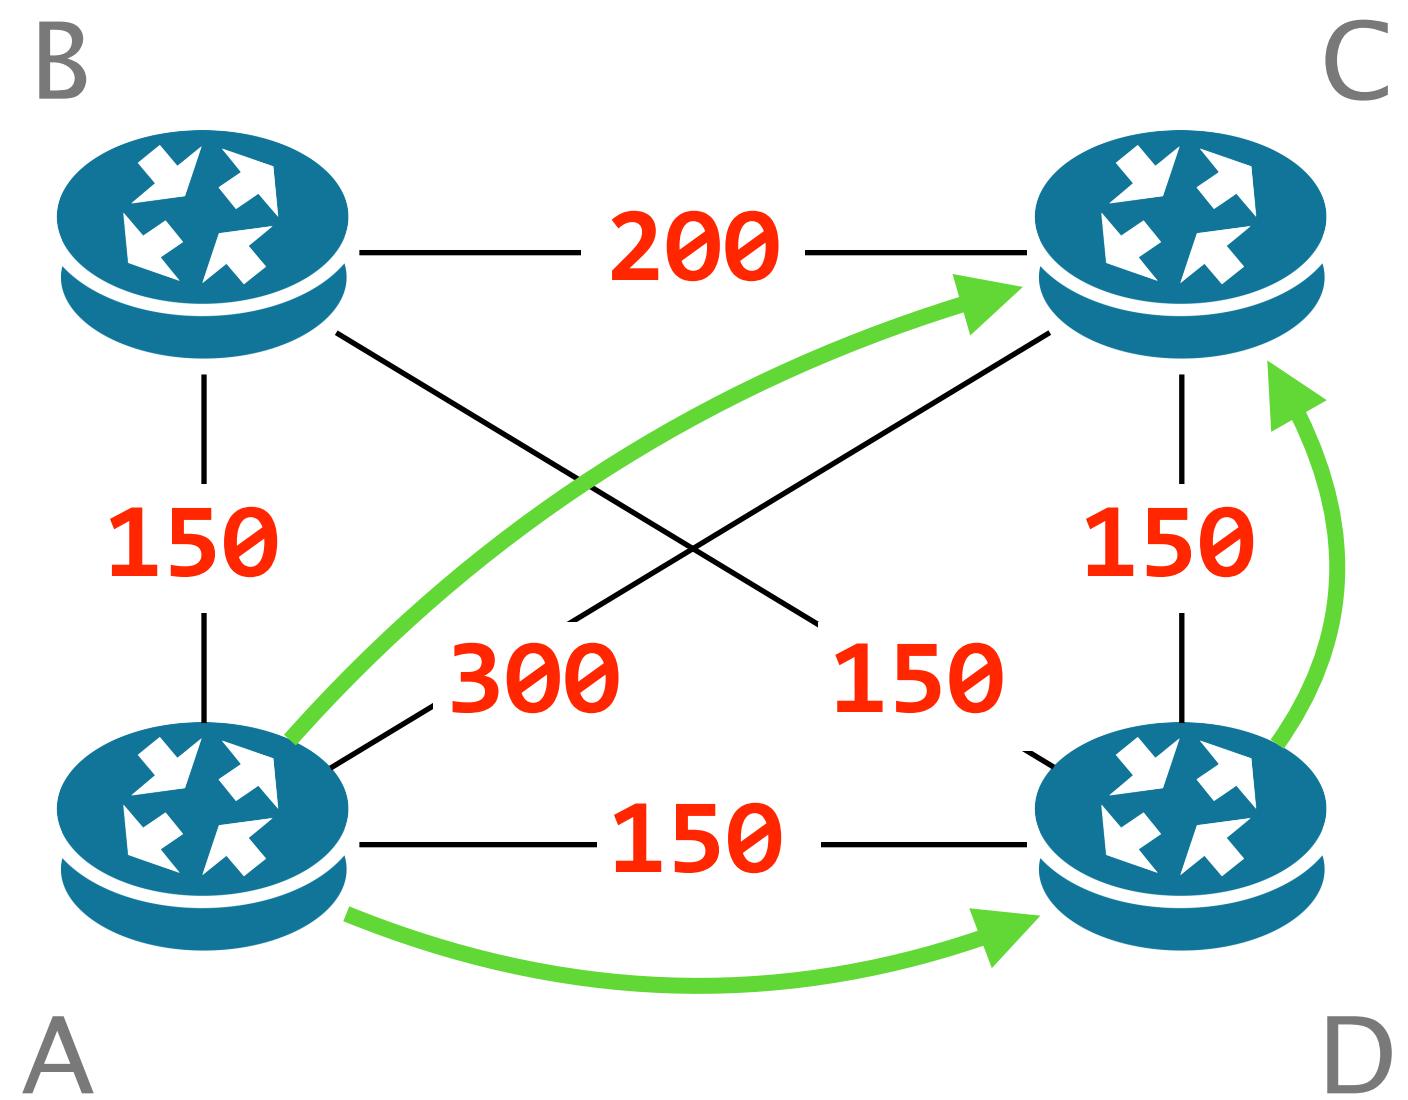
synthesis procedure

$\forall X \in \text{Paths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

input requirements



Synthesized weights

synthesis procedure

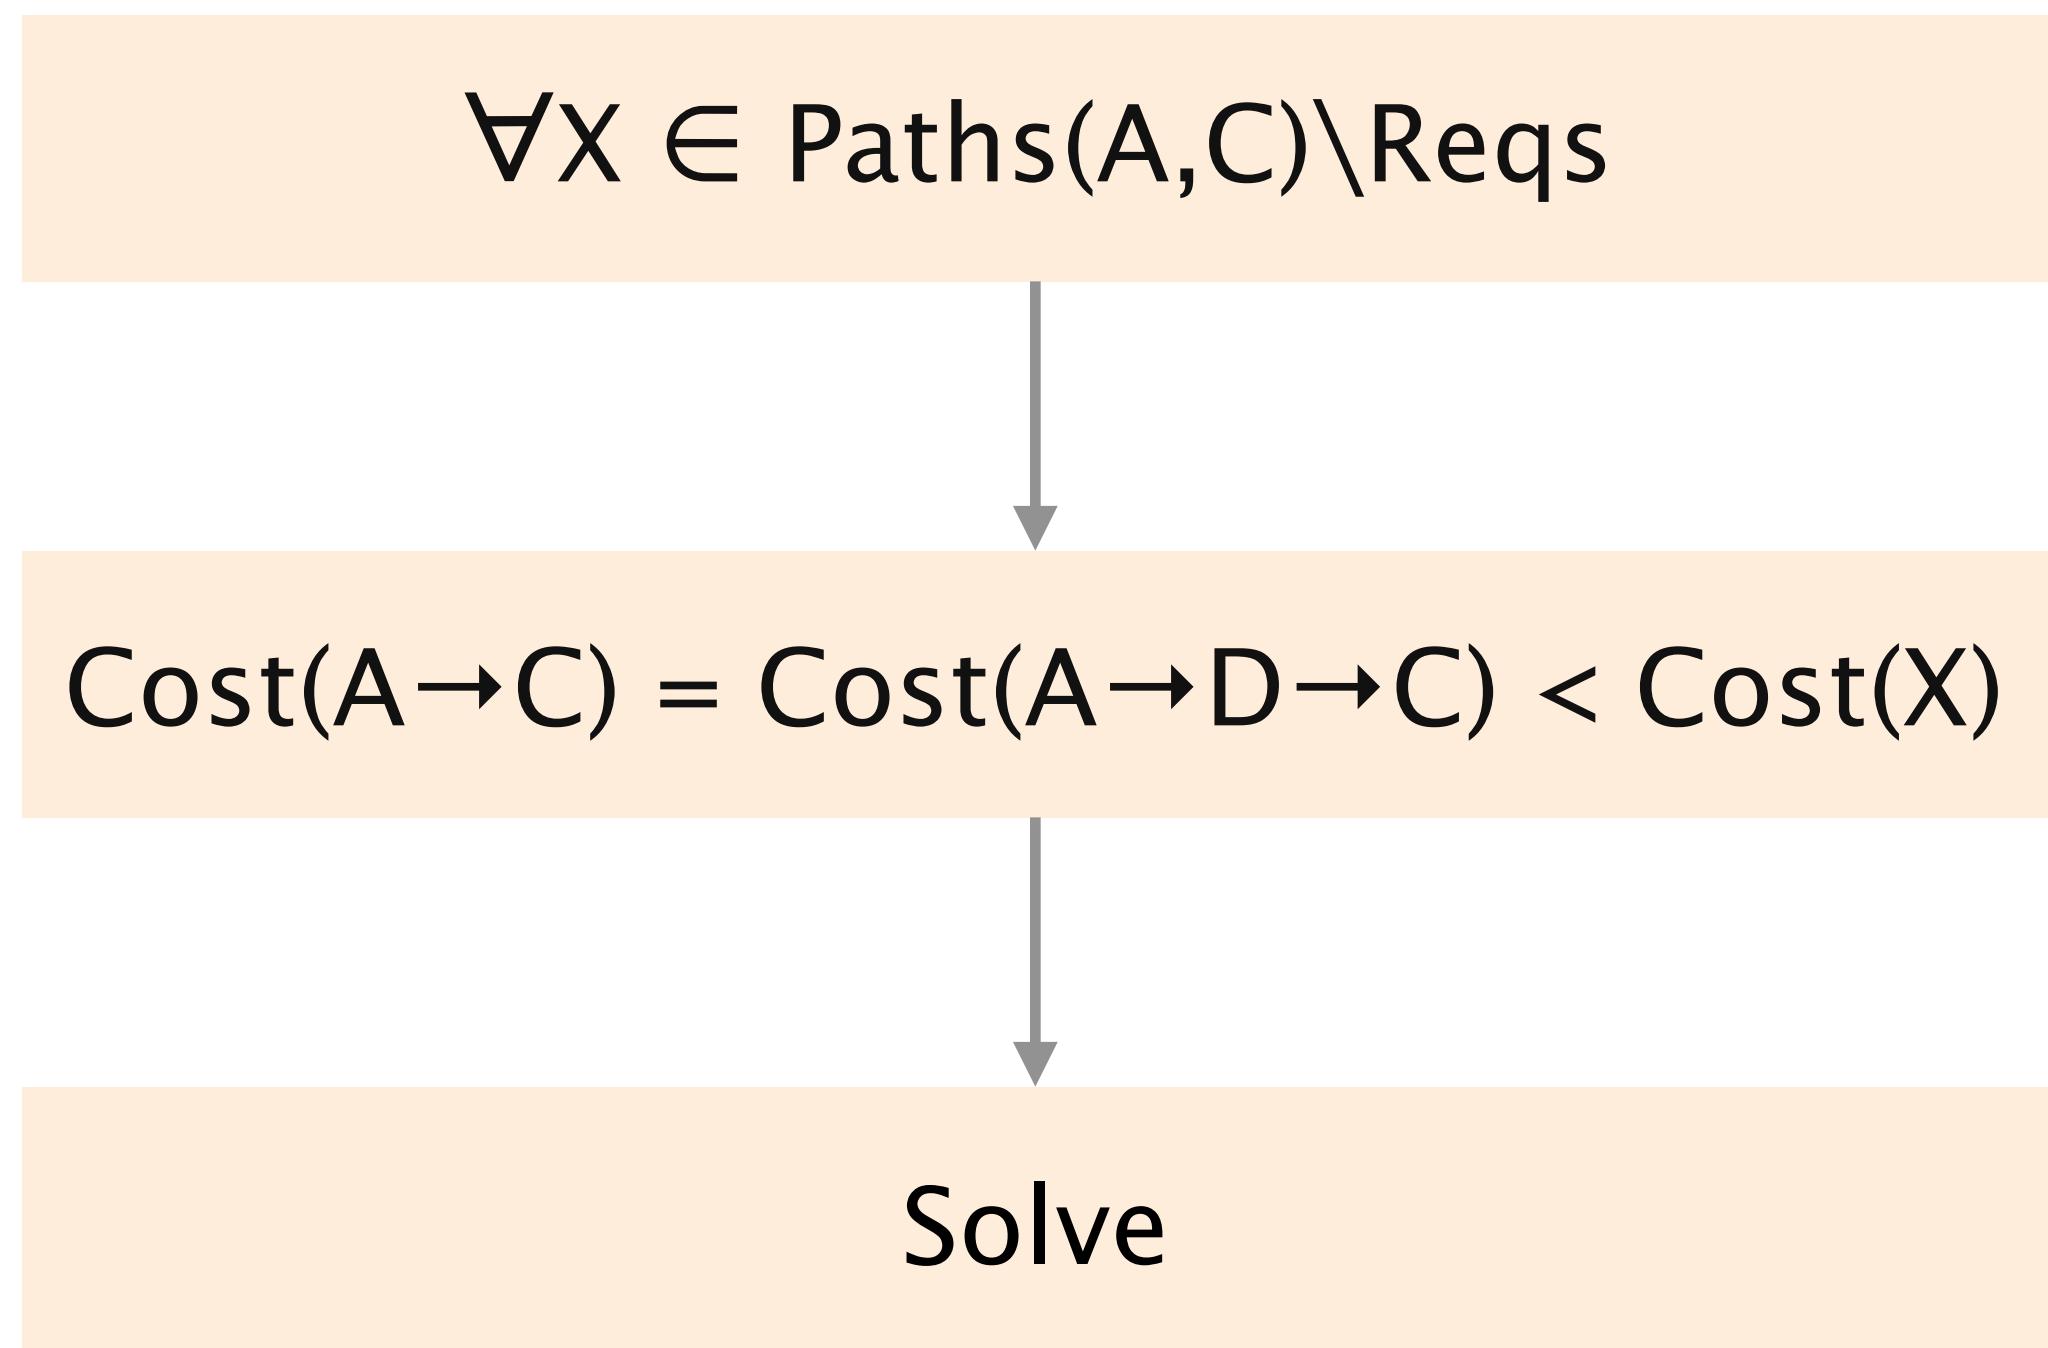
$\forall X \in \text{Paths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

This was easy, but...

it does **not** scale



There can be an exponential number of paths
between A and C...

$\forall X \in \text{Paths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

To scale, NetComplete leverages
Counter-Example Guided Inductive Synthesis (CEGIS)

An contemporary approach to synthesis where
a solution is iteratively learned from counter-examples

While enumerating all paths is hard,
computing shortest paths given weights is easy!

Instead of considering all paths between X and Y

CEGIS
Part 1

Instead of considering all paths between X and Y

Consider a random subset S of them and
synthesize the weights considering S only

Instead of considering all paths between X and Y

CEGIS
Part 1

Consider a random subset S of them and
synthesize the weights considering S only

intuition

Fast as S is small compared to all paths

Instead of considering all paths between X and Y

CEGIS
Part 1

Consider a random subset S of them and
synthesize the weights considering S only

intuition

Fast as S is small compared to all paths
but weights can be wrong

Instead of considering all paths between X and Y

CEGIS
Part 1

Consider a random subset S of them and
synthesize the weights considering S only

CEGIS
Part 2

Check whether the weights found comply
with the requirements over all paths

If so return
Else take a counter-example (a path)
that violates the Req and add it to S

Repeat.

Instead of considering all paths between X and Y

CEGIS
Part 1

Consider a random subset S of them and
synthesize the weights considering S only

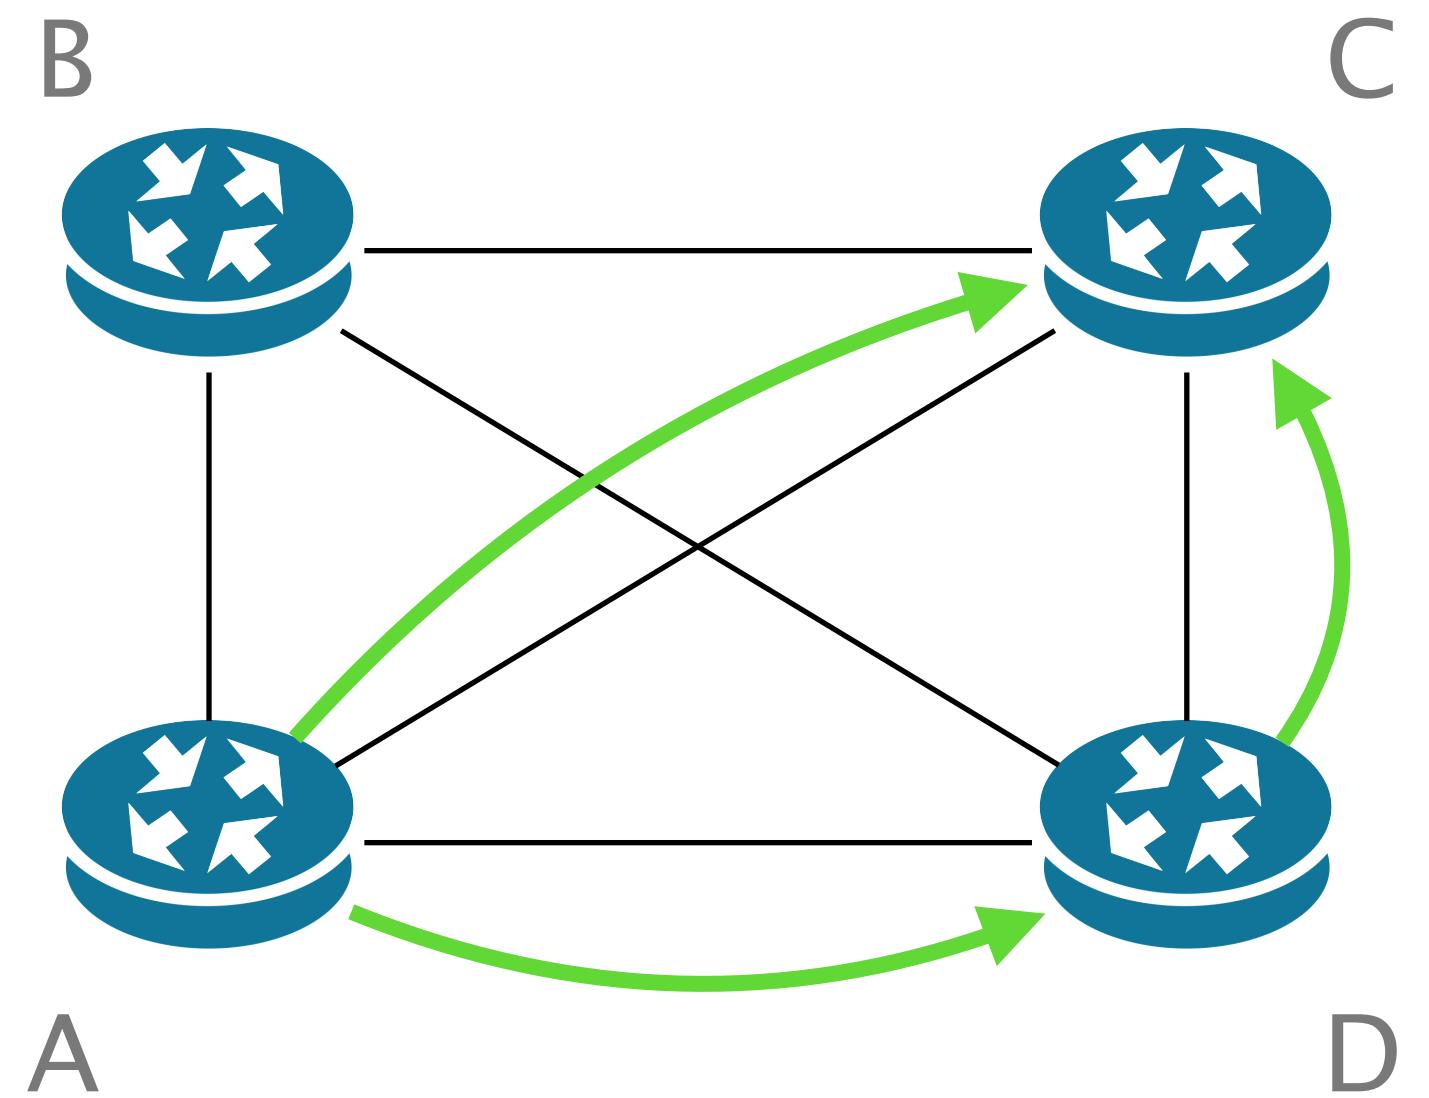
CEGIS
Part 2

Check whether the weights found comply
with the requirements **over all paths**

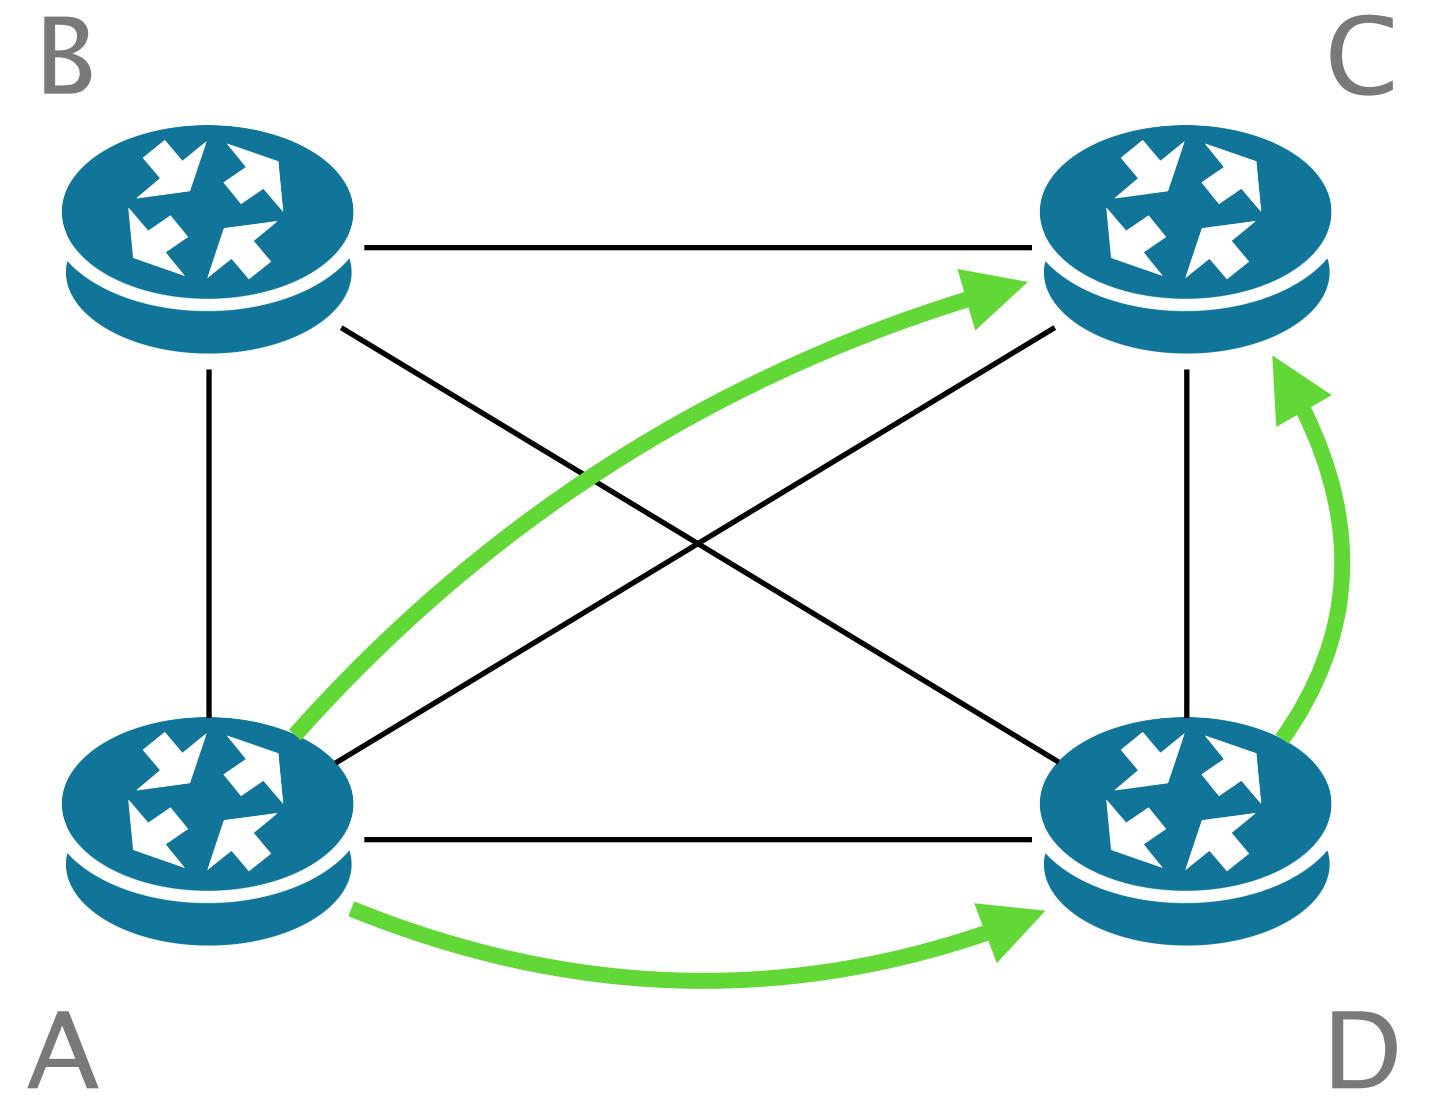
intuition

Fast too
simple shortest-path computation

input requirements

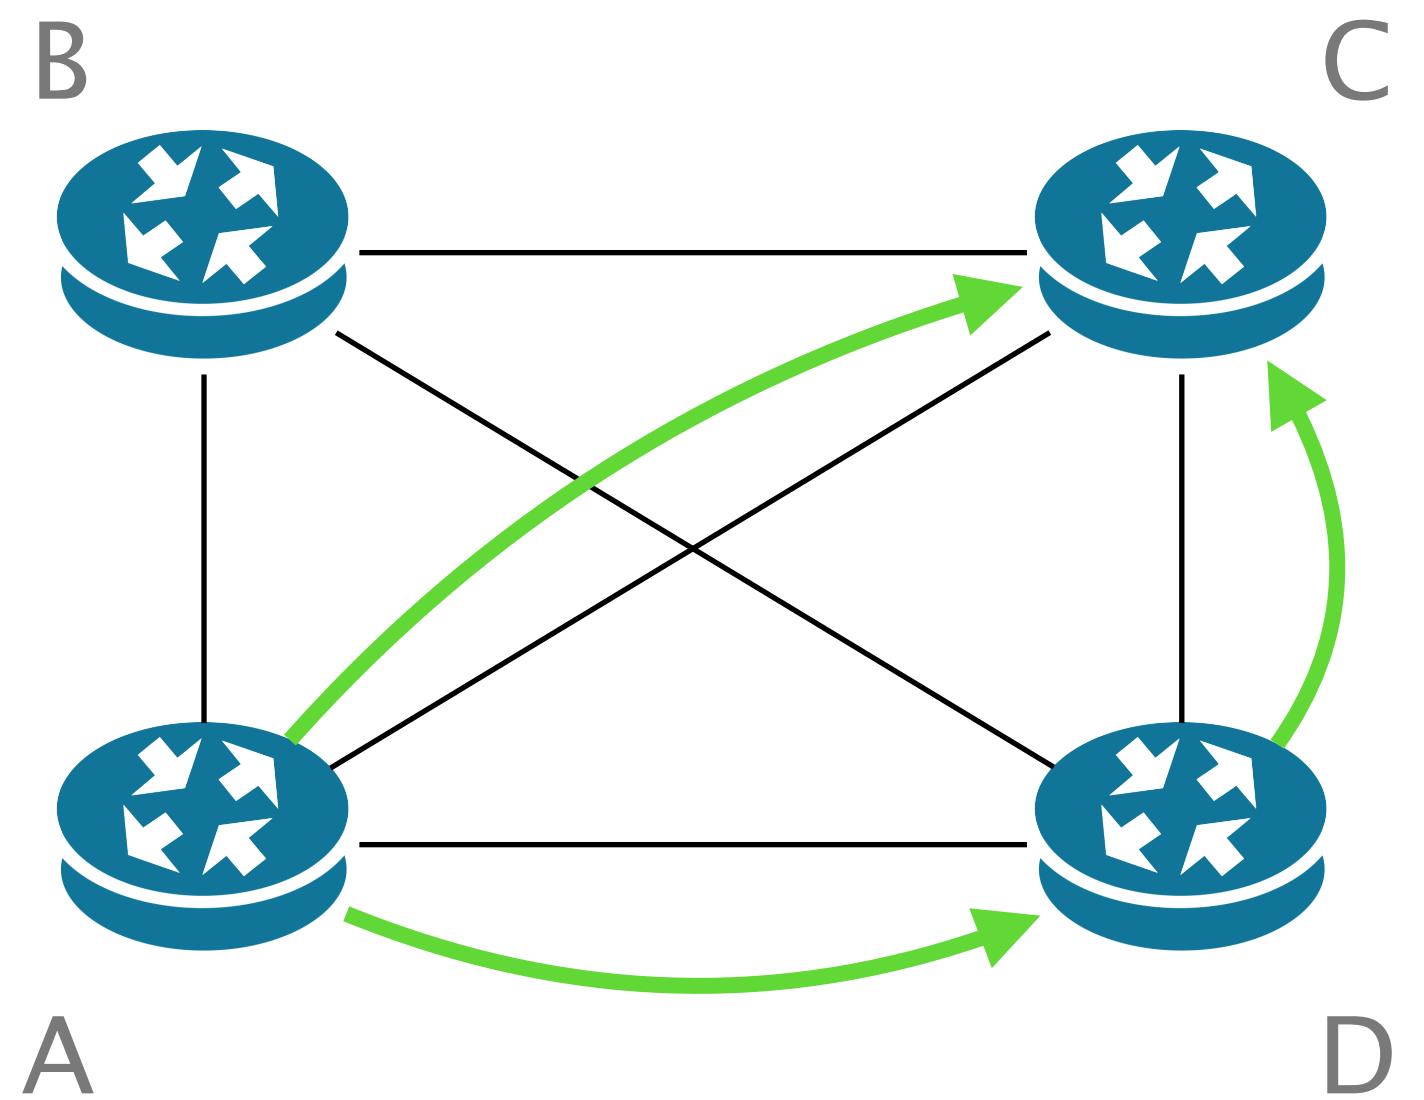


input requirements



synthesis procedure

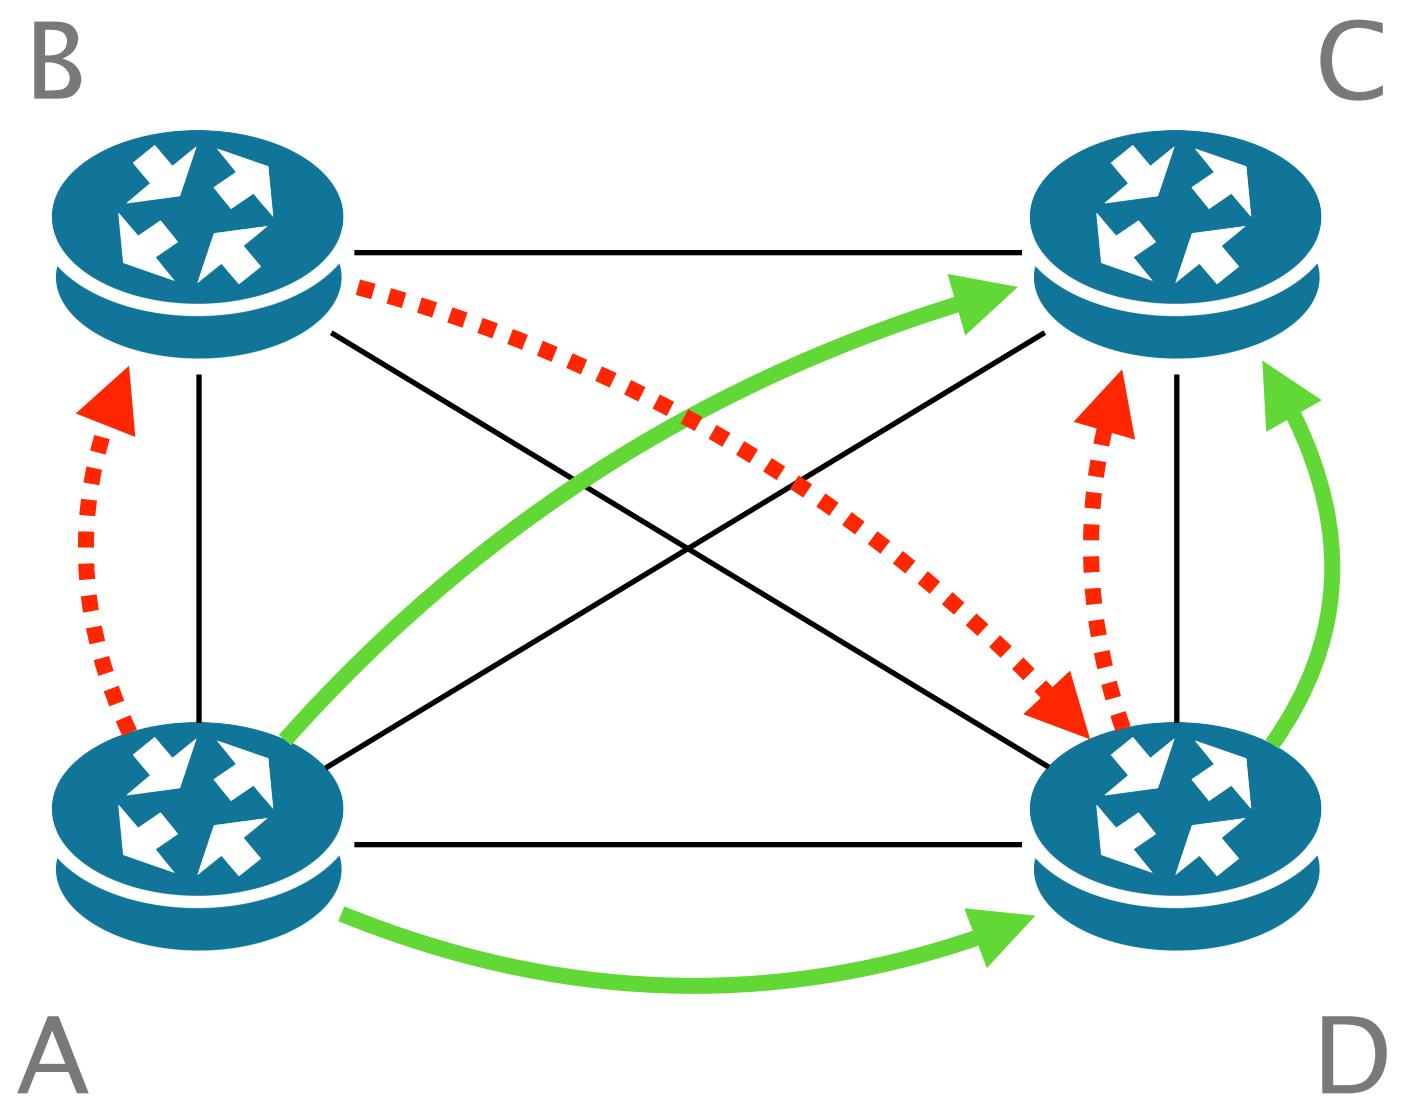
input requirements



synthesis procedure

$\forall x \in \text{SamplePaths}(A, C) \setminus \text{Reqs}$

input requirements

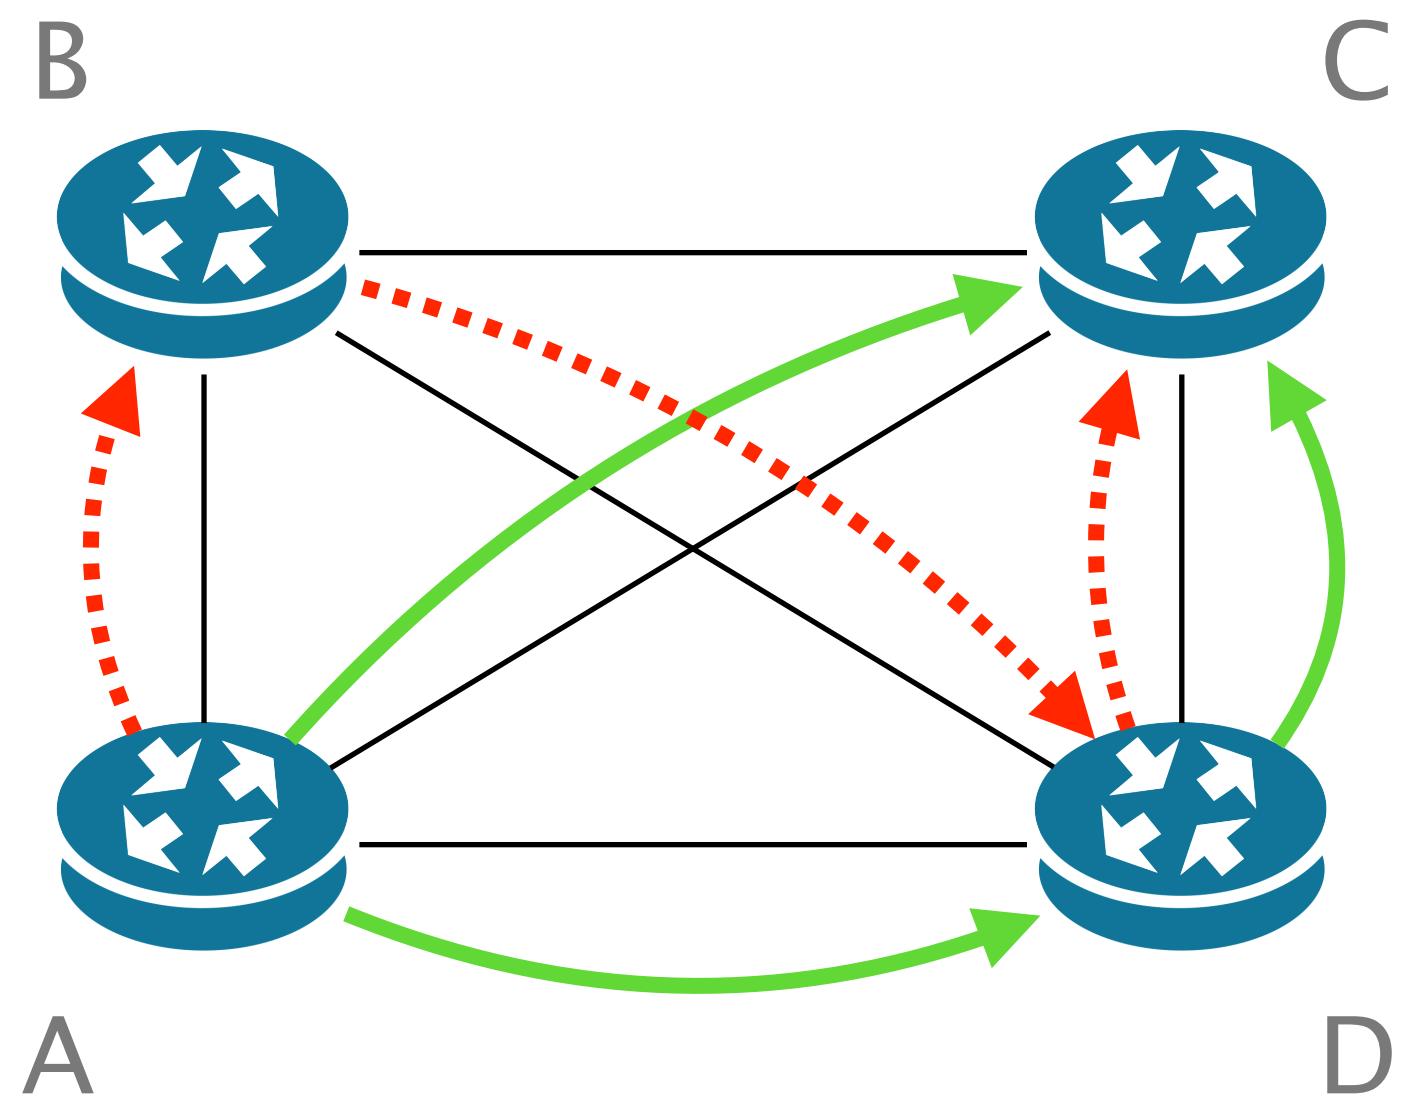


synthesis procedure

$$\forall x \in \text{SamplePaths}(A, C) \setminus \text{Reqs}$$

Sample: { [A,B,D,C] }

input requirements

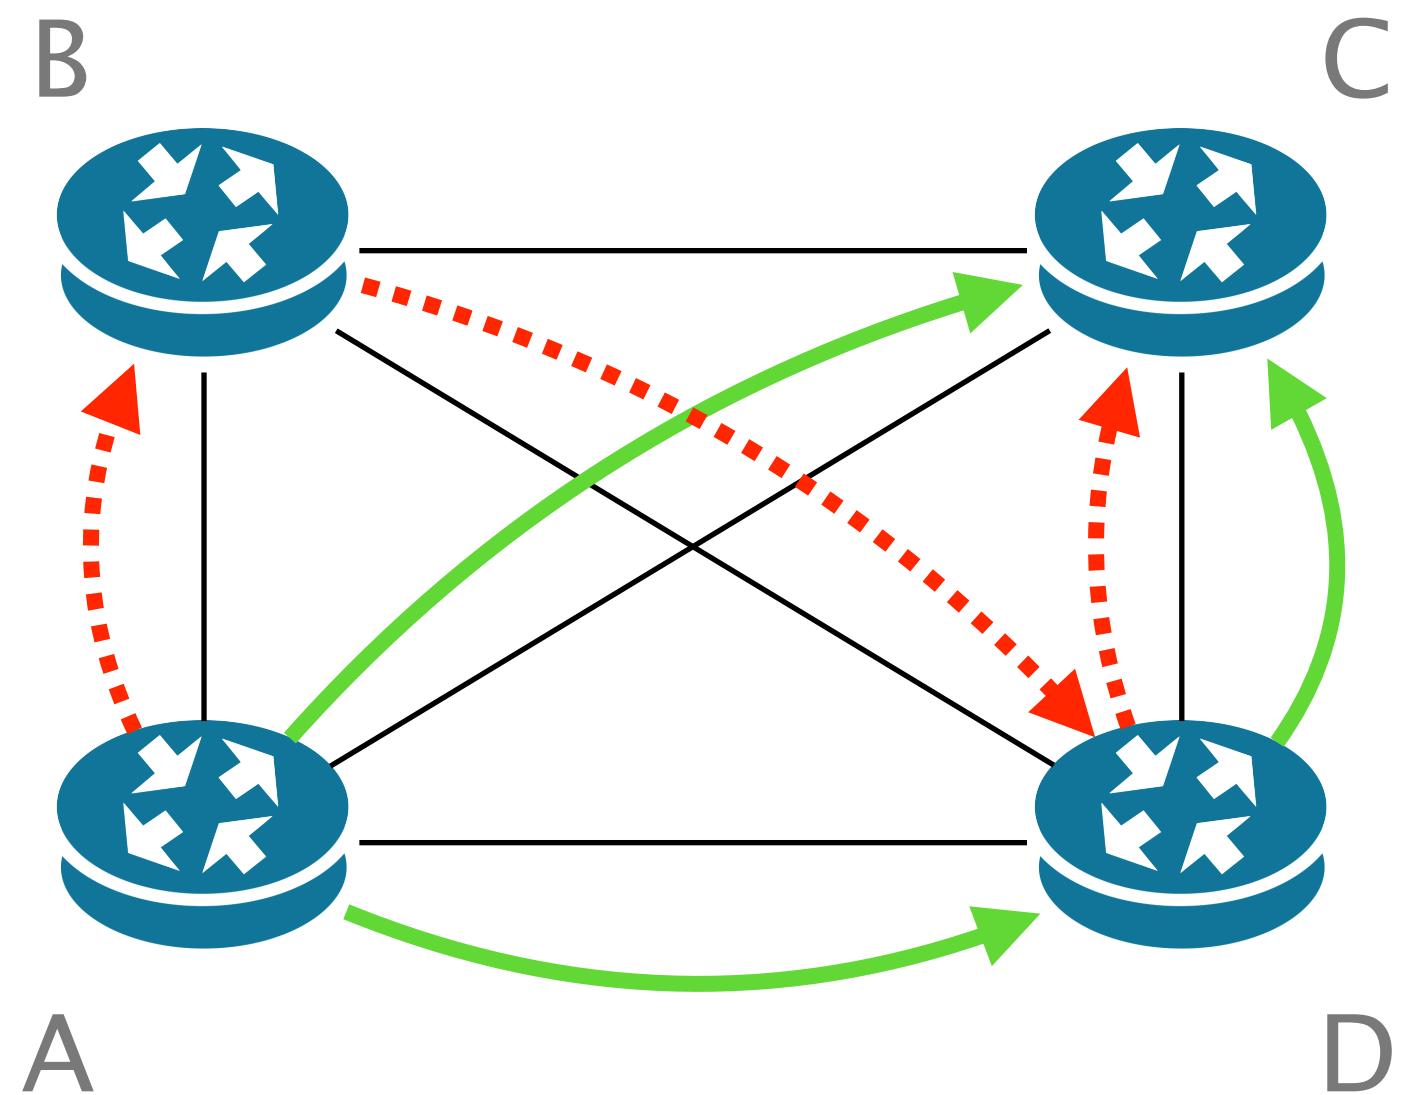


synthesis procedure

$$\forall X \in \text{SamplePaths}(A, C) \setminus \text{Reqs}$$

$$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$$

input requirements



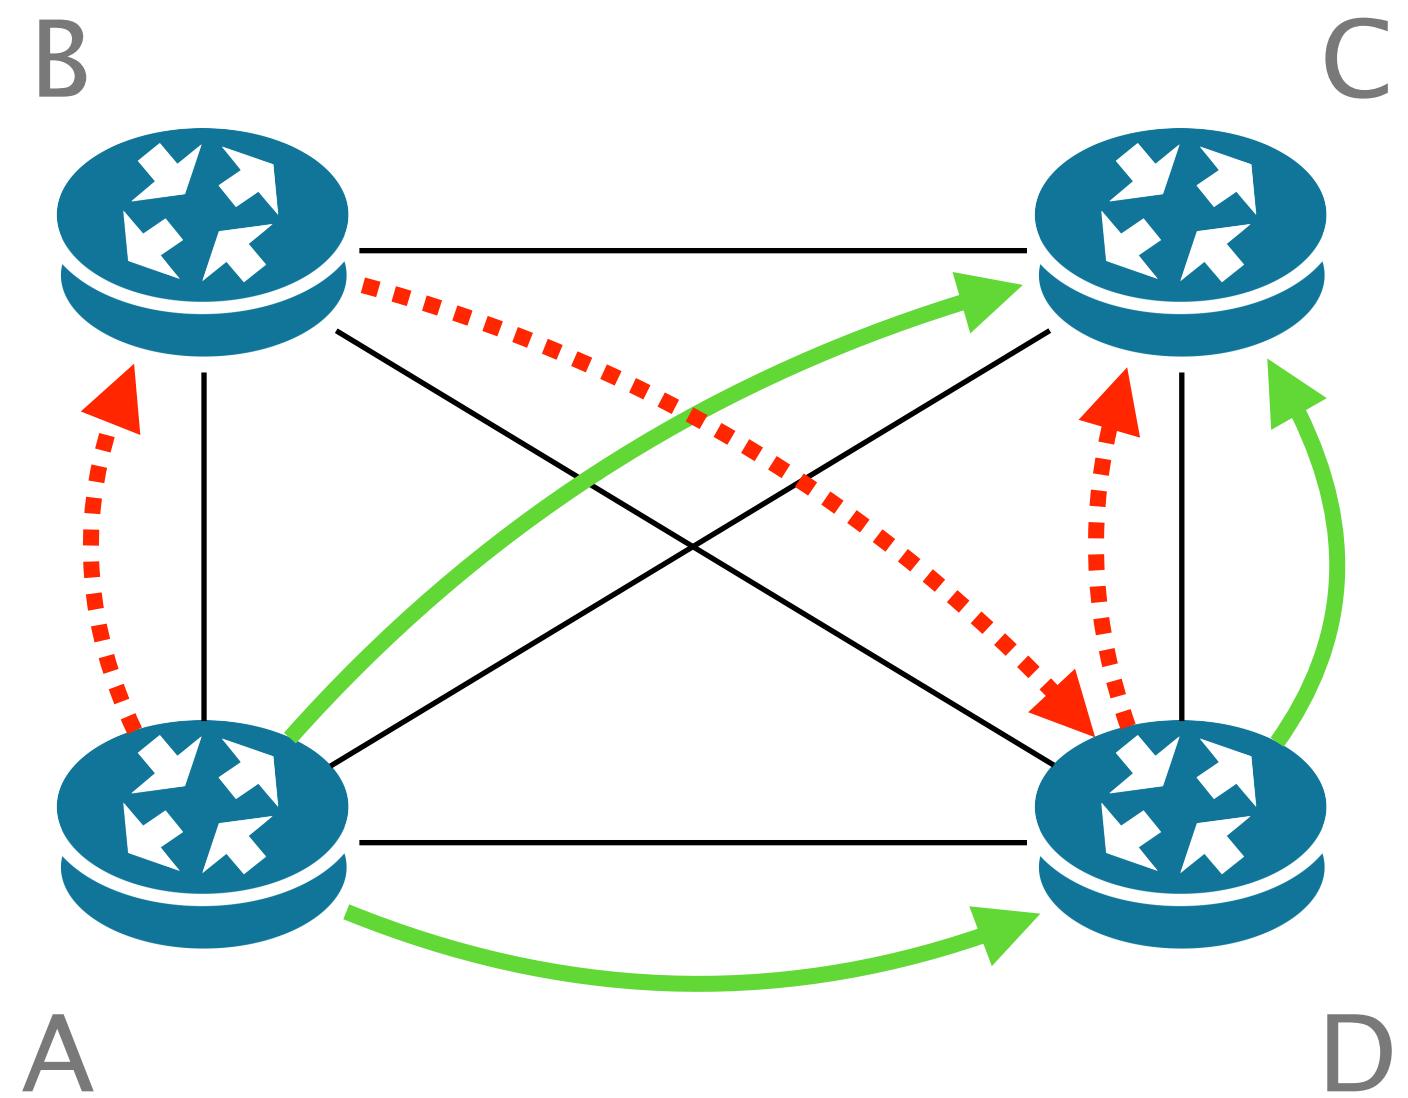
synthesis procedure

$\forall X \in \text{SamplePaths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

input requirements



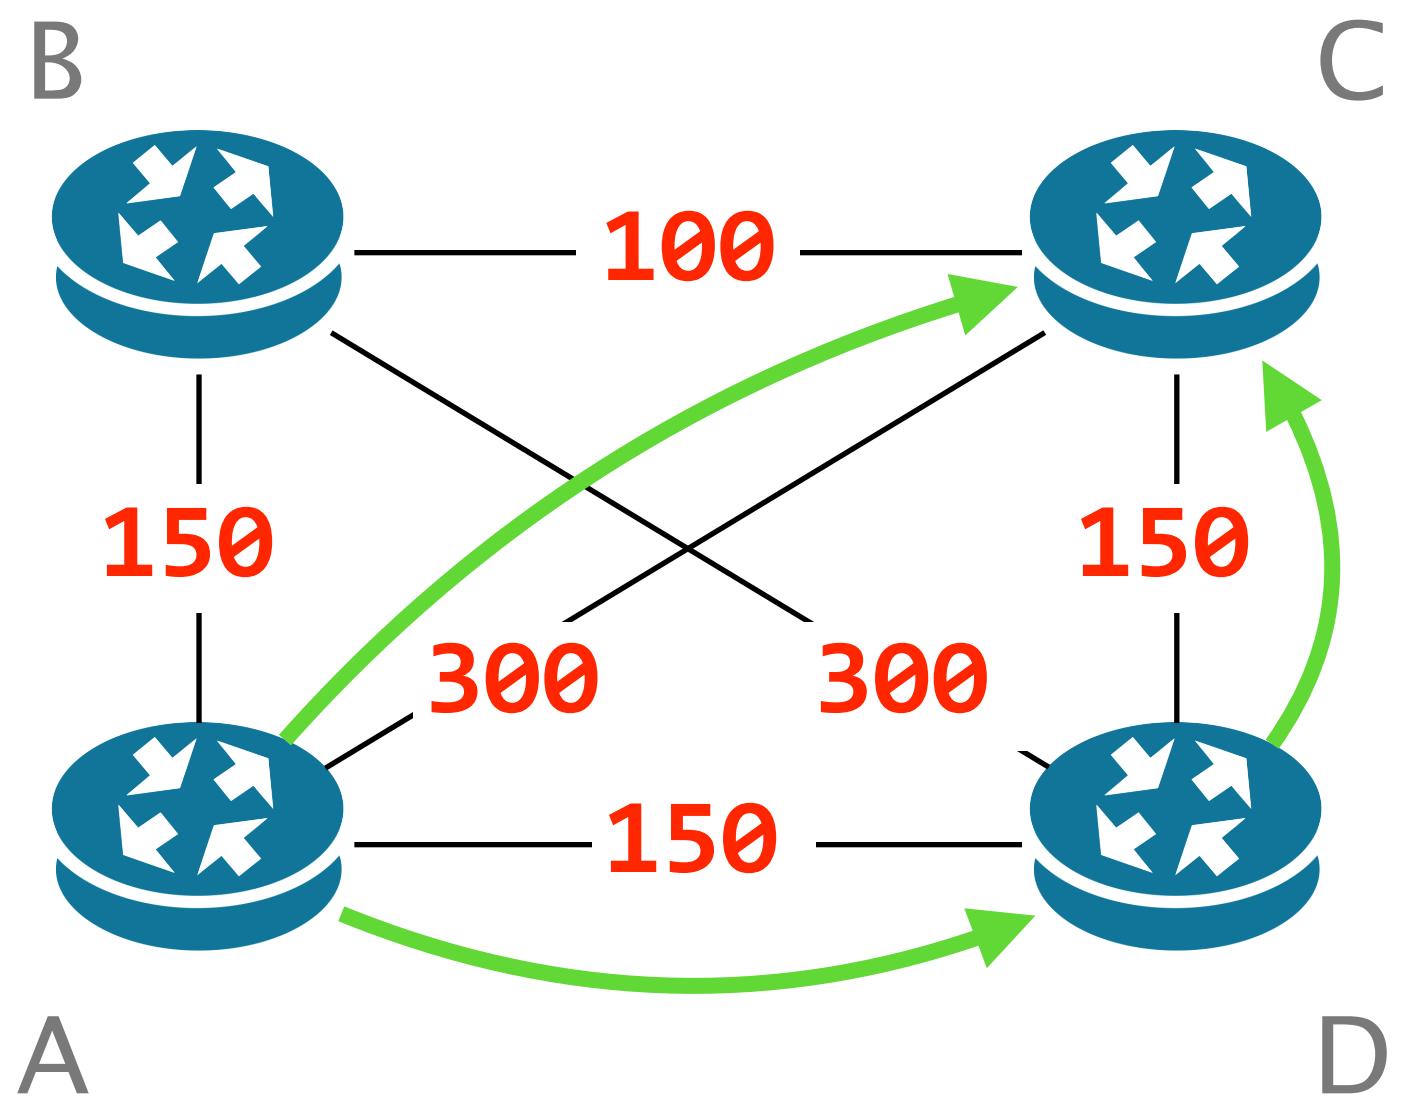
synthesis procedure

$\forall X \in \text{SamplePaths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

input requirements



Synthesized weights

synthesis procedure

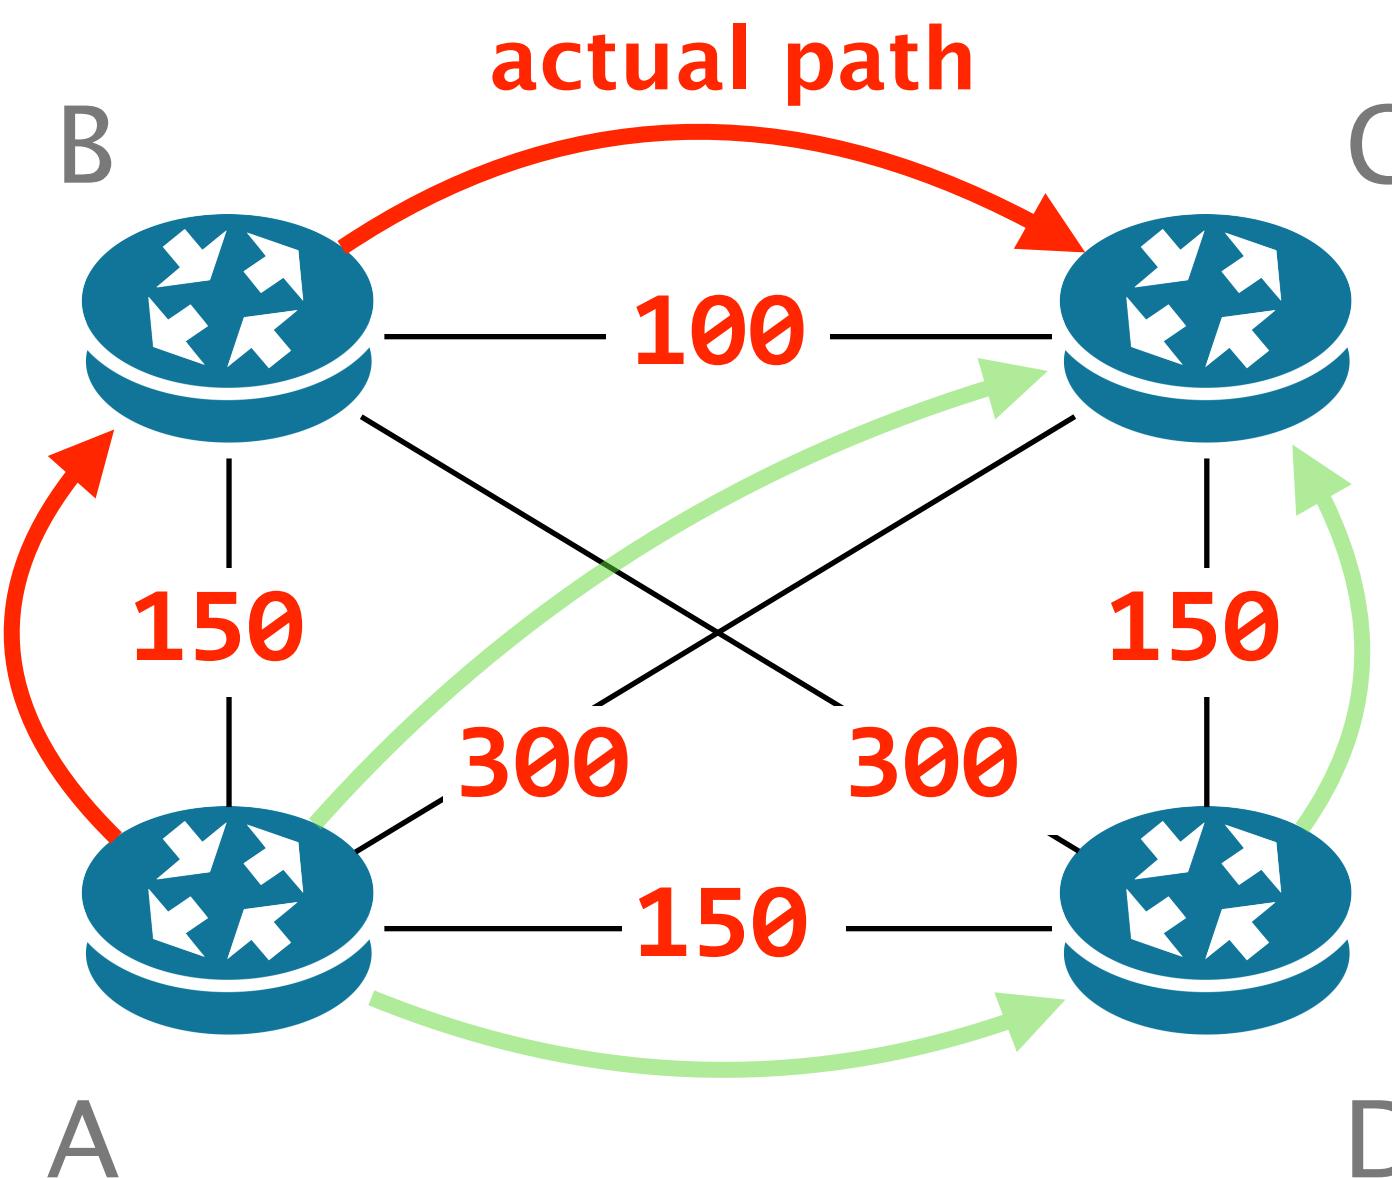
$\forall X \in \text{SamplePaths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

The synthesized weights are incorrect:

$$\text{cost}(A \rightarrow B \rightarrow C) = 250 < \text{cost}(A \rightarrow C) = 300$$

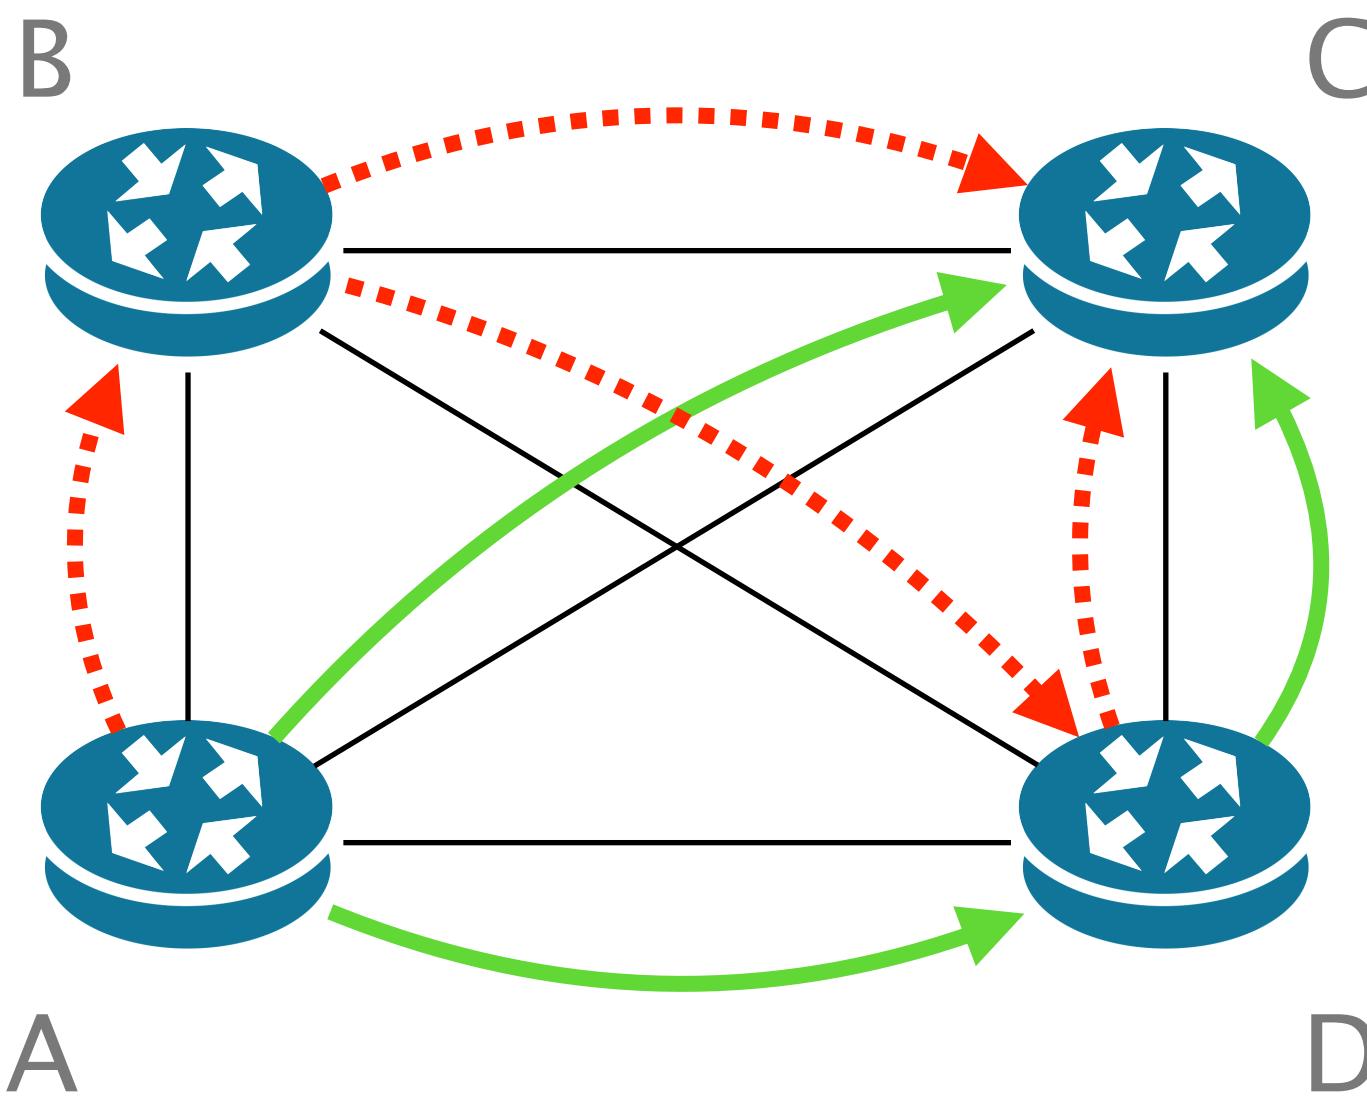


$\forall X \in \text{SamplePaths}(A,C) \setminus \text{Reqs}$

$\text{Cost}(A \rightarrow C) = \text{Cost}(A \rightarrow D \rightarrow C) < \text{Cost}(X)$

Solve

We simply add the counter example to
SamplePaths and repeat the procedure


$$\forall x \in \text{SamplePaths}(A,C) \setminus \text{Reqs}$$

Sample: { [A,B,D,C] } \cup { [A,B,C] }

The entire procedure usually converges in few iterations
making it very fast in practice

Self-Driving Networks

Breaking new ground in network automation



- 1 operator assistance
- 2 partial automation

Part I
assisting operators

- 3 conditional automation
- 4 high automation

Part II
taking control

- 5 full automation

too futuristic? 😊

Part II: Taking control

applications

frameworks

methods

data-driven
convergence

monitoring/inference

some of our challenges
lying ahead

Part II: Taking control

applications

frameworks

methods

data-driven
convergence

Blink

Fast Connectivity Recovery Entirely in the Data Plane



Thomas Holterbach
Costa Molero



Edgar
Costa Molero



Maria Apostolaki



Alberto Dainotti

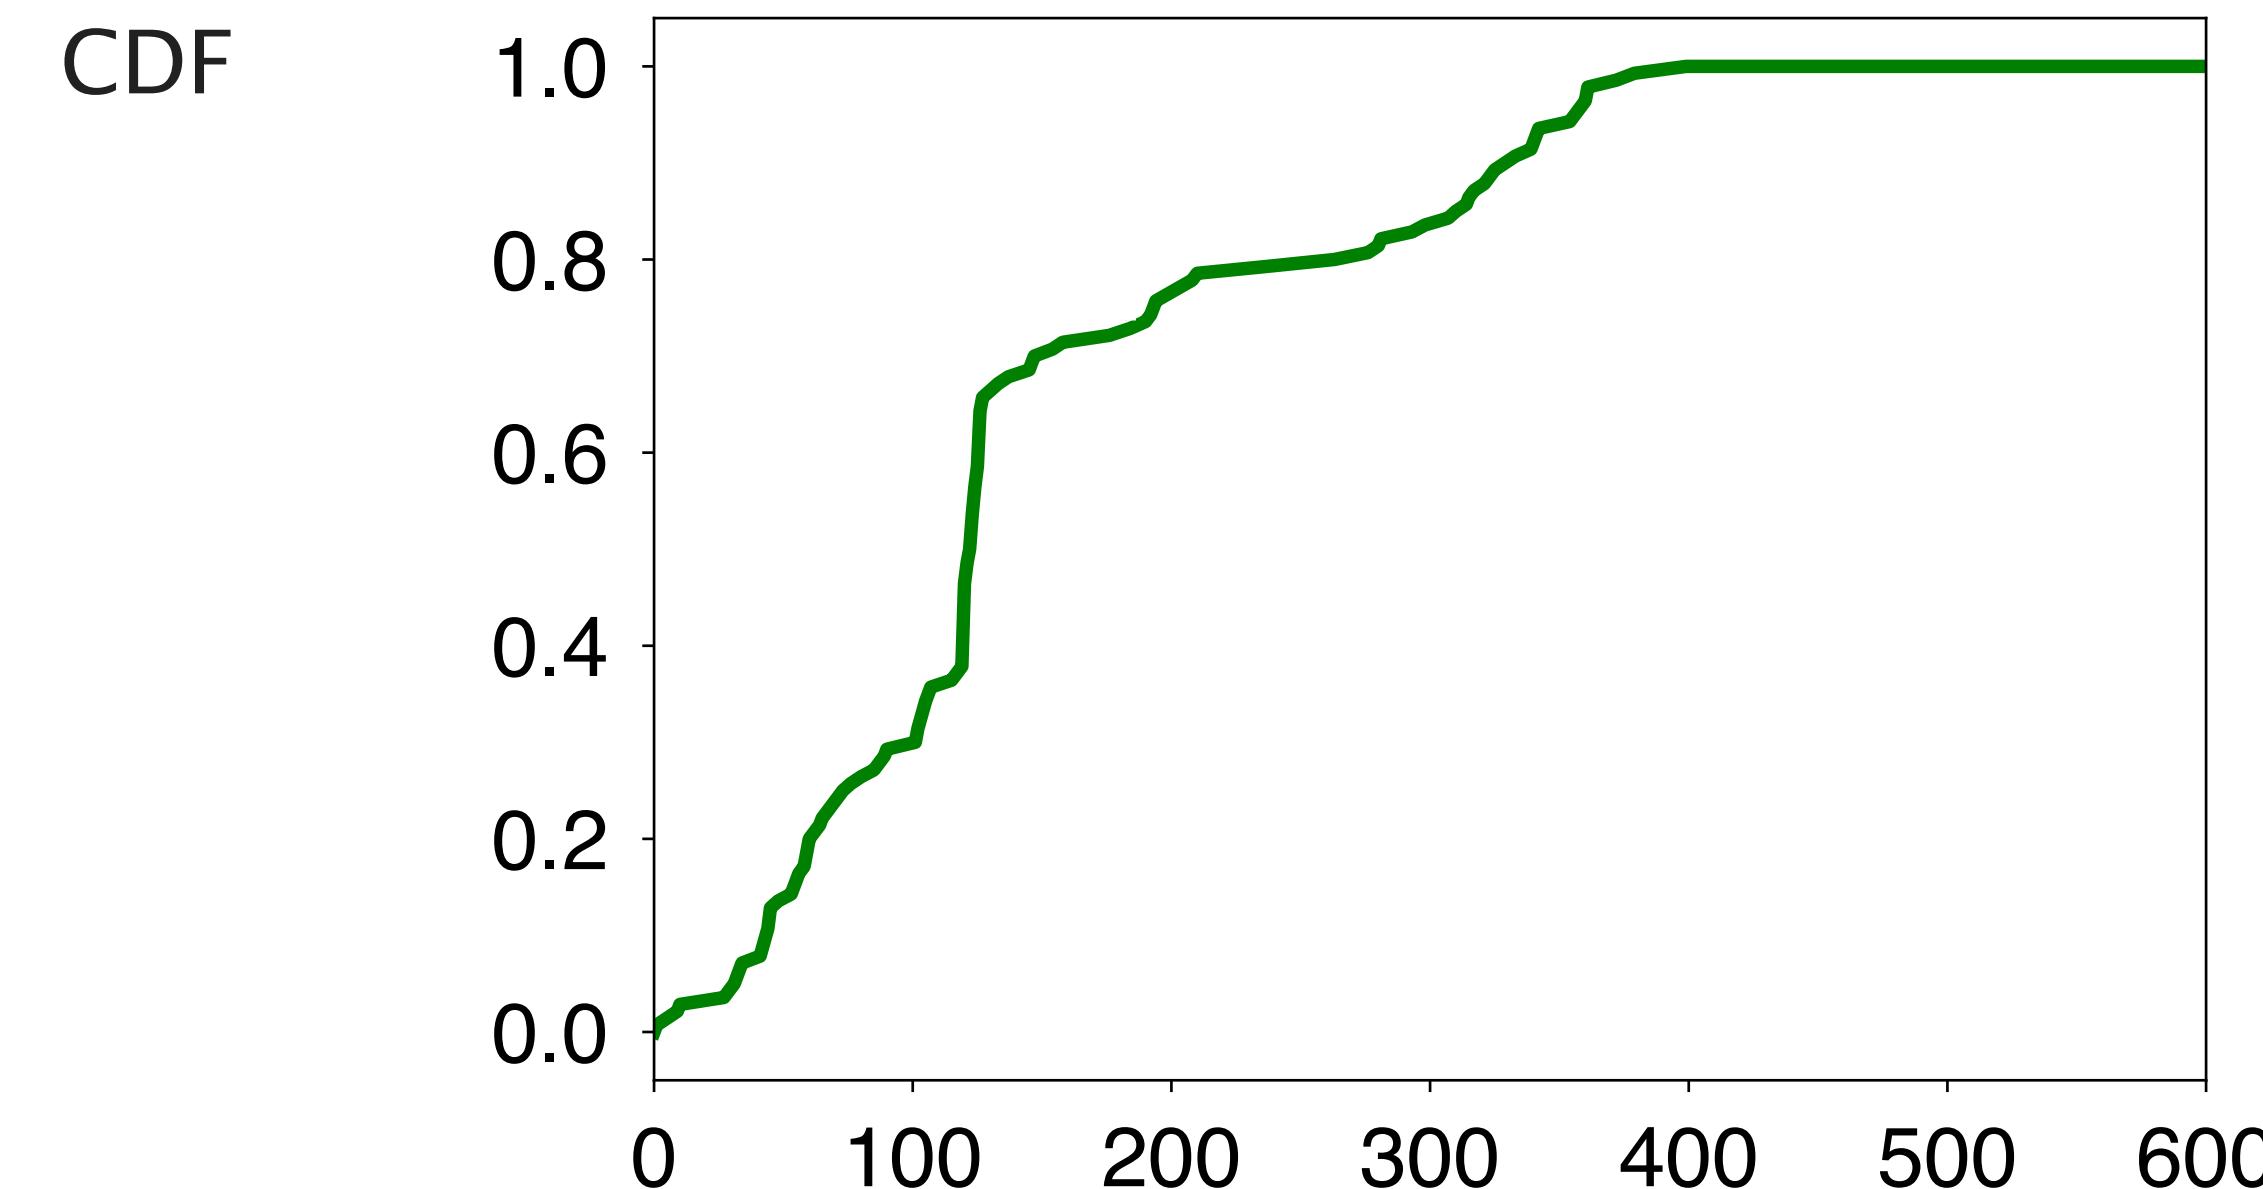


Stefano Vissicchio



Laurent Vanbever

Internet control plane can take **minutes** to converge
upon remote failures

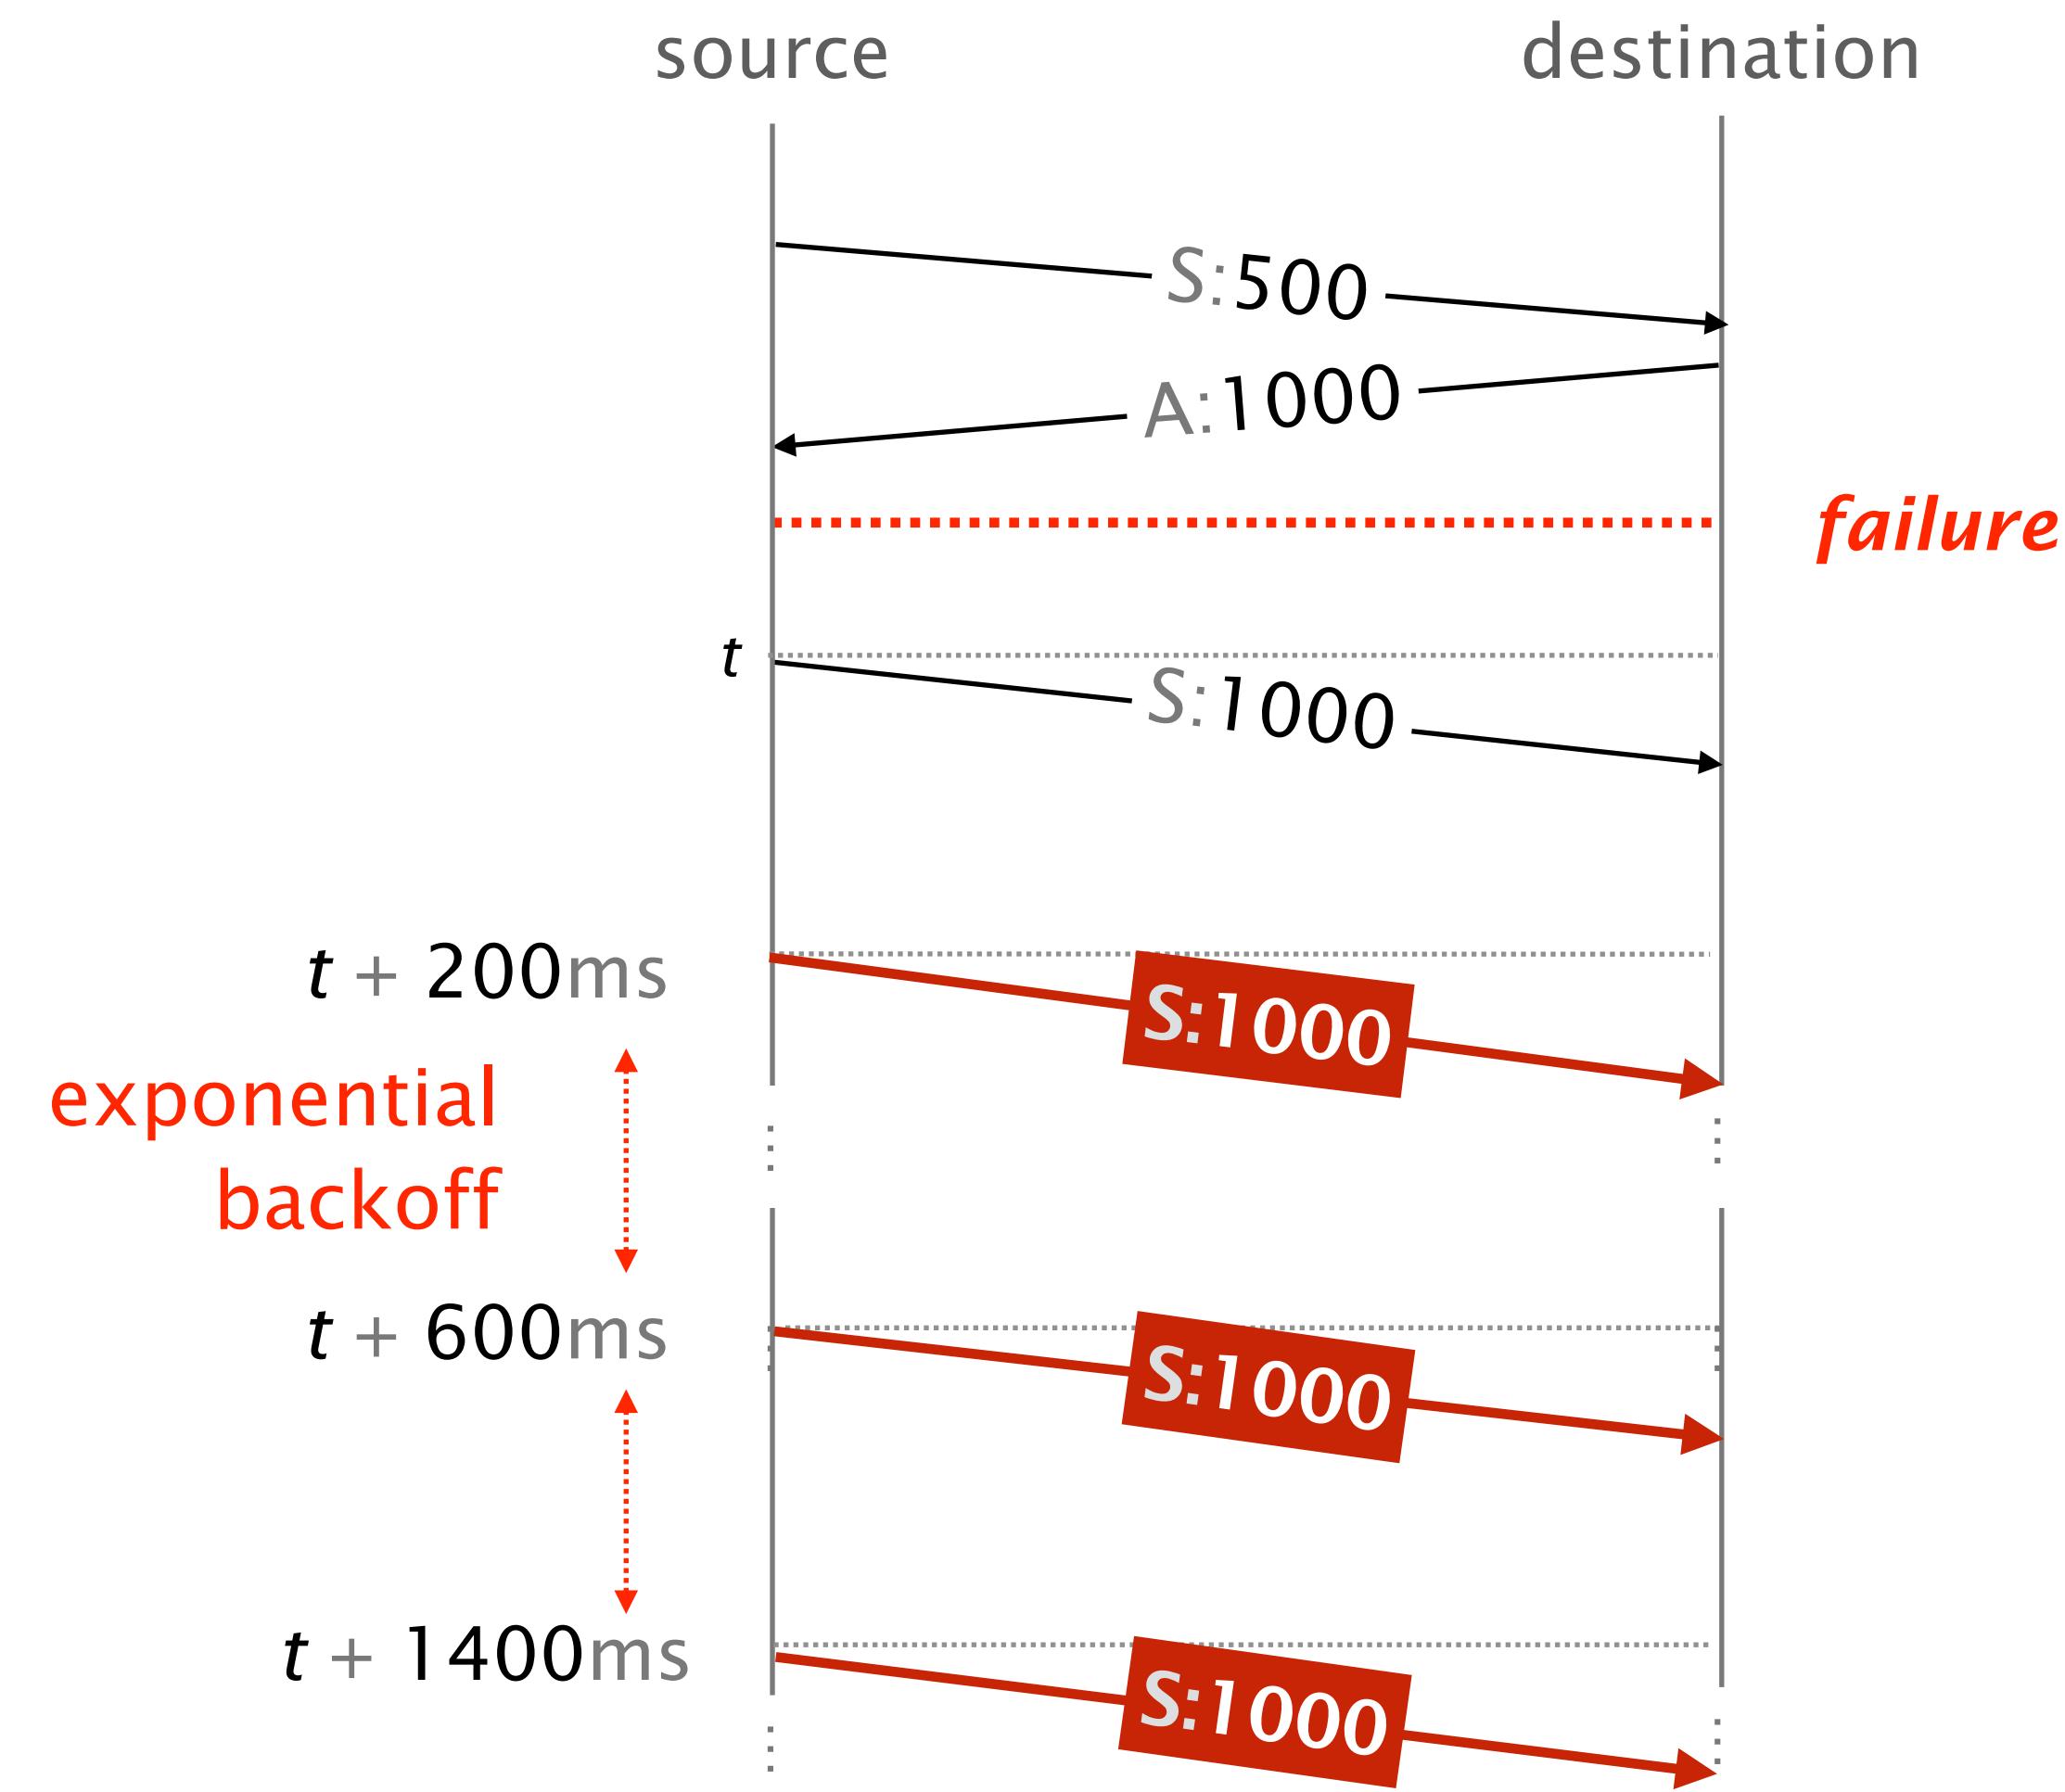


Second elapsed between the data plane failure
and the control plane learning about it

While the control plane can take minutes to learn about a failure,
the actual Internet traffic is affected almost instantaneously

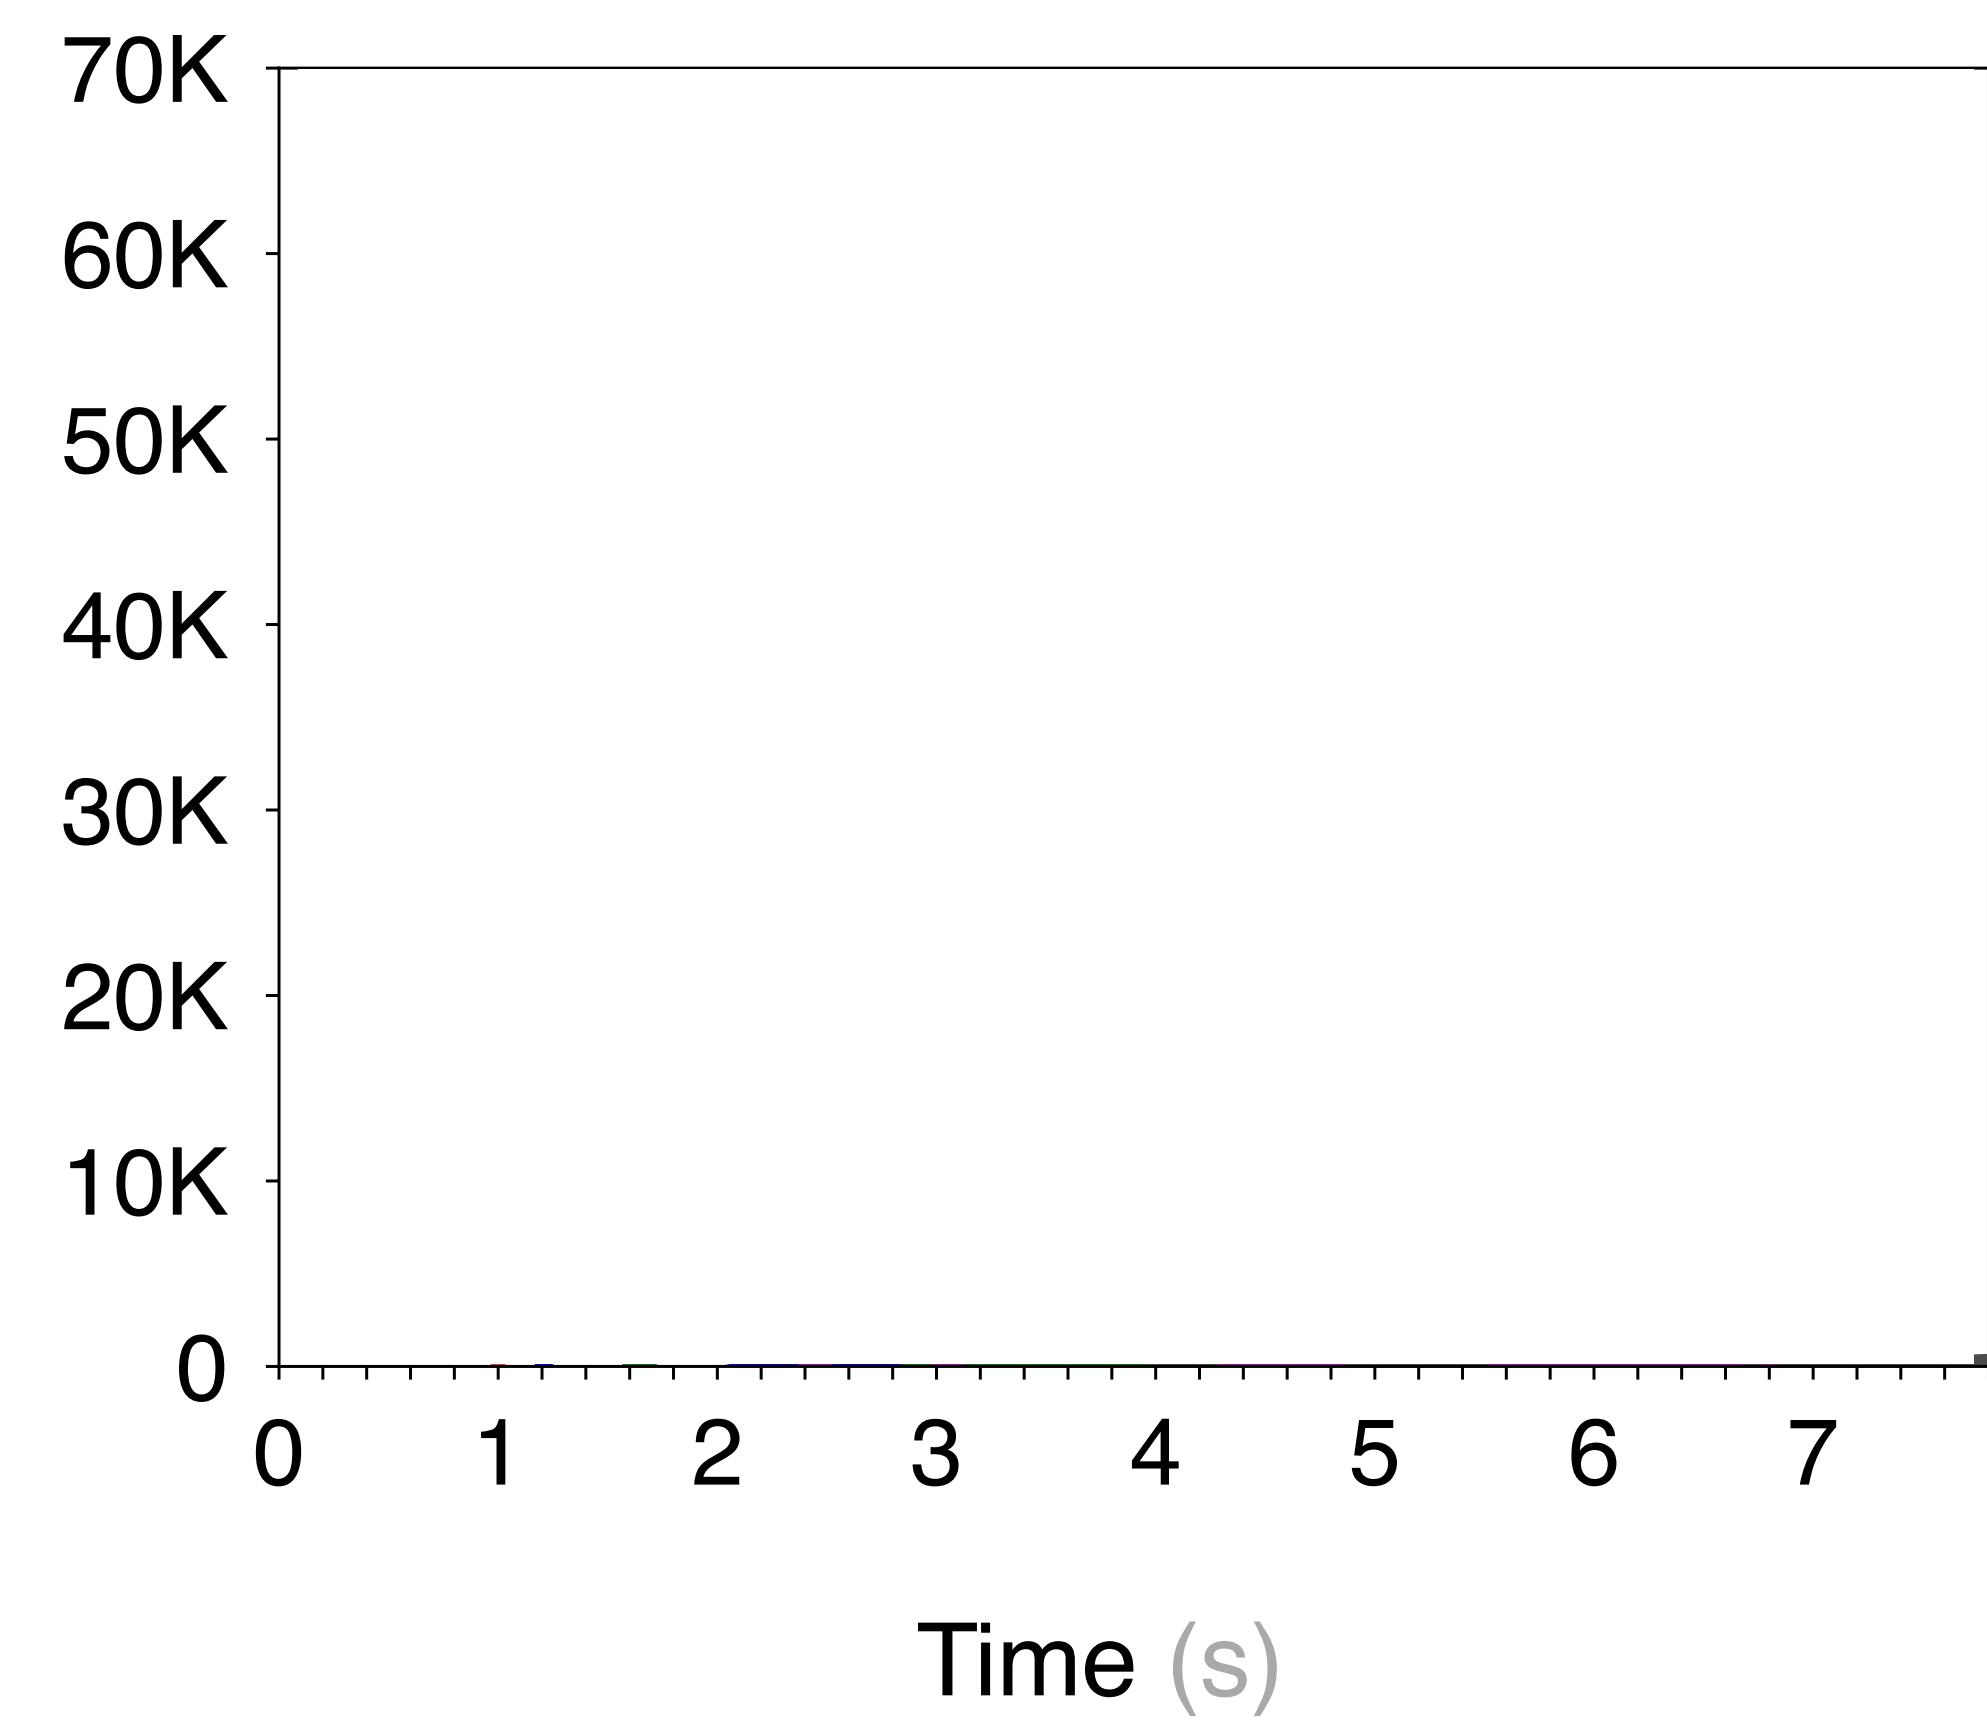
How about we track this signal instead?

Internet traffic end-points retransmit packets
upon experiencing packet drops

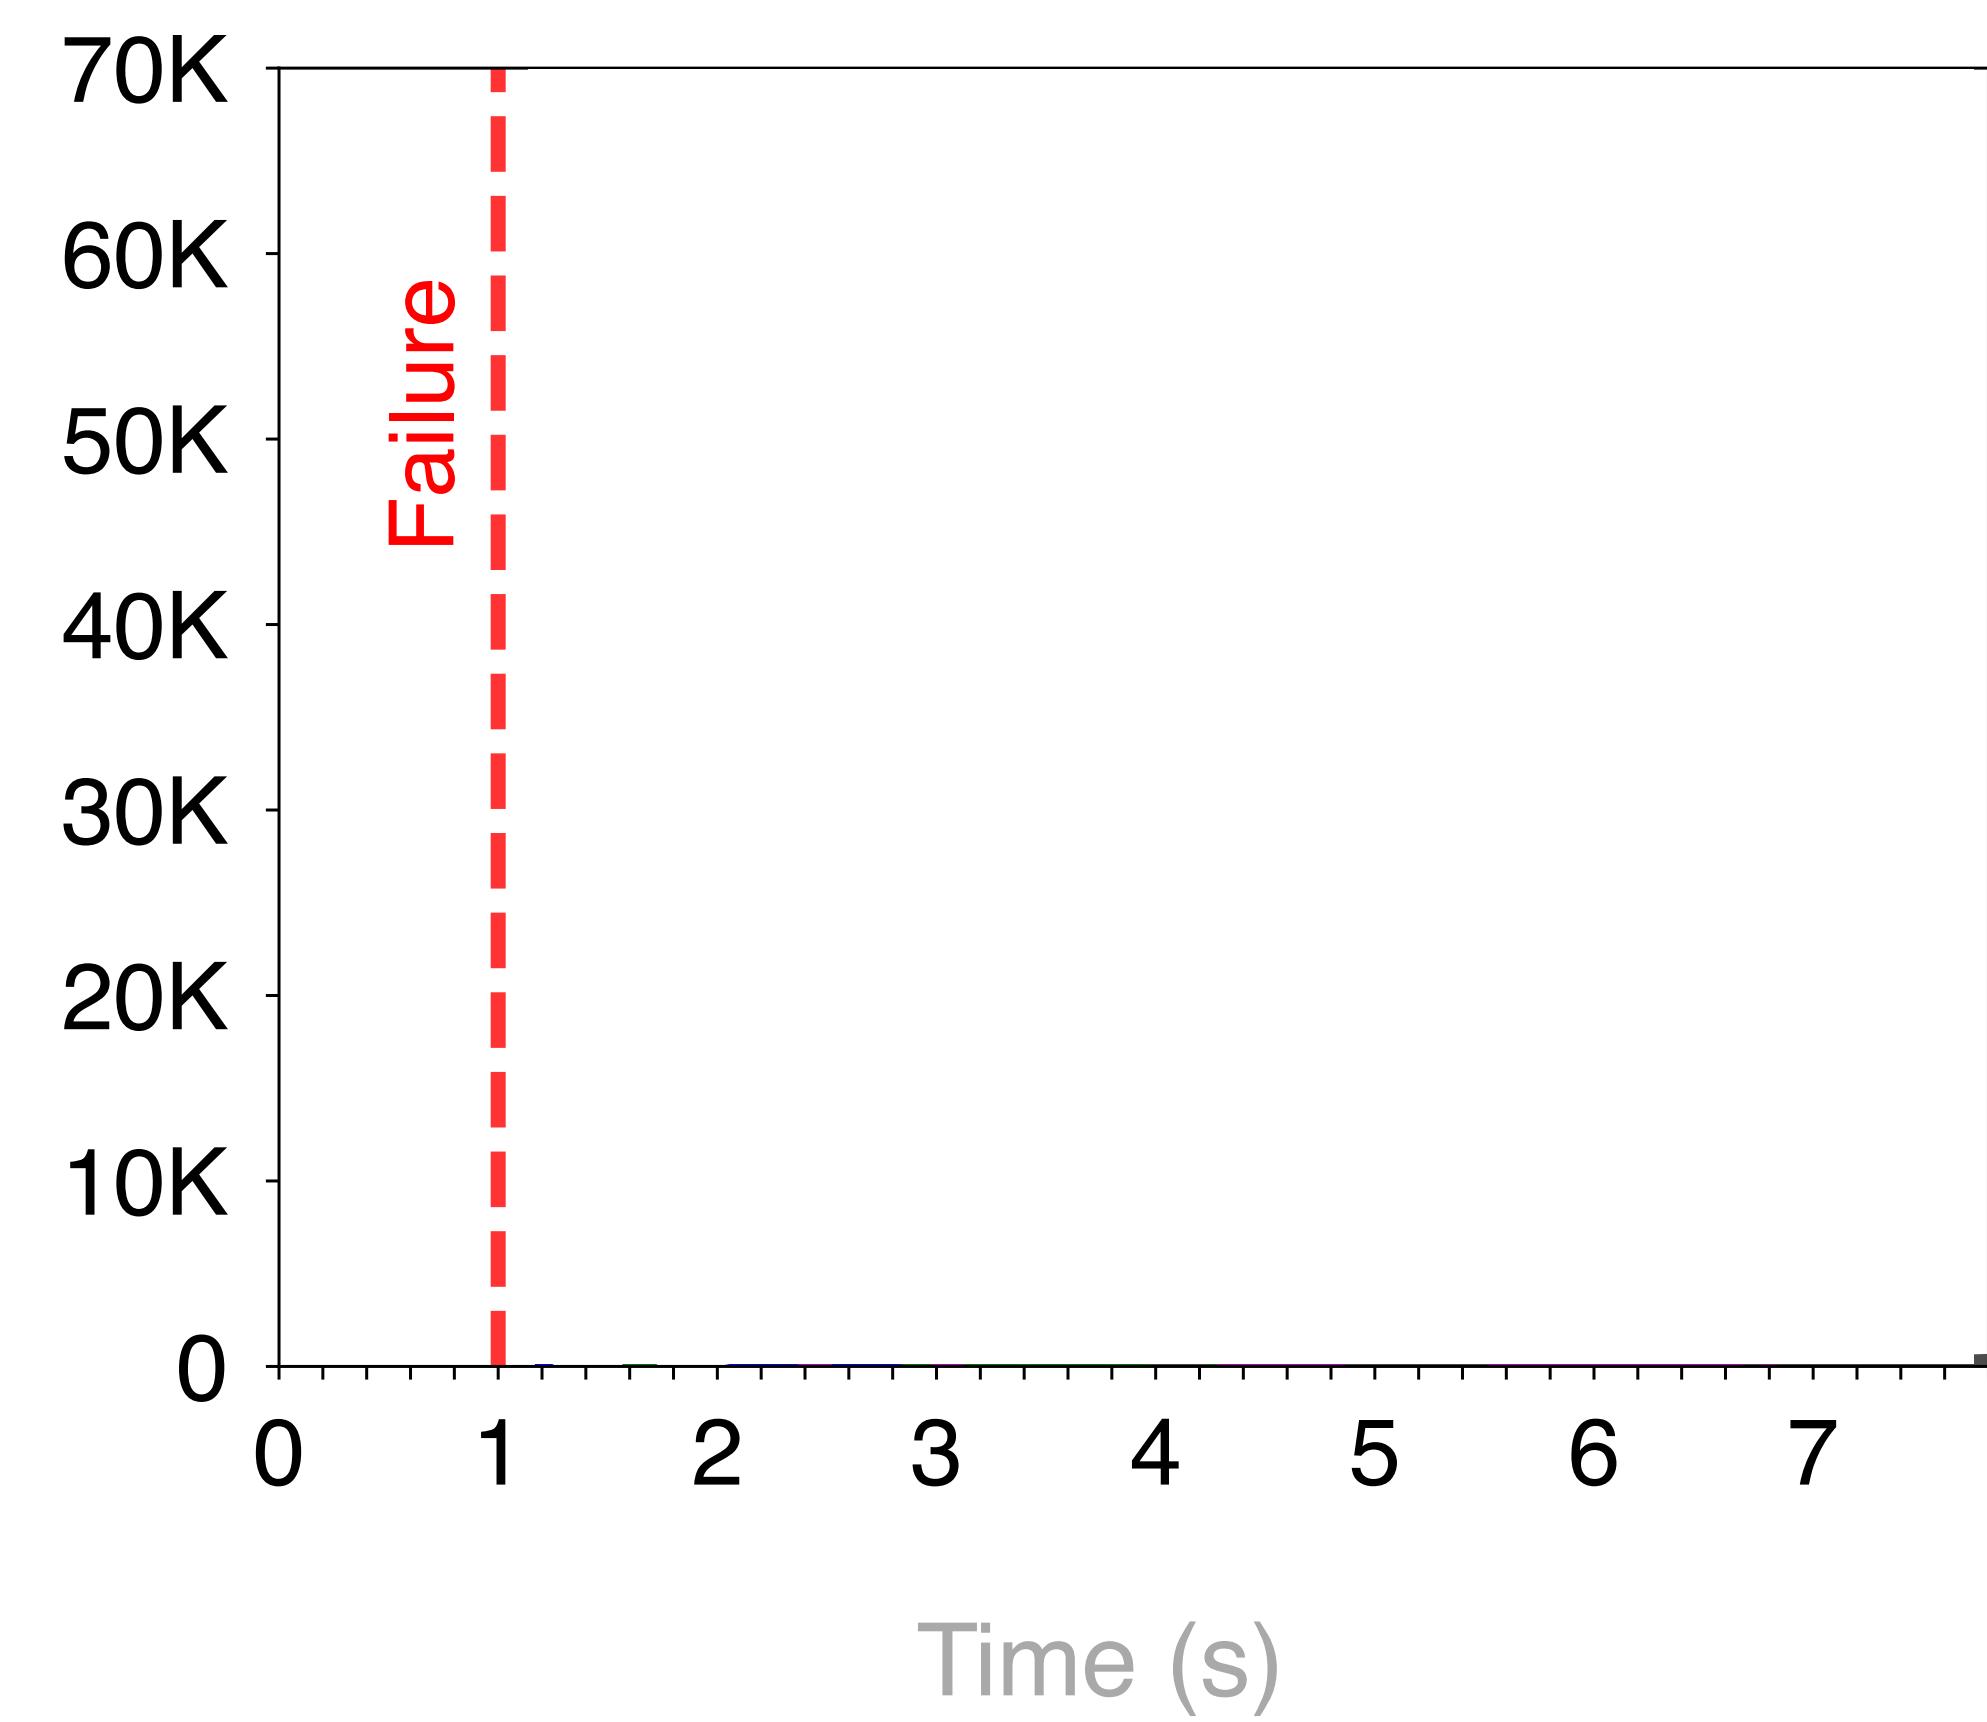


Many end-points retransmitting leads to
waves of retransmission

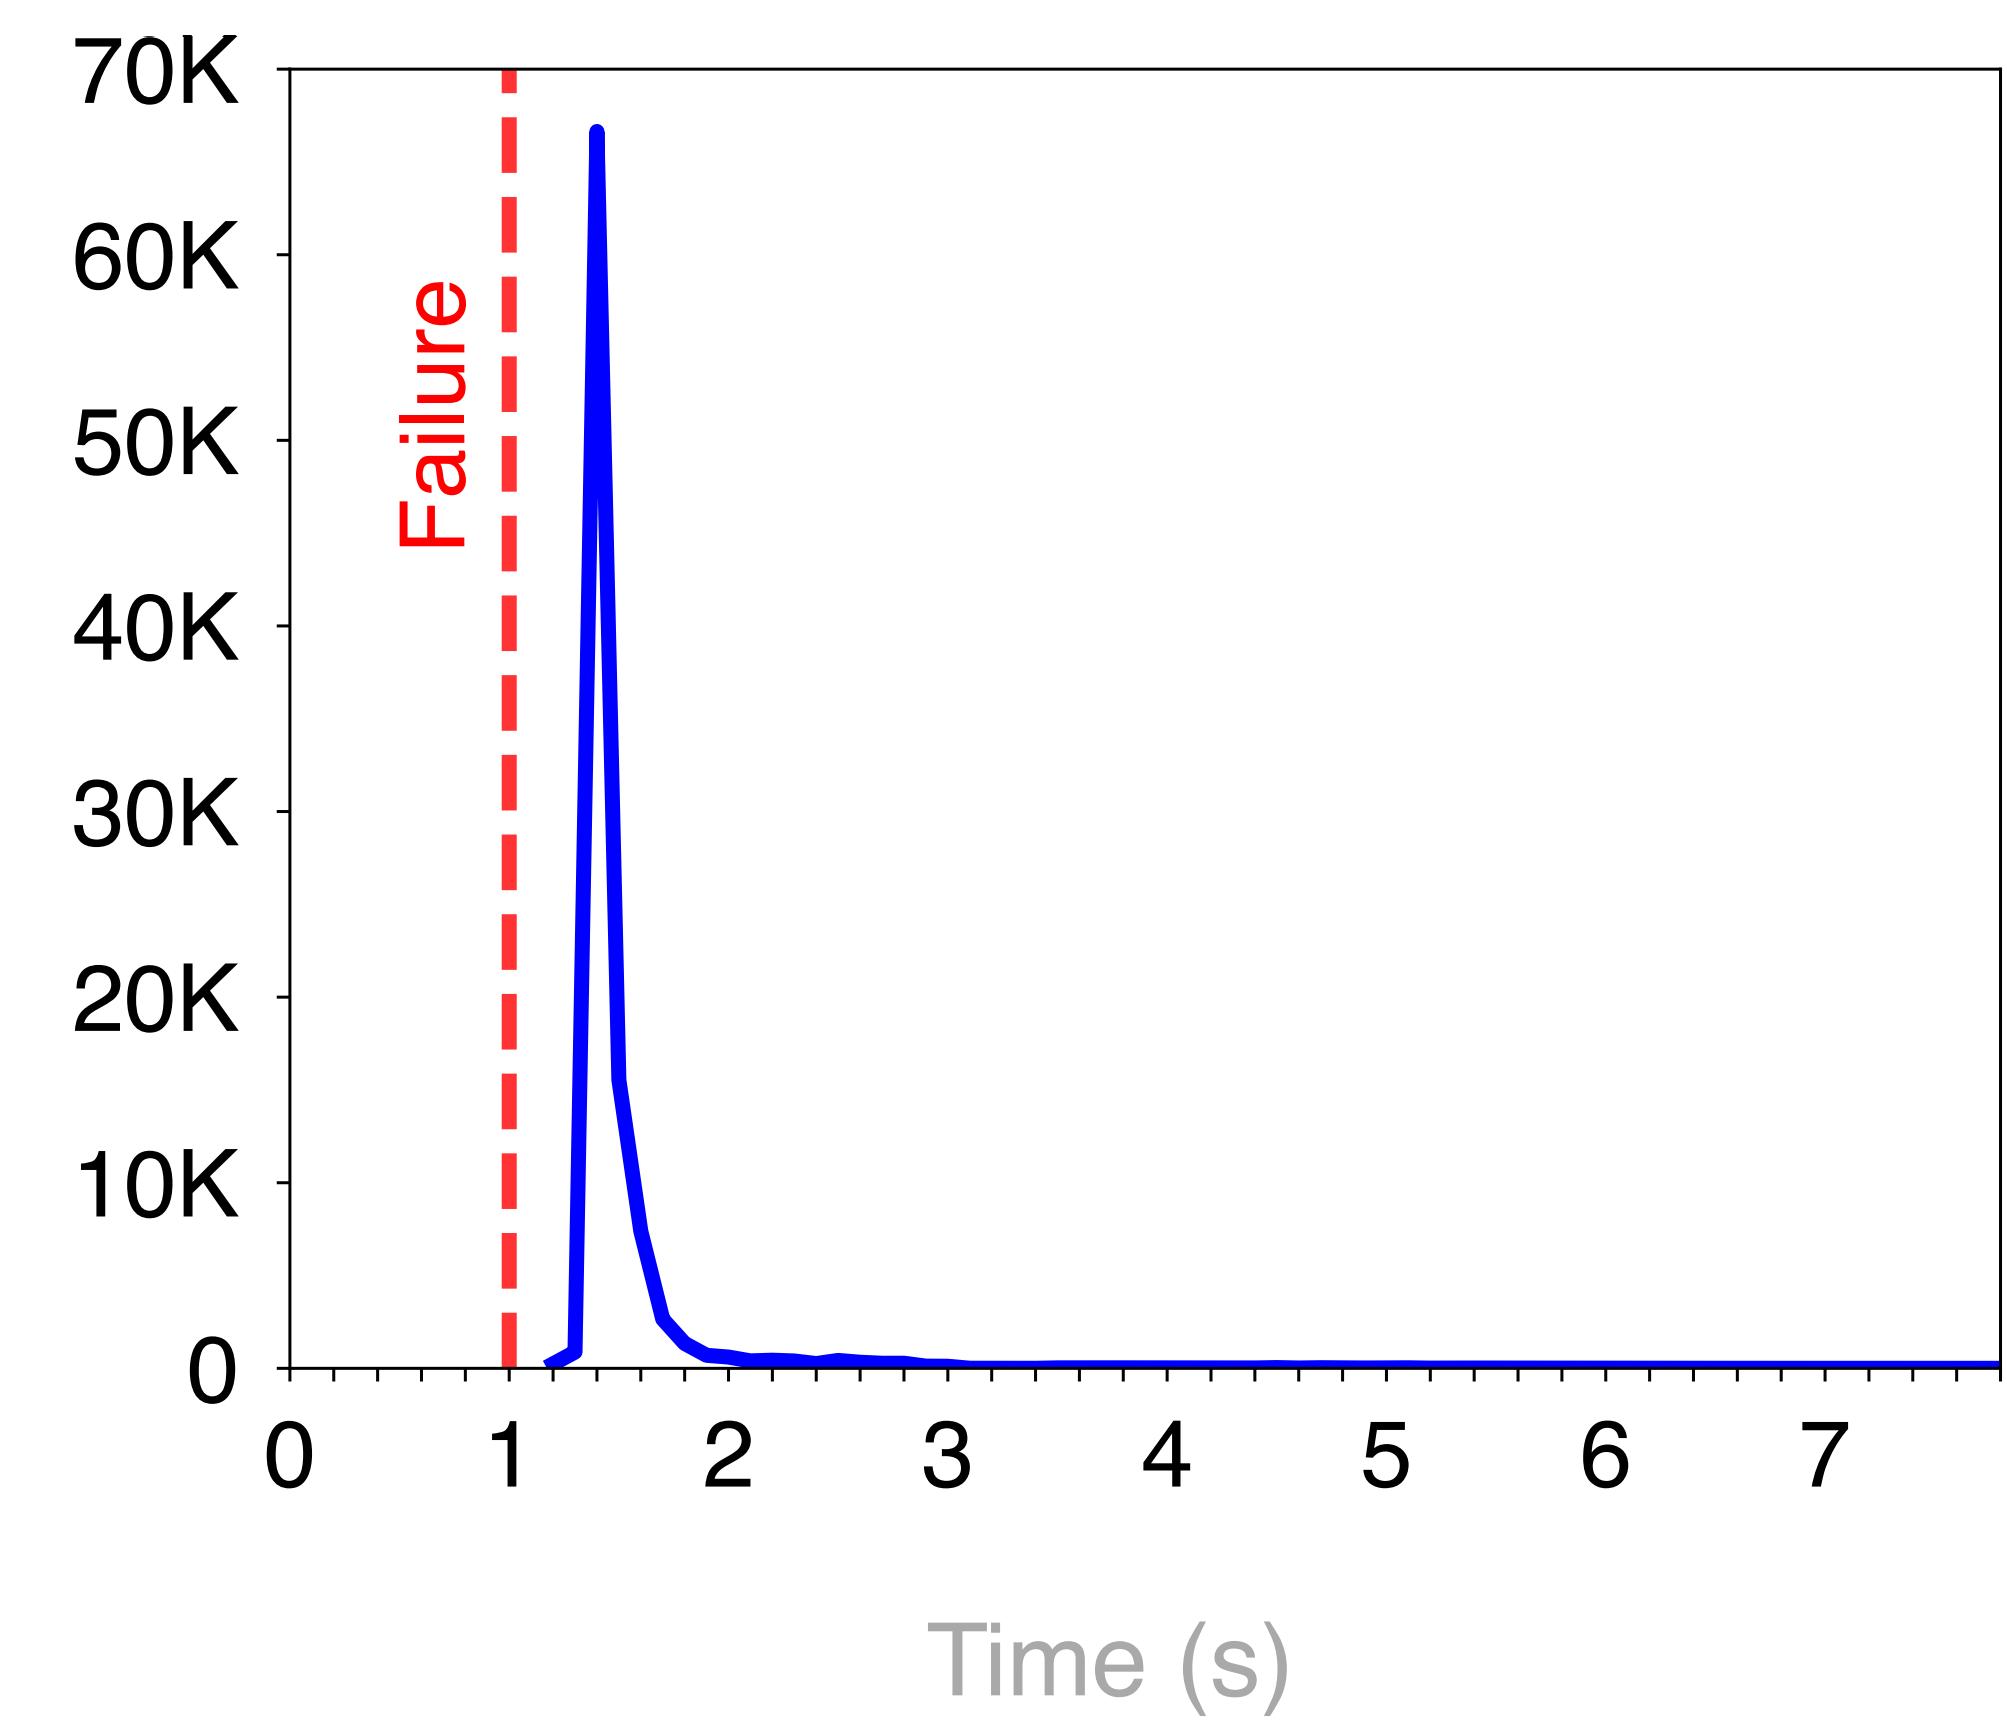
Number of
retransmissions



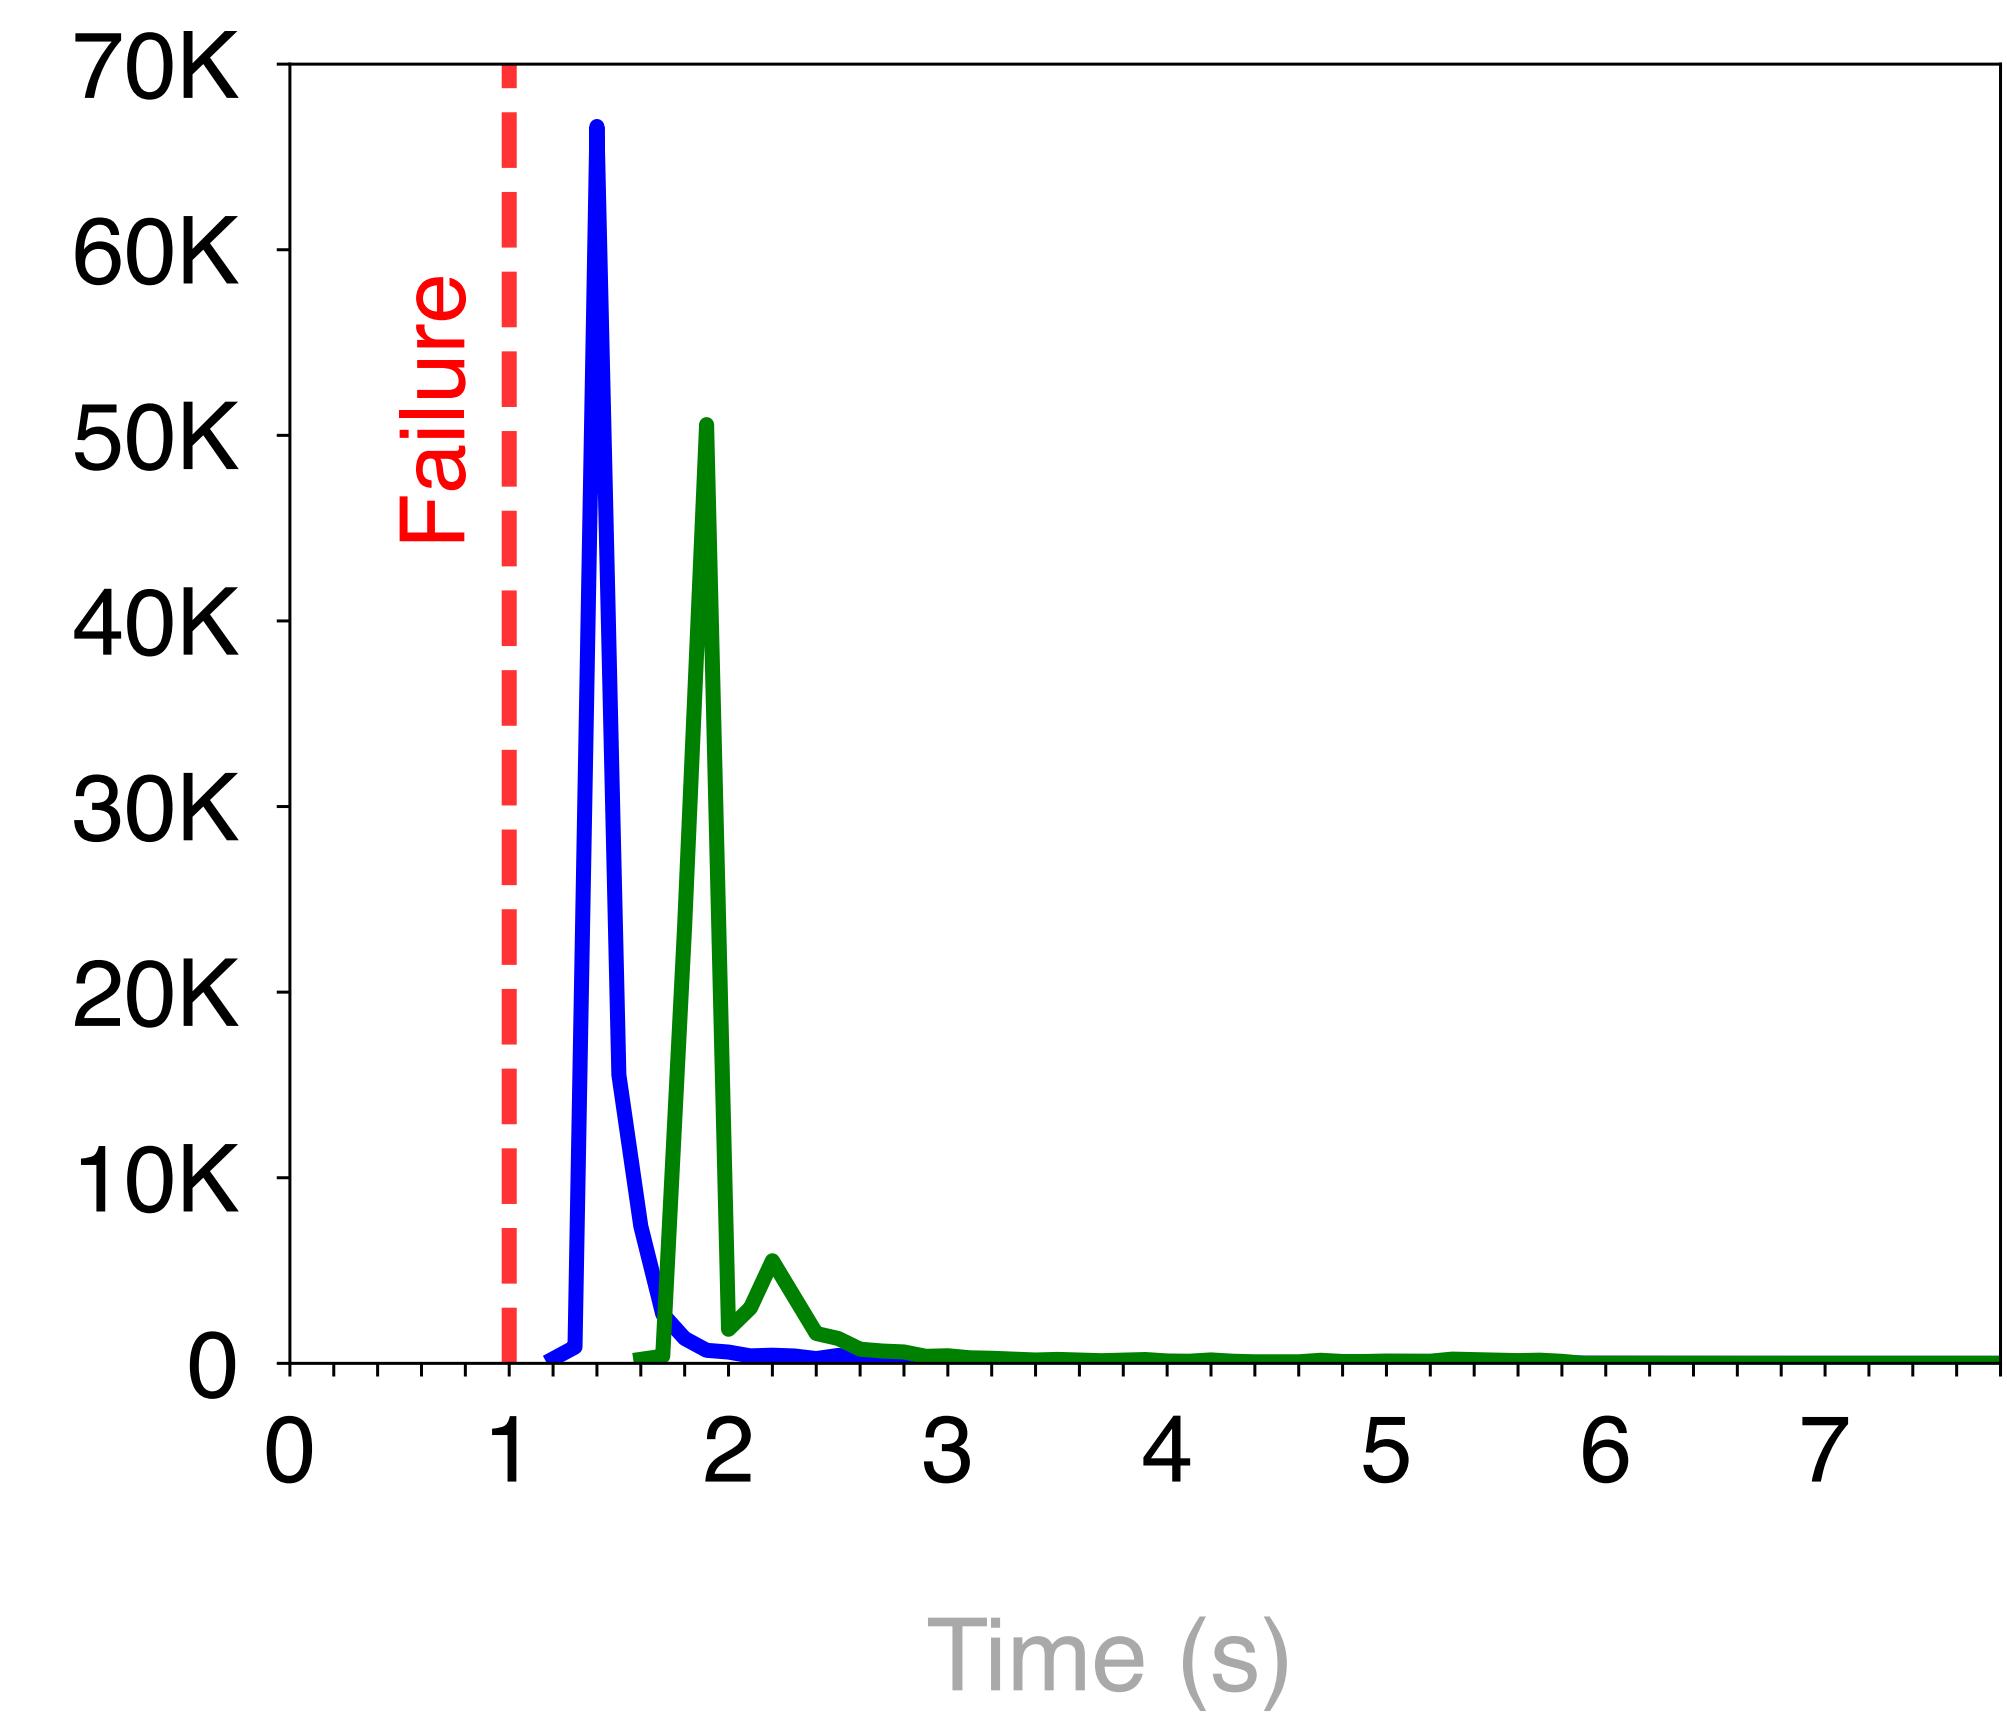
Number of
retransmissions



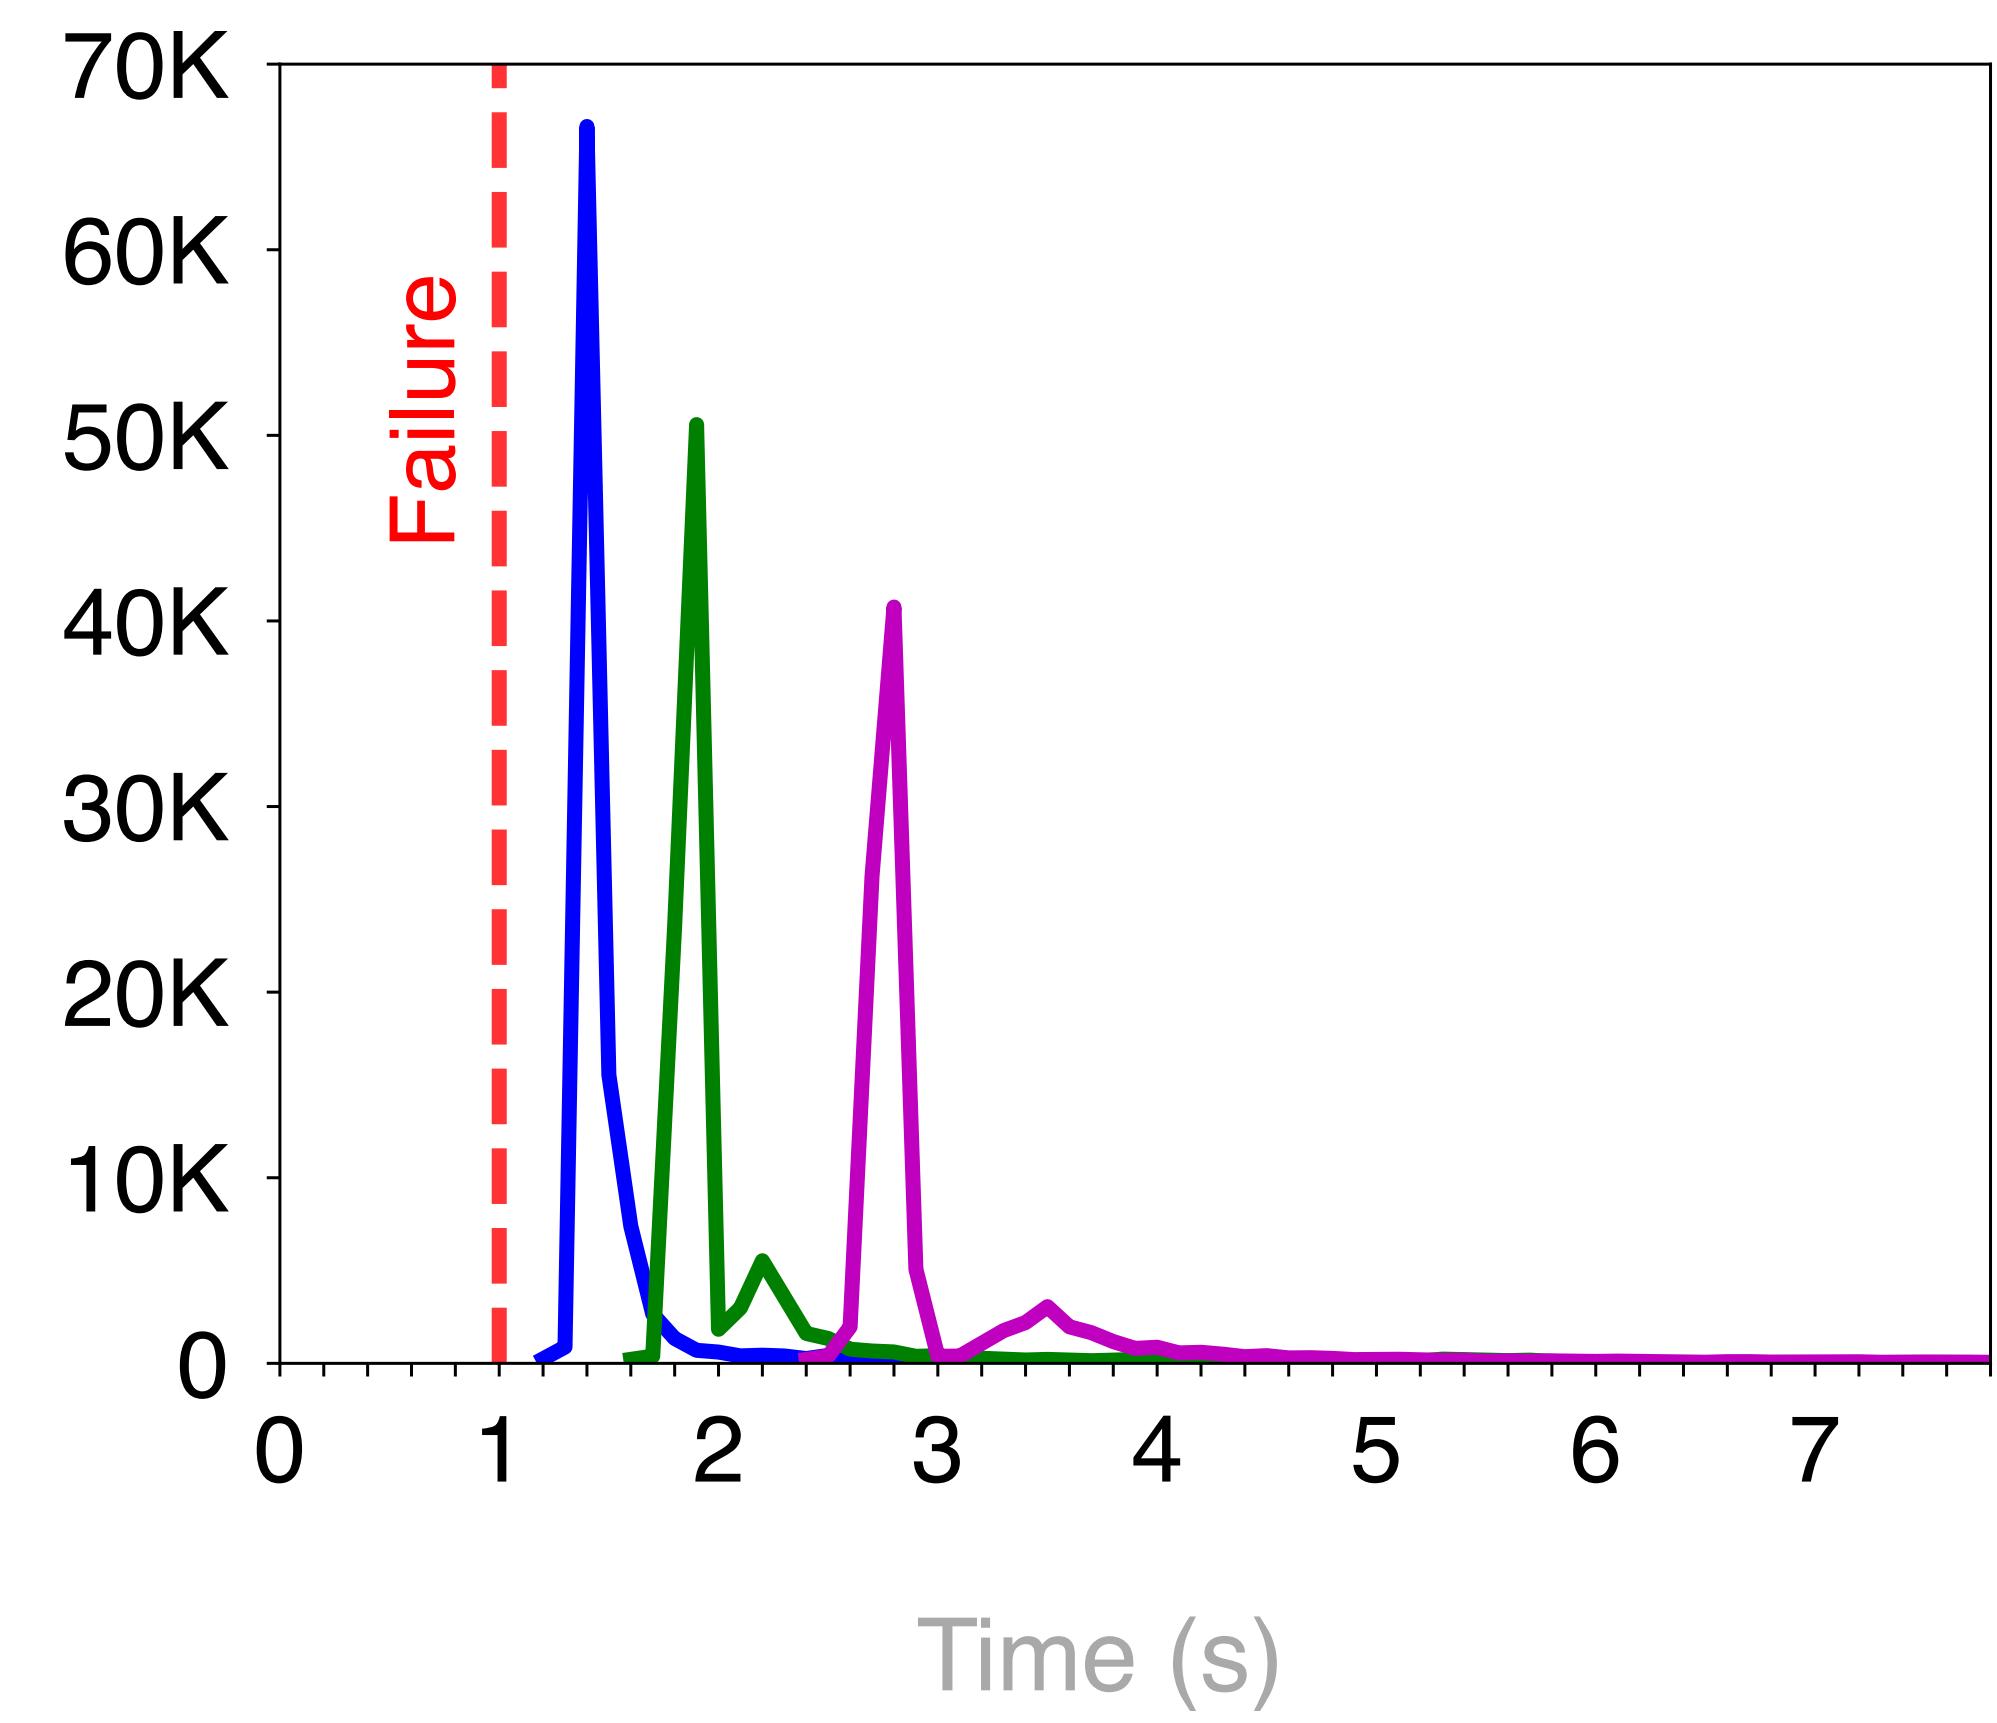
Number of
retransmissions



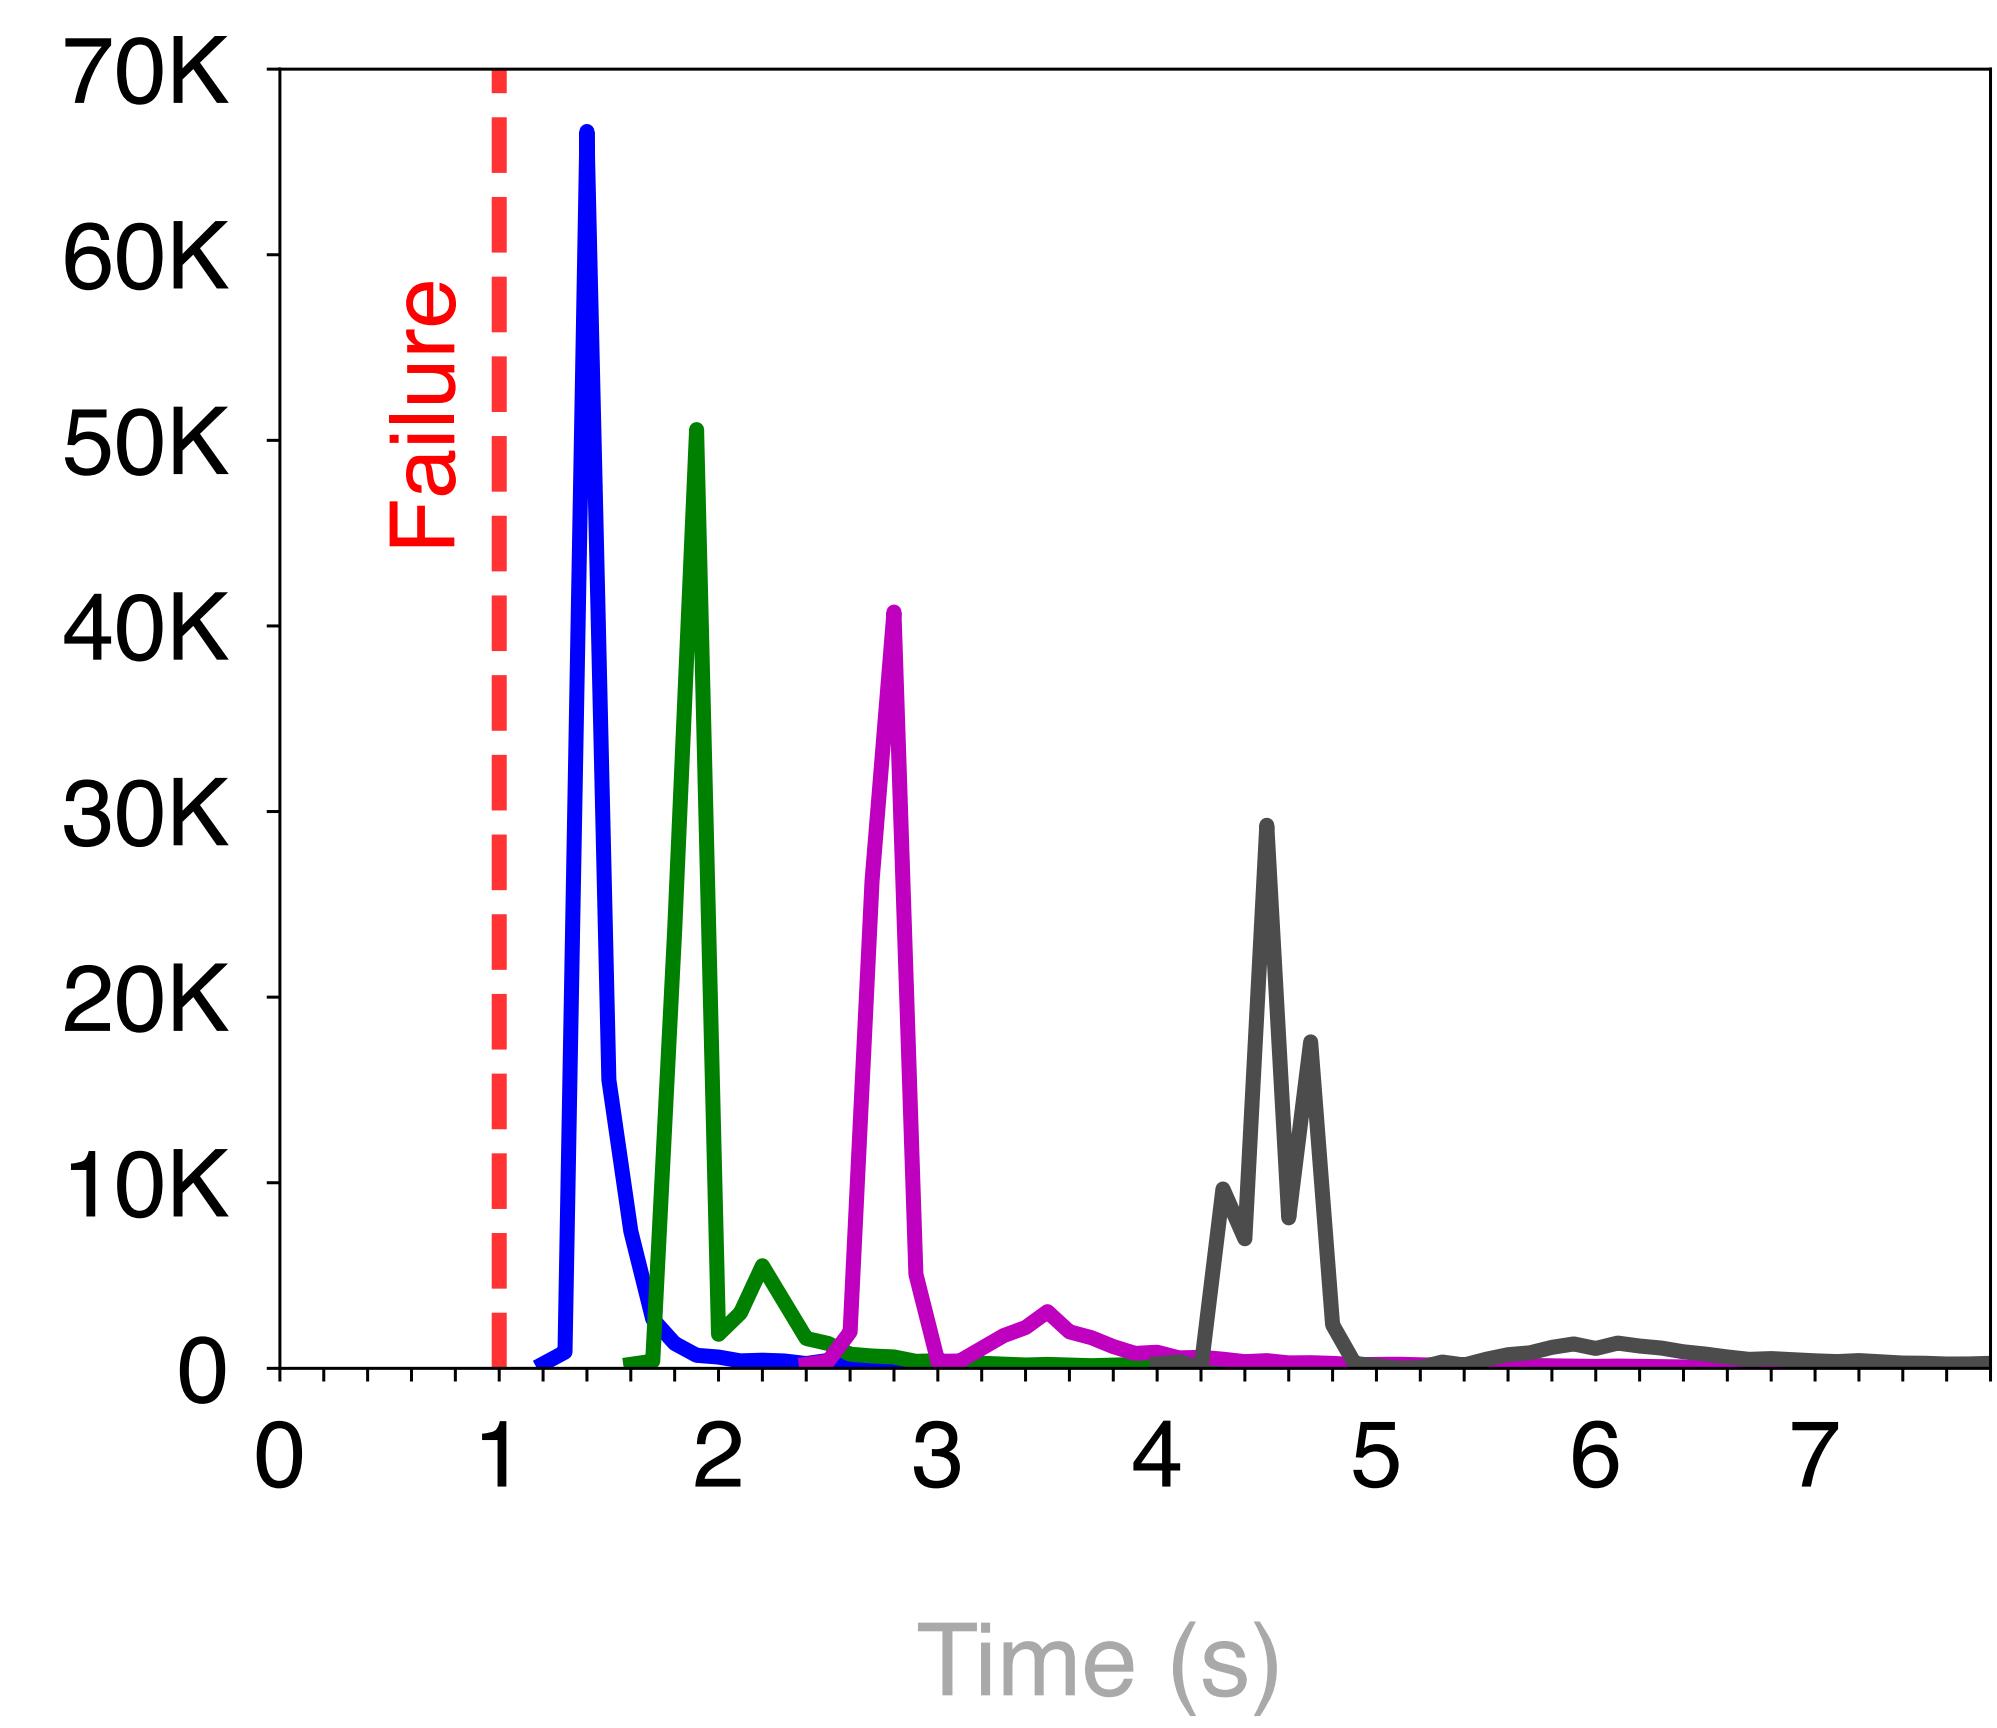
Number of
retransmissions



Number of
retransmissions



Number of
retransmissions



Tracking this signal in the data plane is challenging

Challenges

Signal is noisy
packets loss are routinely observed

Signal fades away quickly
due to the exponential backoff

Signal is compounded over many small ones
requires per-connection tracking, which is hard to scale

We solve these challenges by considering a subset of the signal that we carefully craft for maximal signal-to-noise ratio

Solutions

Signal is noisy

Focus on retransmissions caused by bursty losses

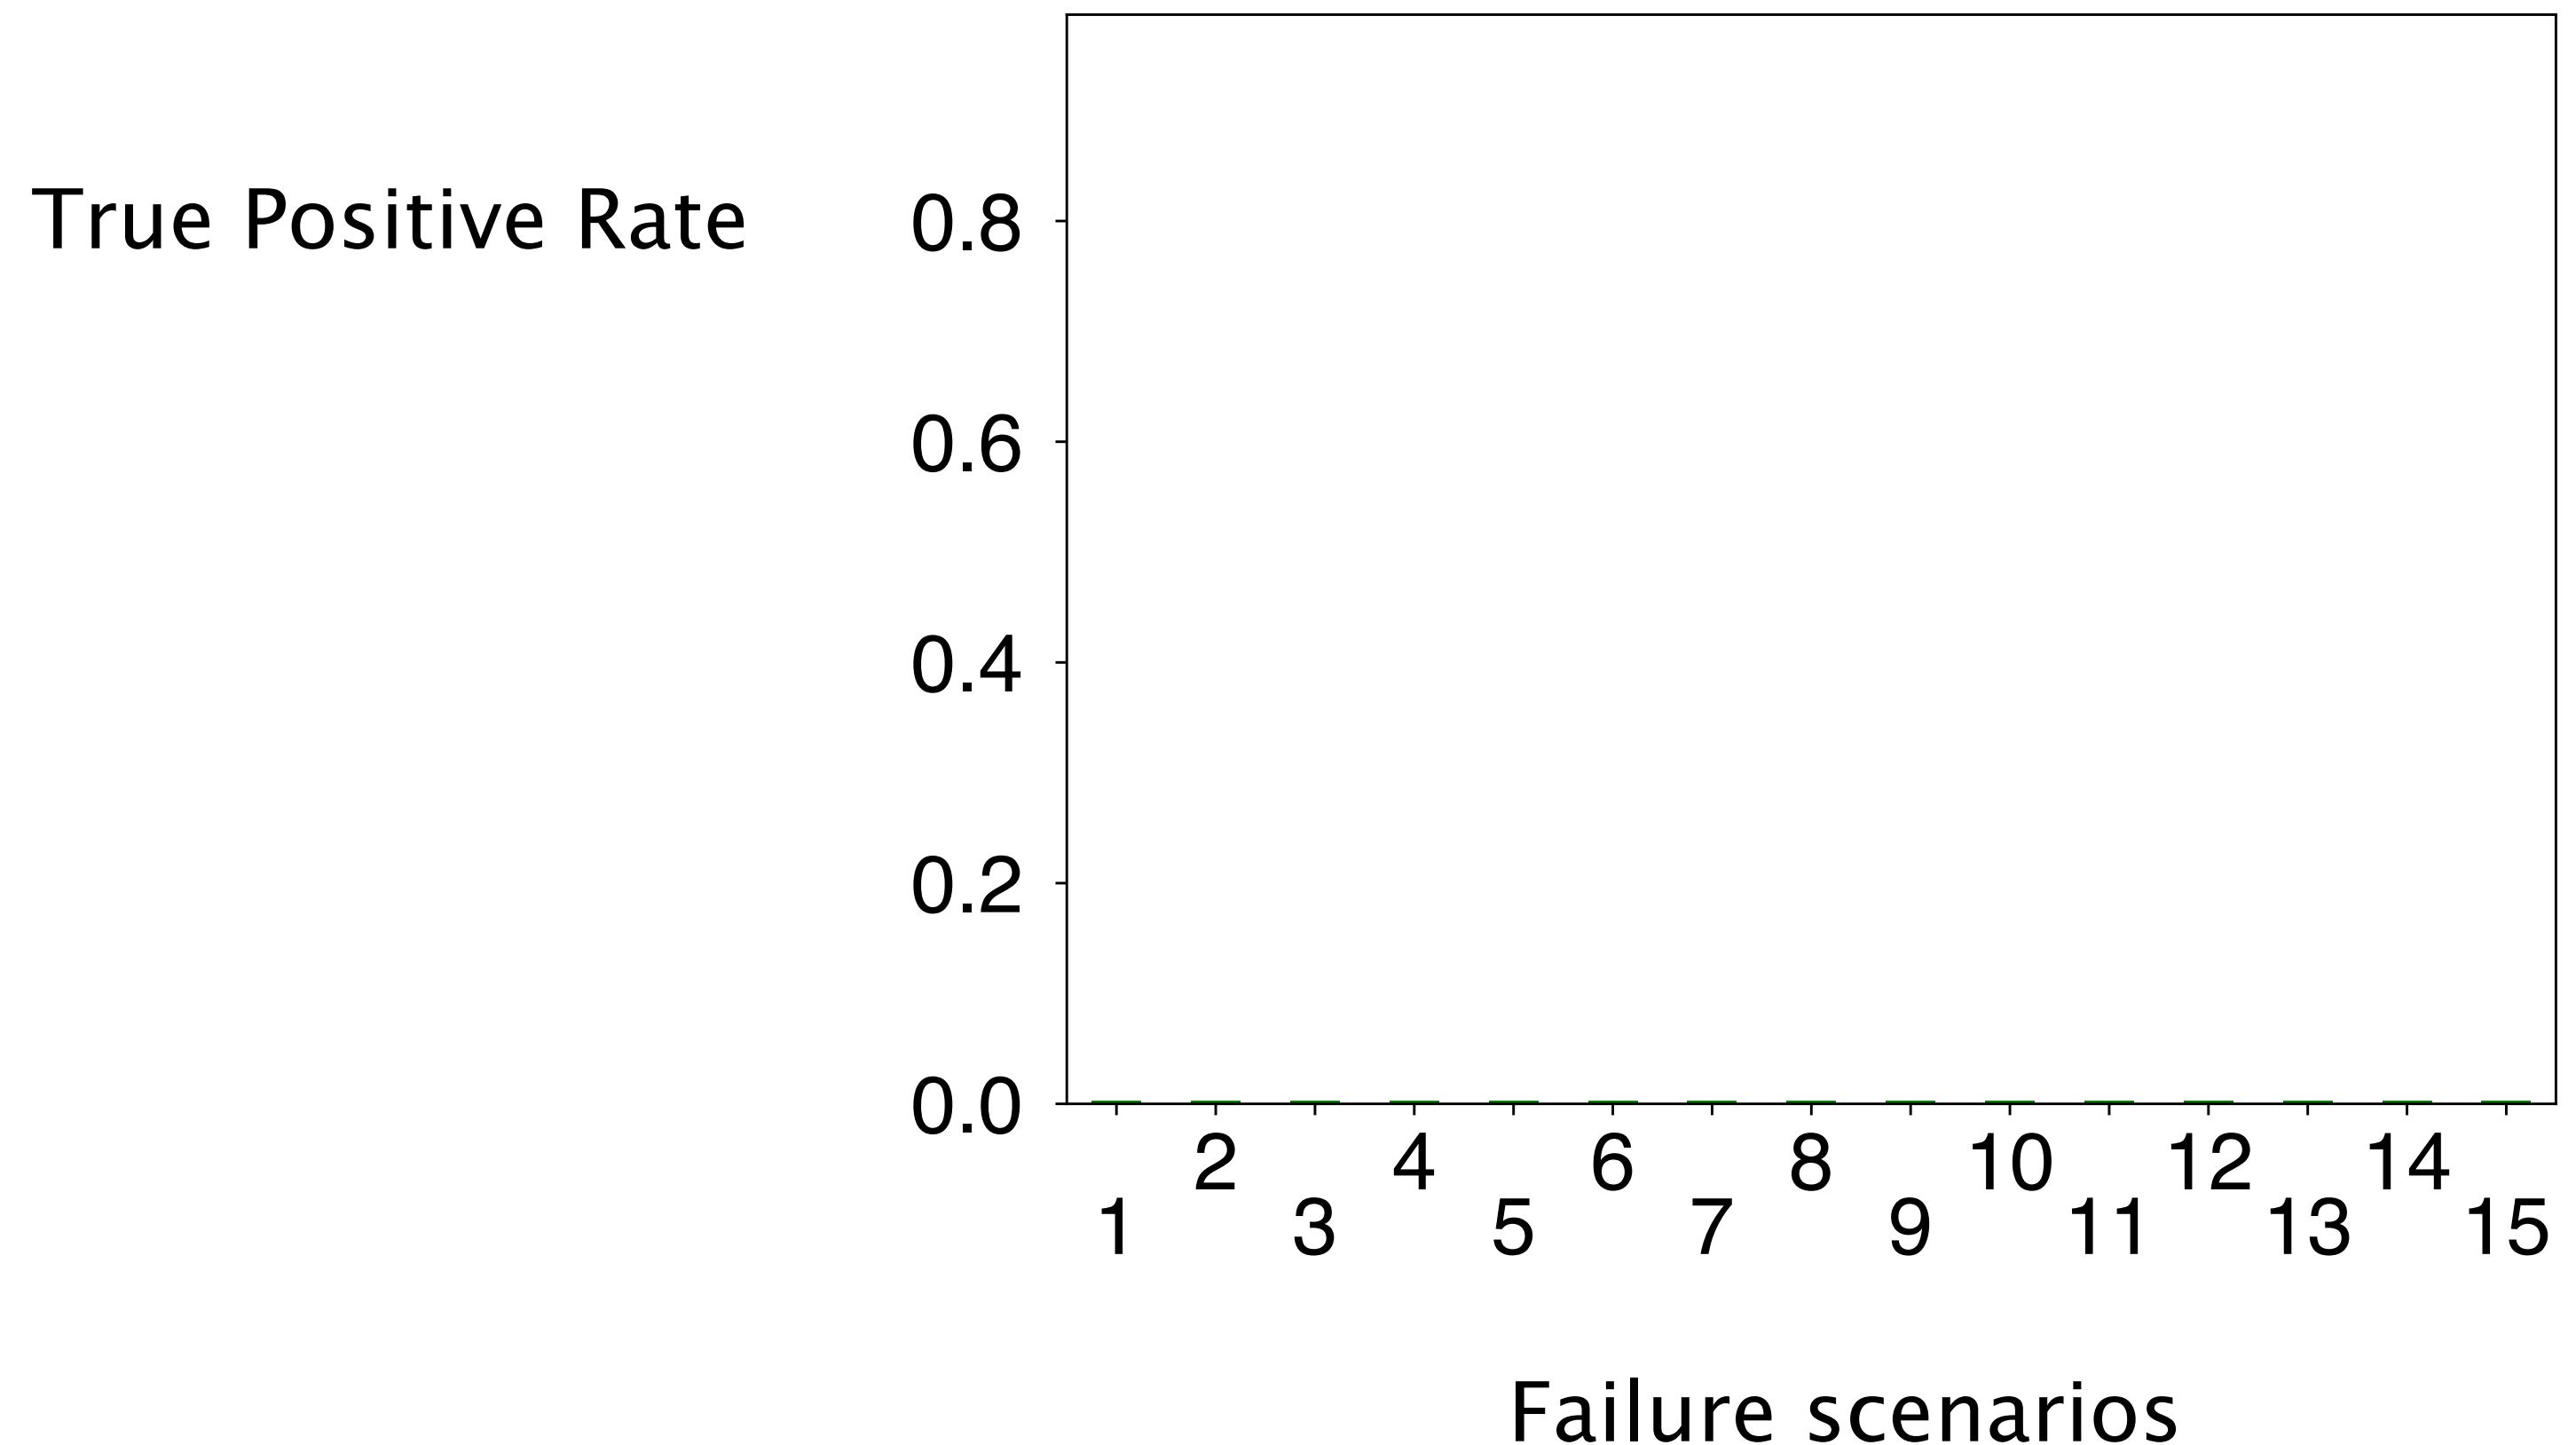
Signal fades away quickly

Focus on active flows (fast to retransmit after a failure)

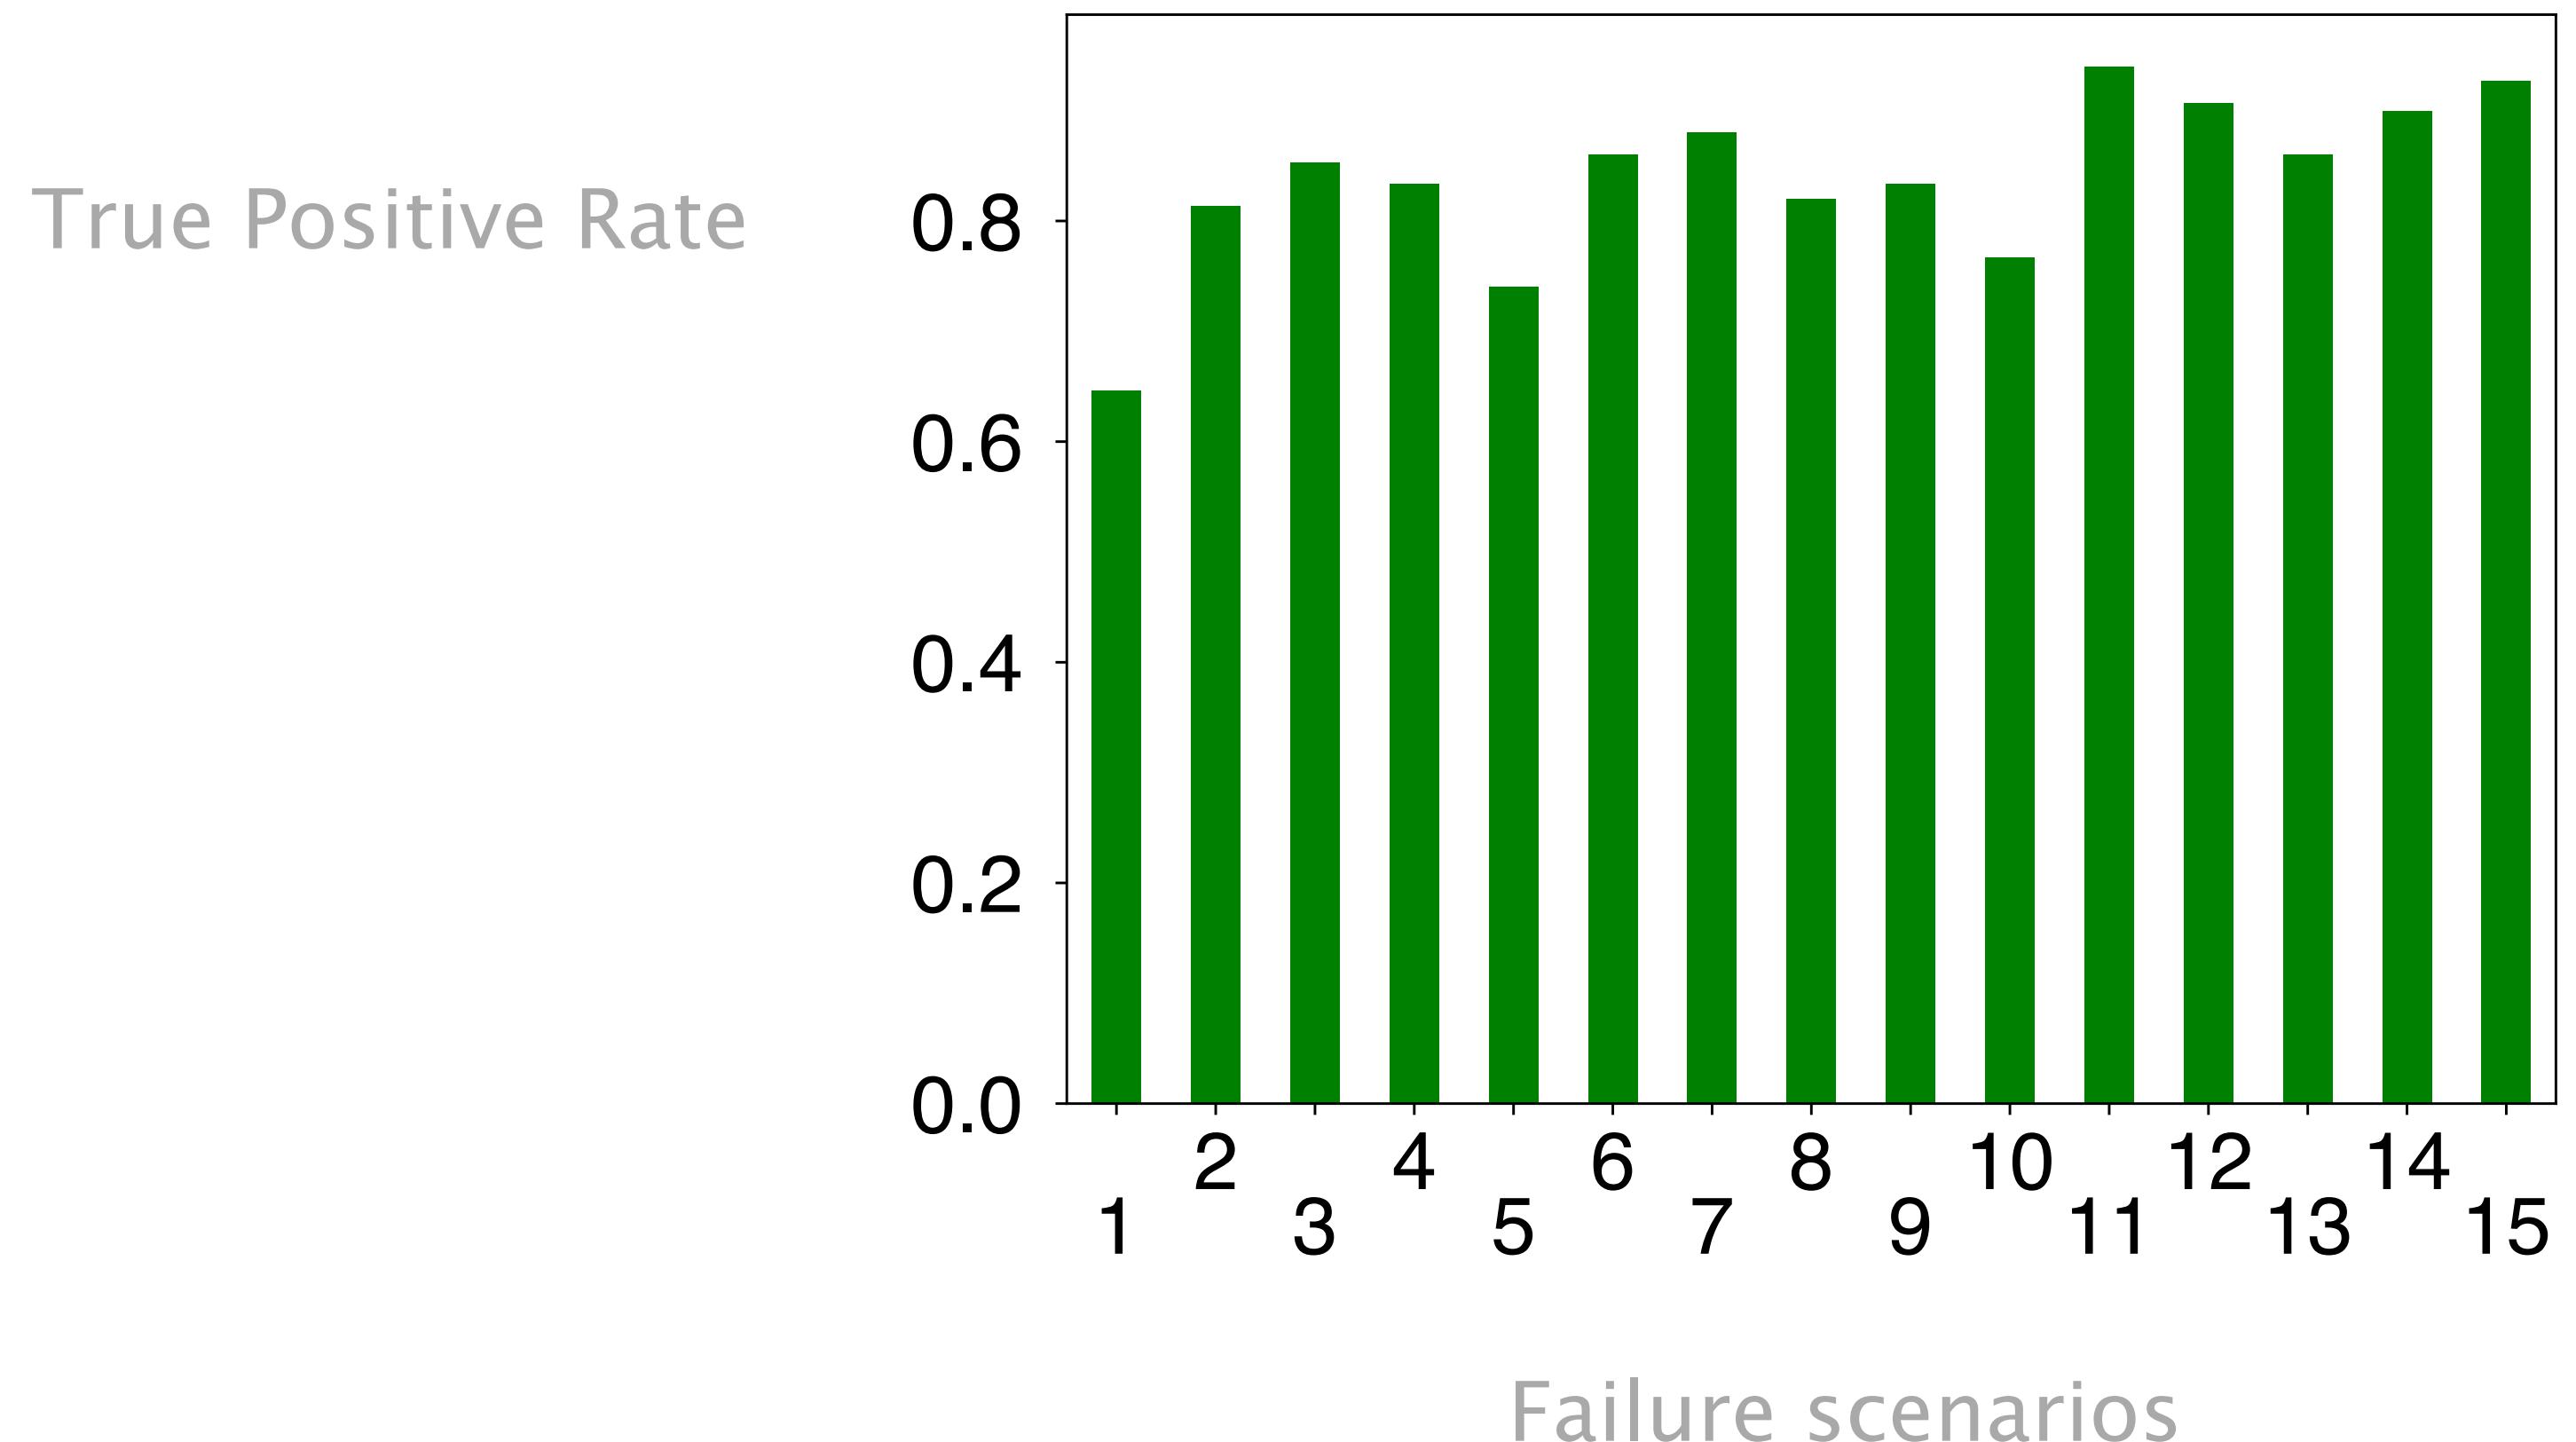
Signal is compounded over many small ones

Rely on scalable data structures and sampling

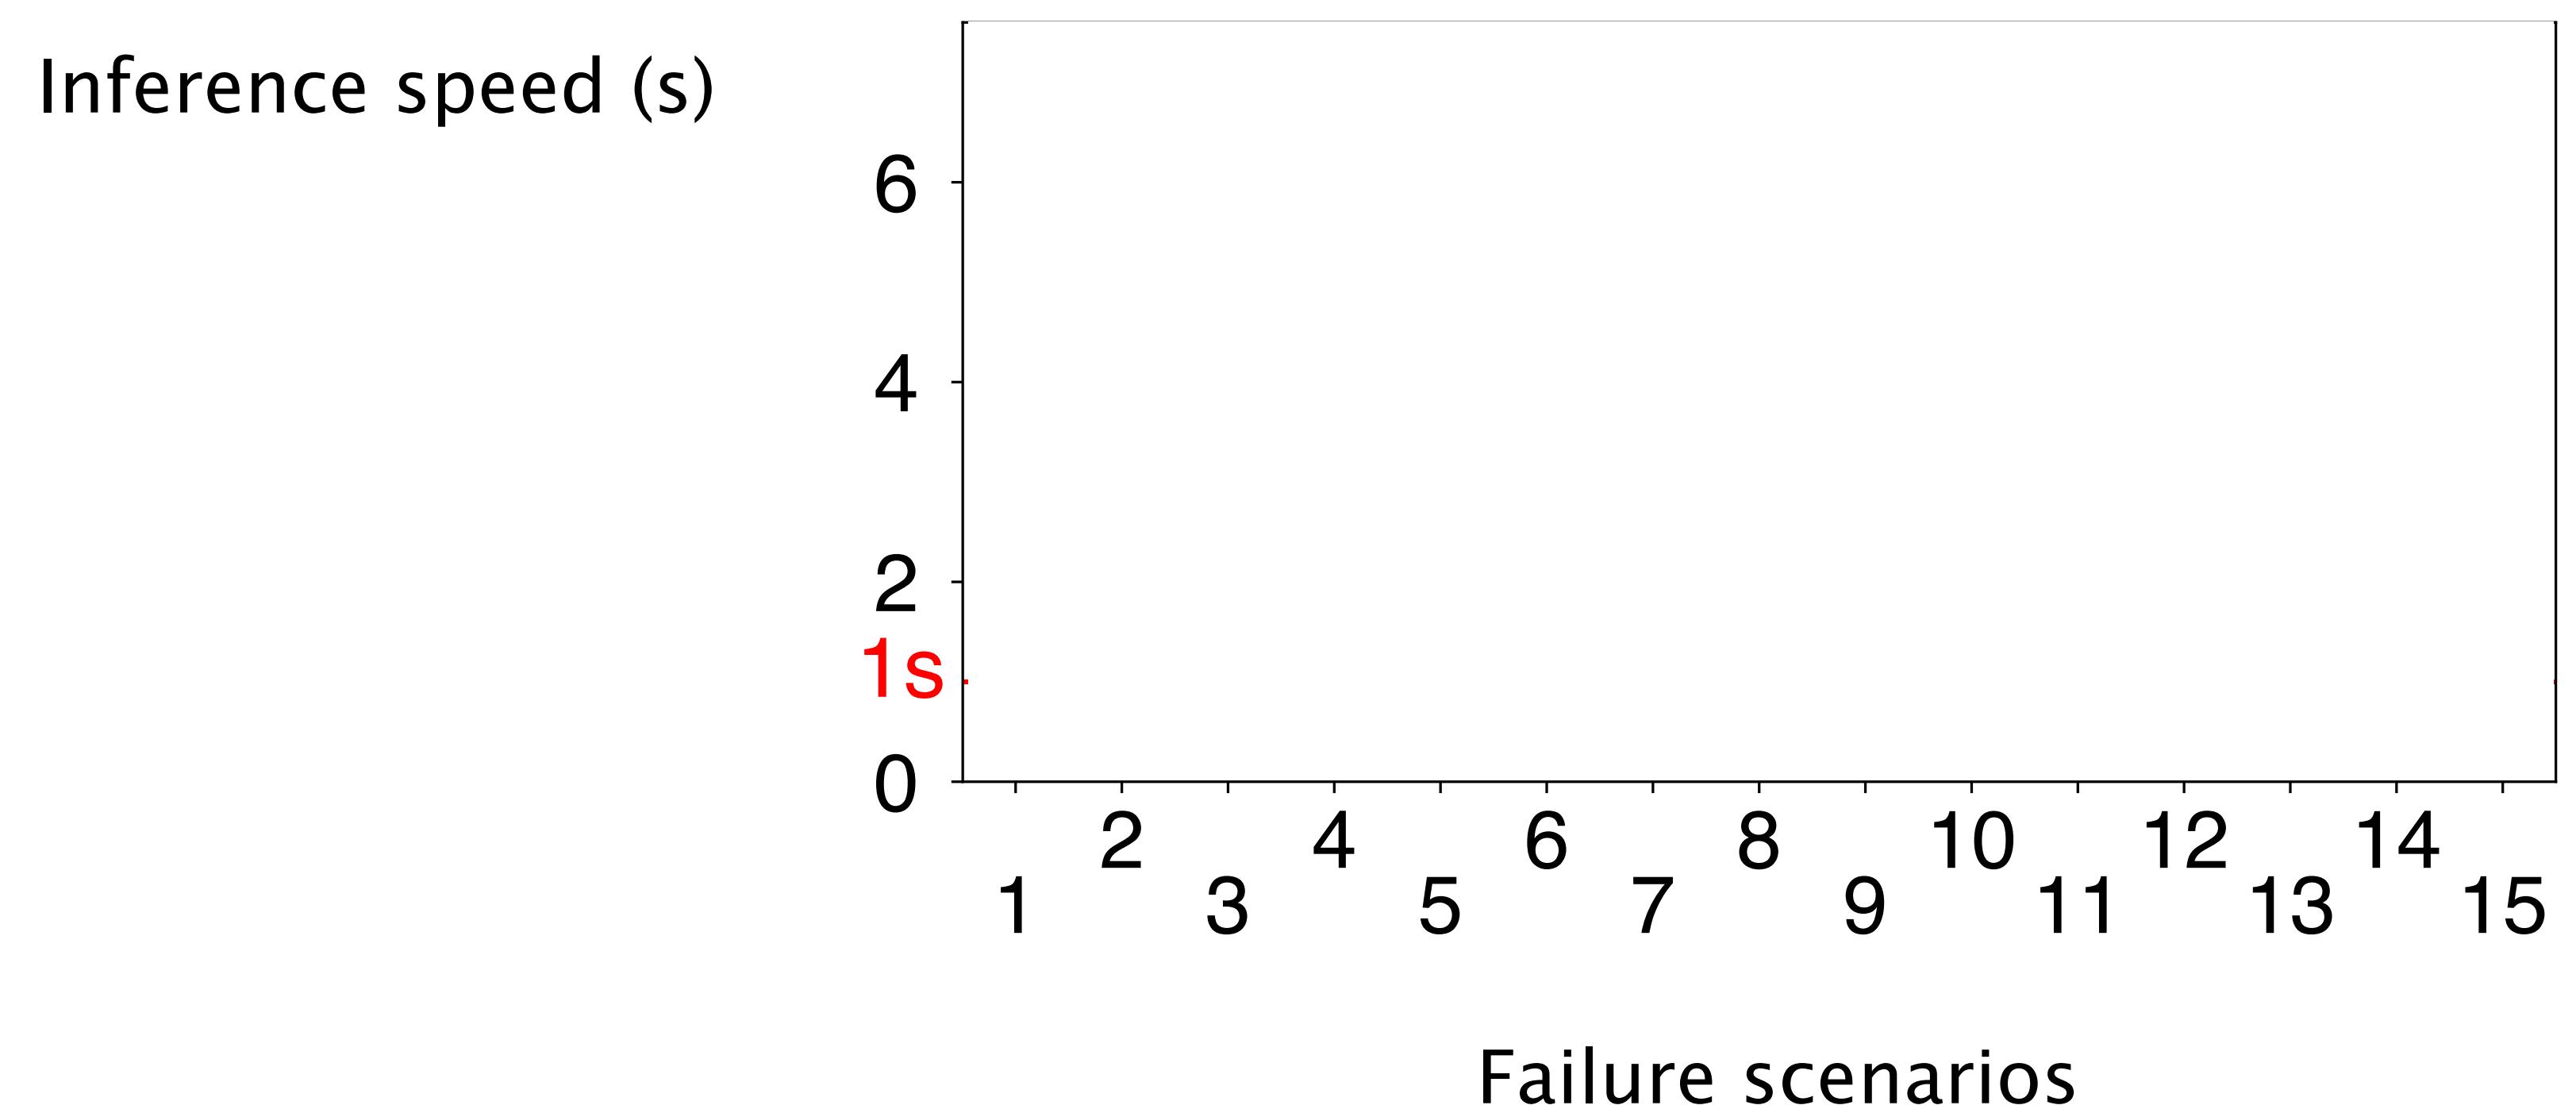
Blink works well in practice,
with an inference accuracy above 80% in most cases



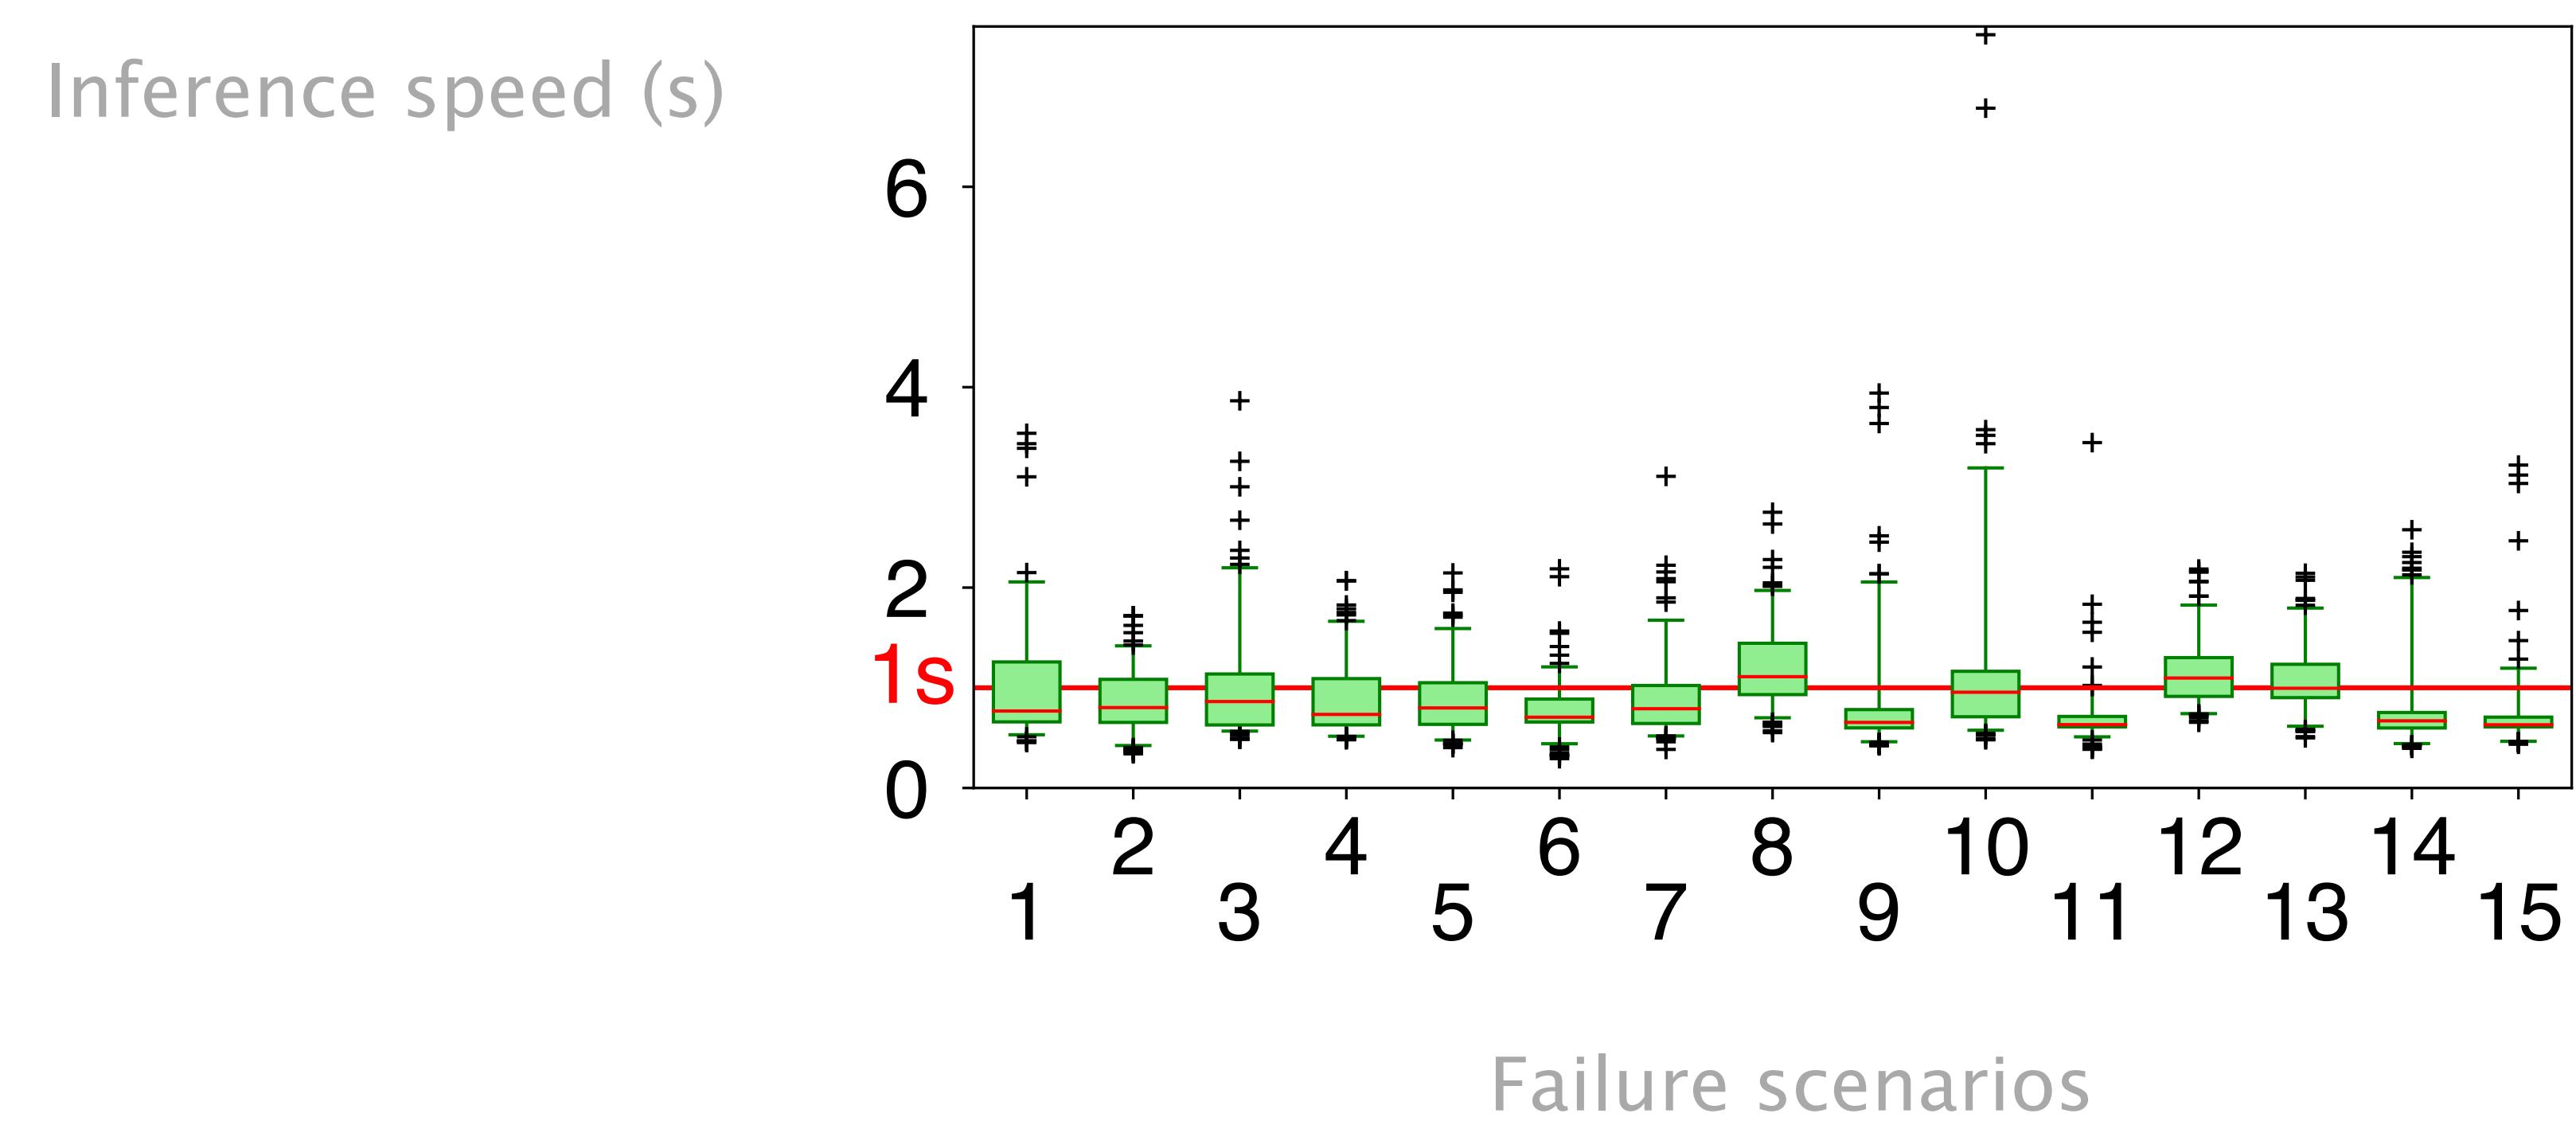
Blink works well in practice,
with an inference accuracy above 80% in most cases



Blink works well in practice,
with an inference speed below 1 second in most cases



Blink works well in practice,
with an inference speed below 1 second in most cases



We intend to generalize the signals we track,
for a wider variety of applications

reliability

reroute around fine-grained failures

performance

steer traffic along the most performant path

security

detect unwanted traffic redirection

Part II: Taking control

applications

frameworks

methods

monitoring/inference

FitNets: An Adaptive Framework to Learn Accurate Traffic Distributions



Alexander Dietmüller



Albert
Gran Alcoz

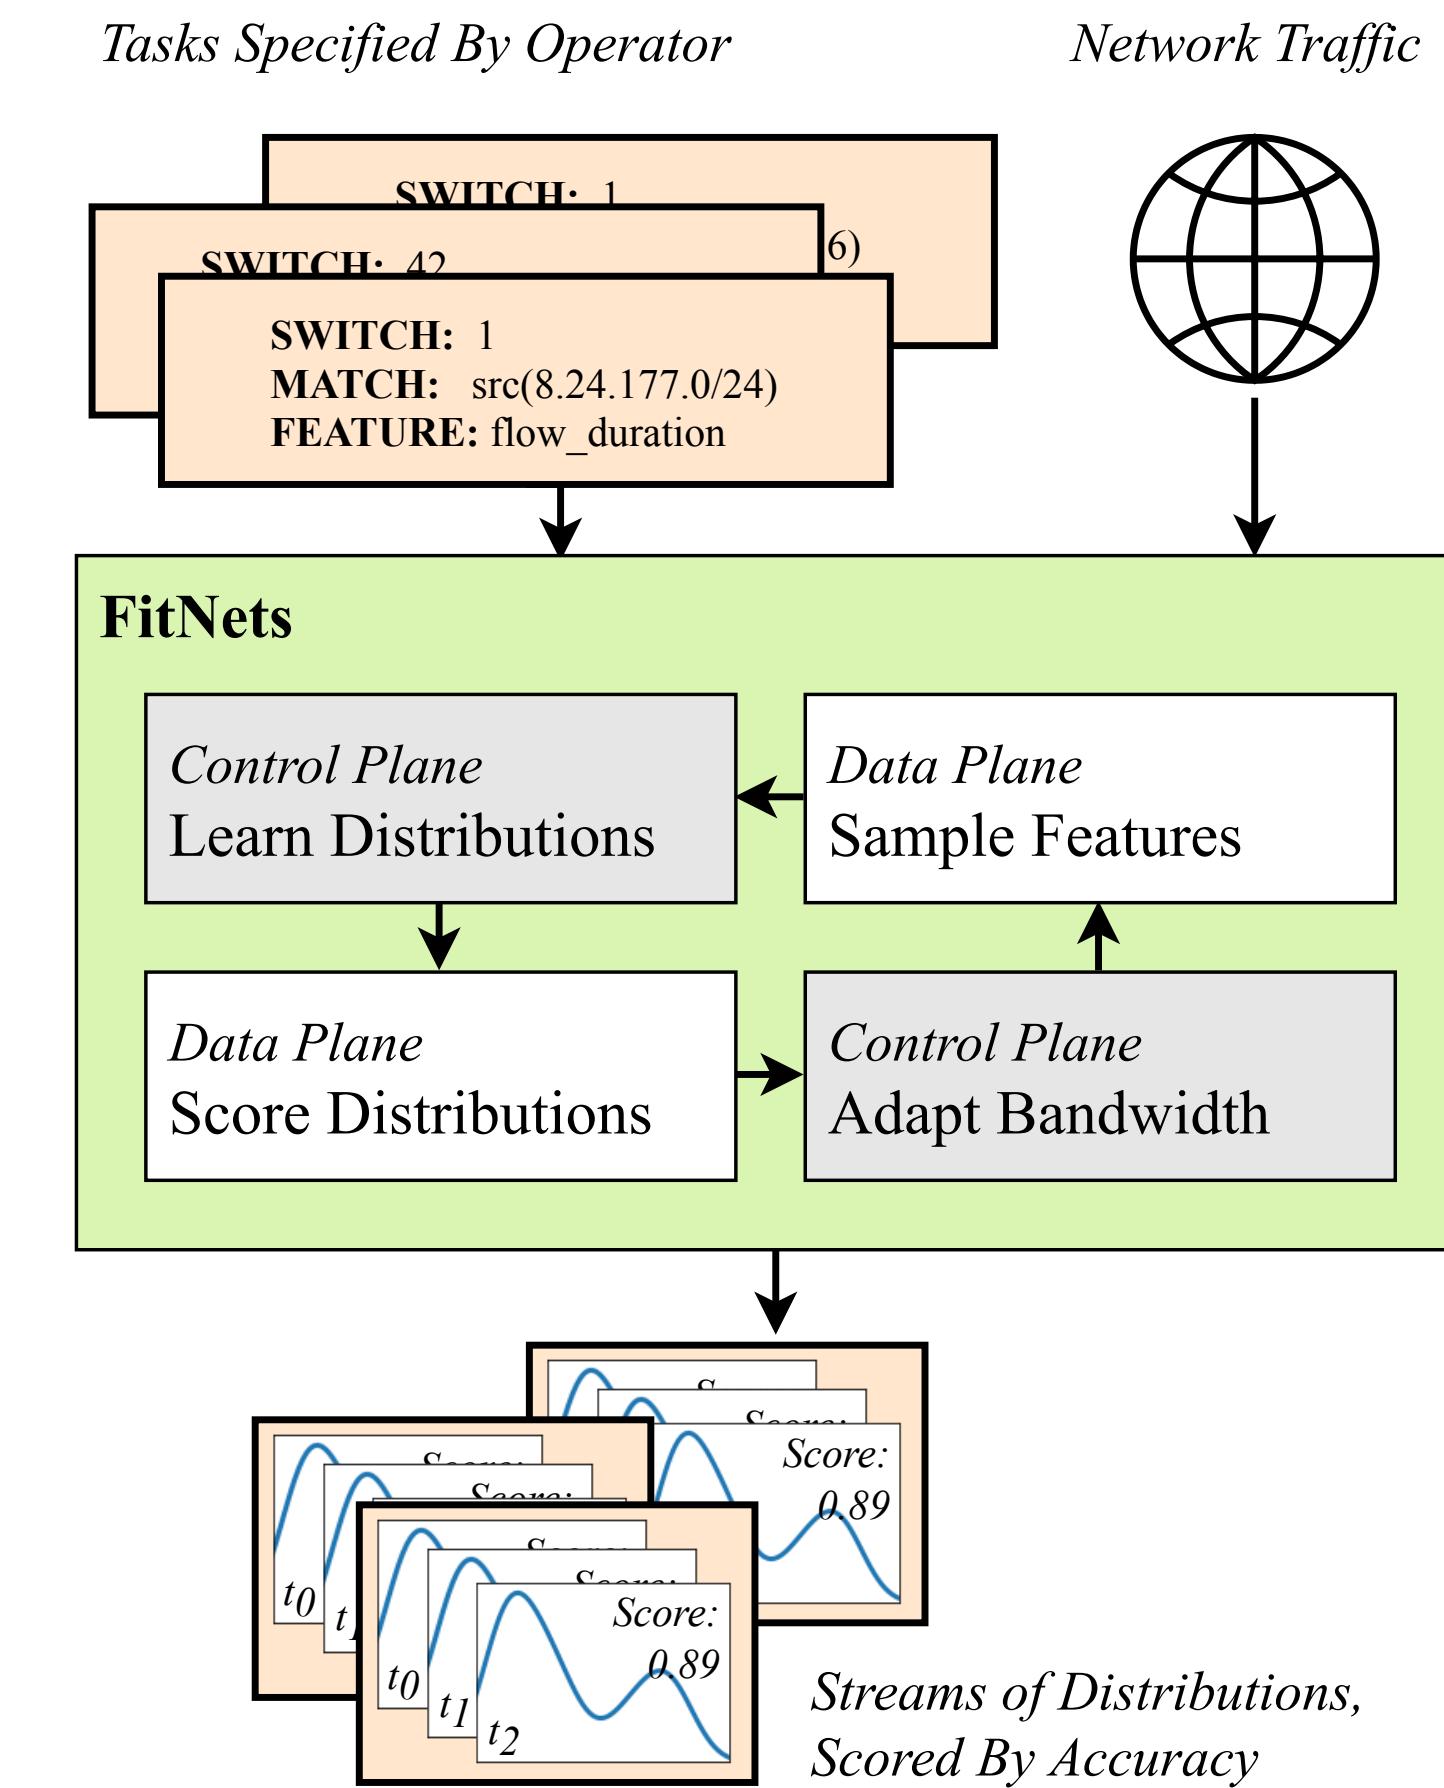


Laurent Vanbever

FitNets

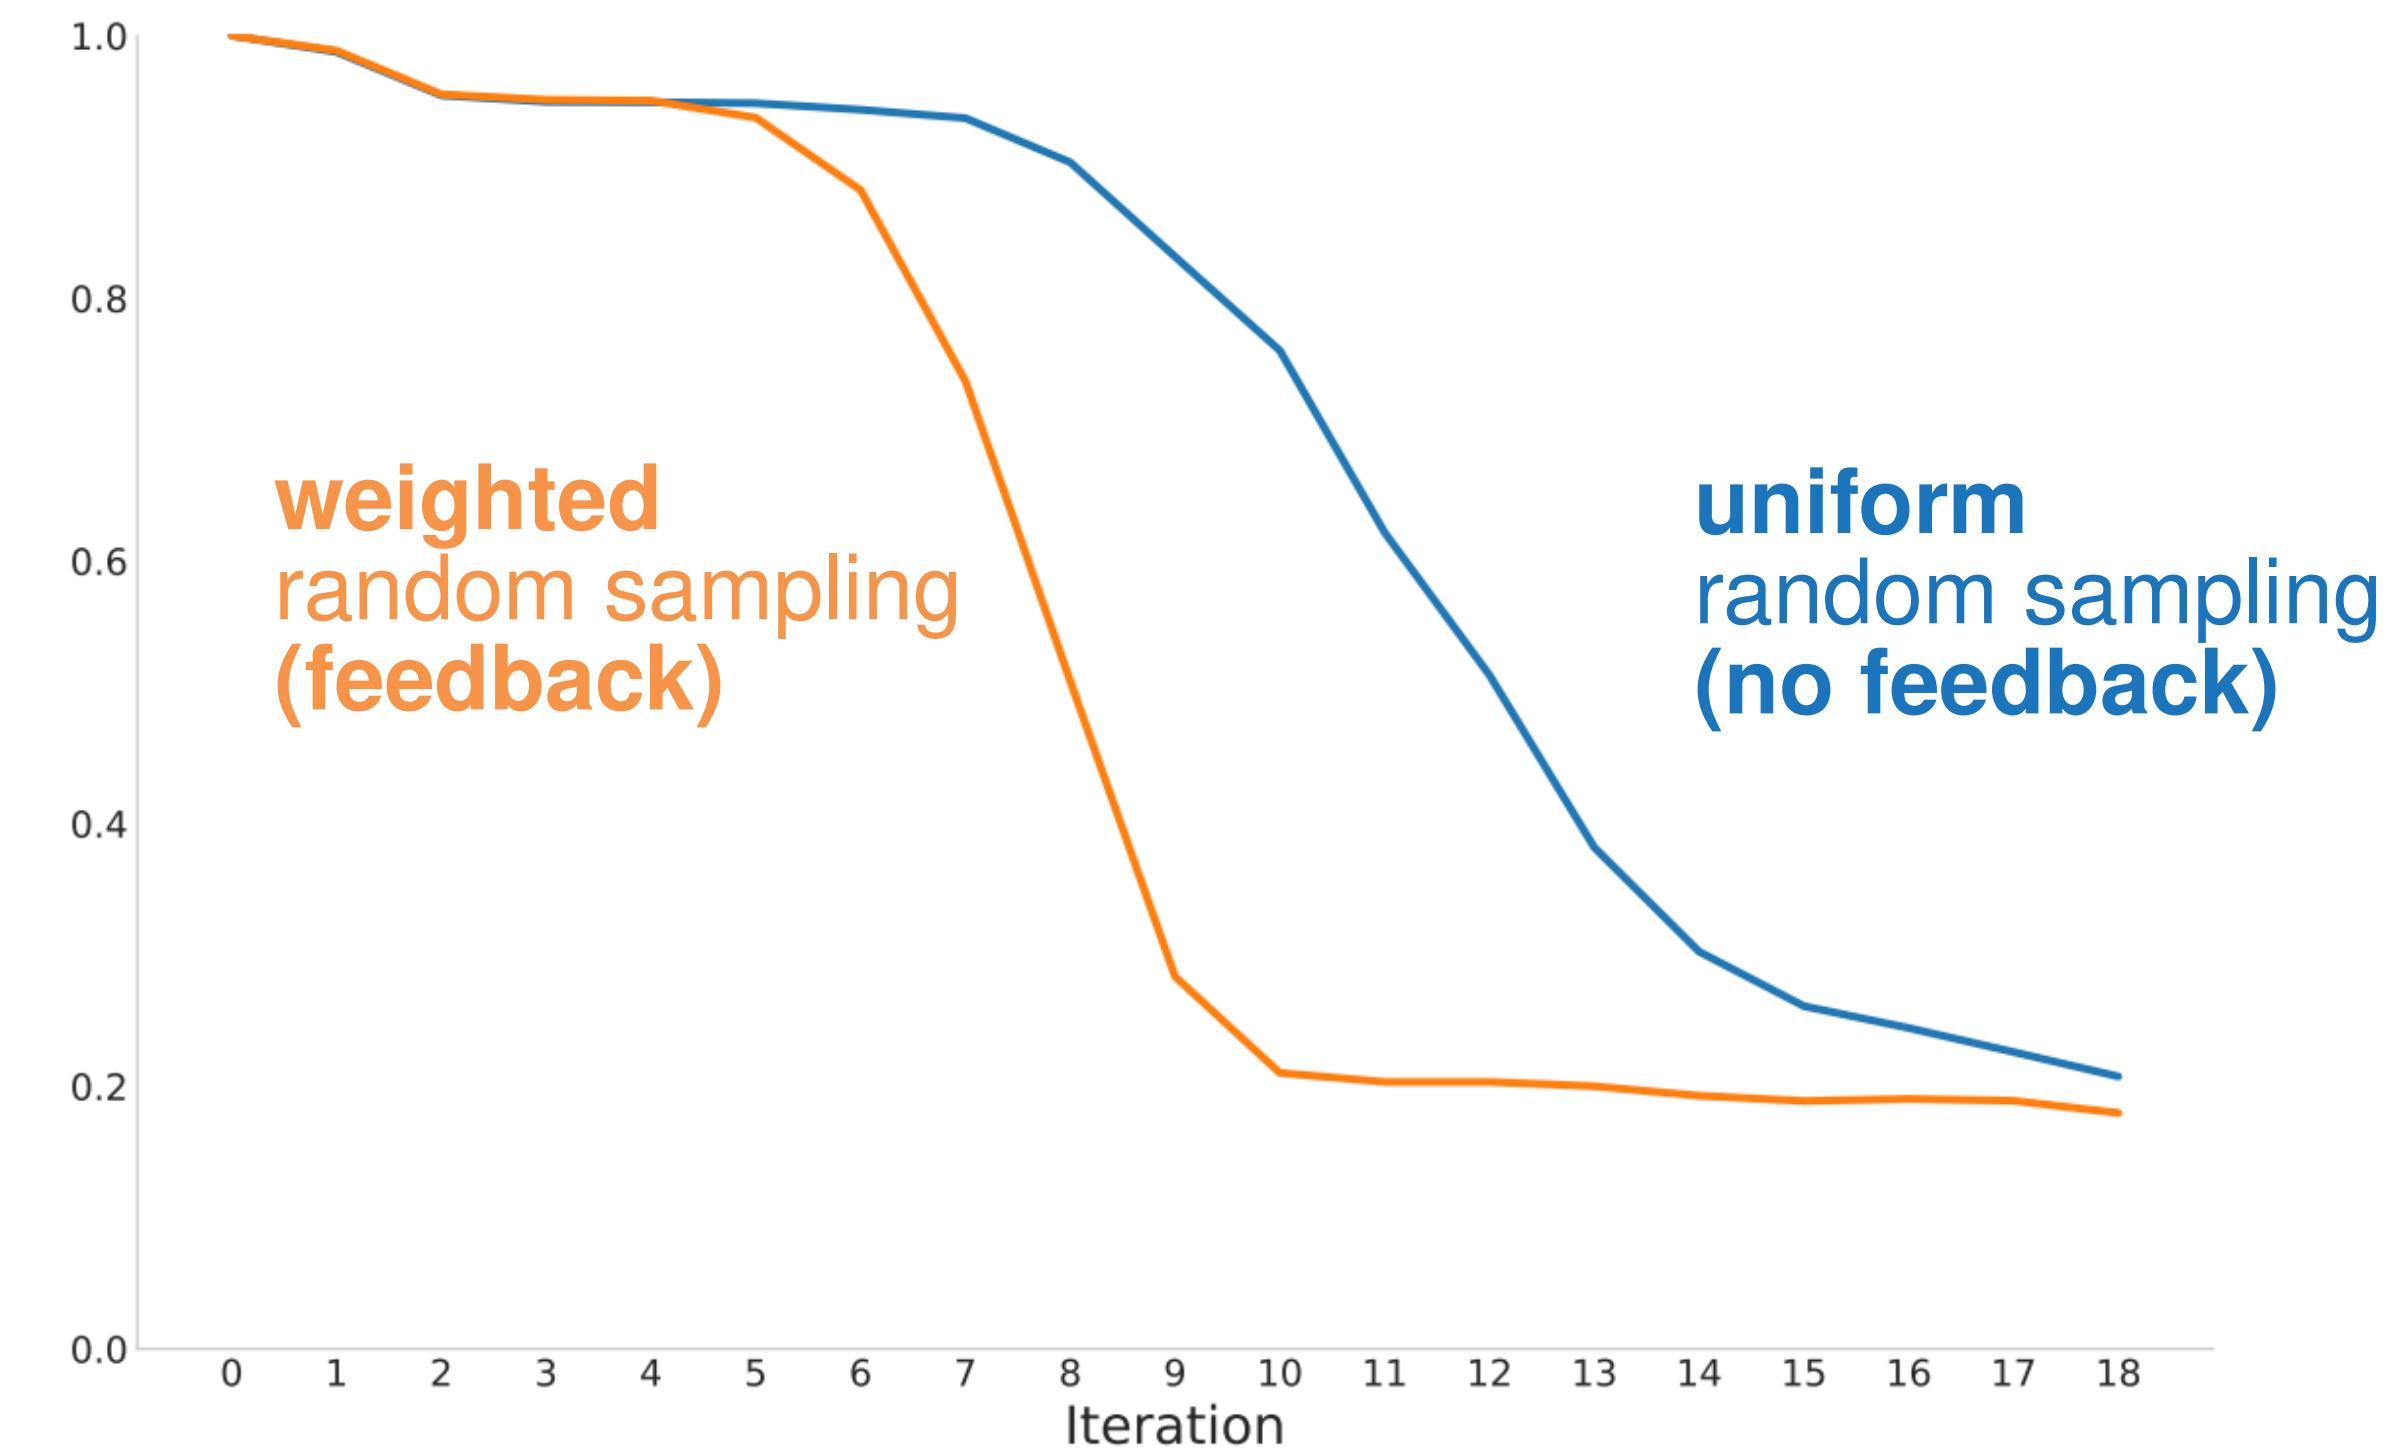
An adaptive network monitoring framework

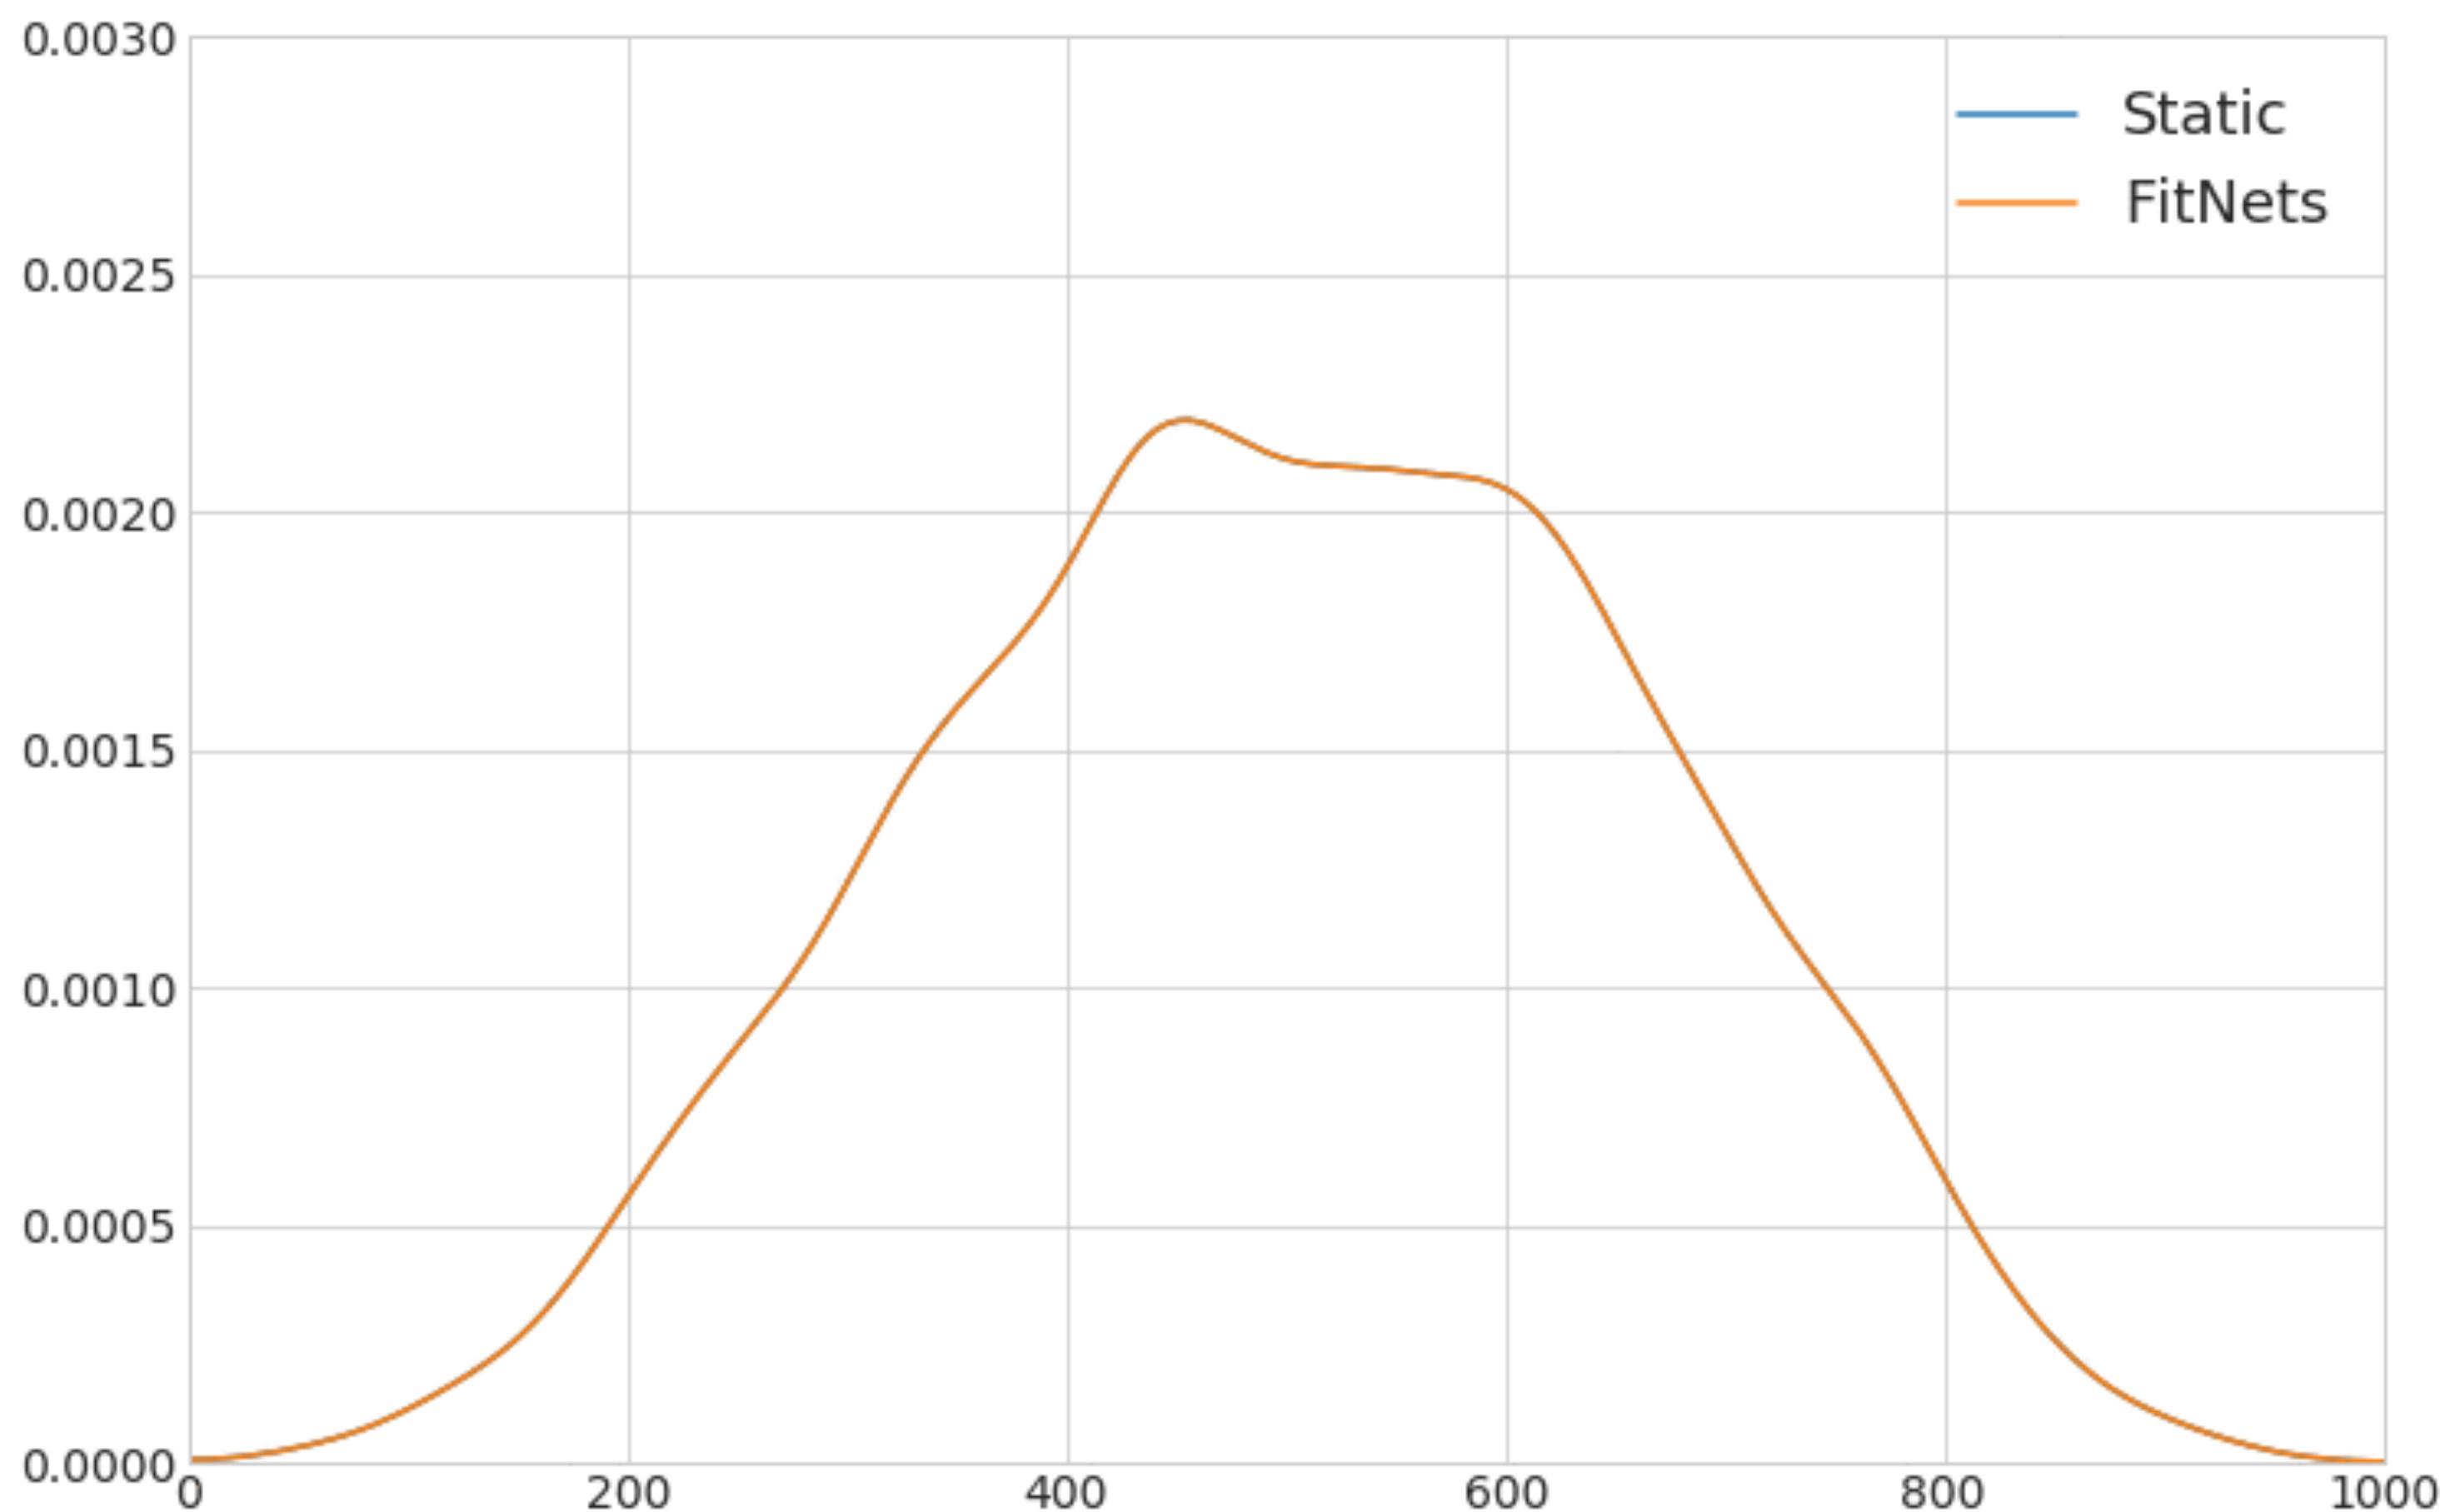
- computes traffic distribution in the control plane, on traffic samples
- scores these distributions in the data plane, on *all* traffic
- adapts the sampling rate according to the score



FitNets learns distribution better and faster
than simply randomly sampling traffic

Distance from
true distribution
(Jensen-Shannon)





Part II: Taking control

applications

frameworks

methods

monitoring/inference

pForest: In-Network Inference with Random Forests



Coralie Busse-Grawitz



Roland Meier



Alexander Dietmüller



Tobias Bühler



Laurent Vanbever

Why (and How) Networks Should Run Themselves

Nick Feamster and Jennifer Rexford
Princeton University

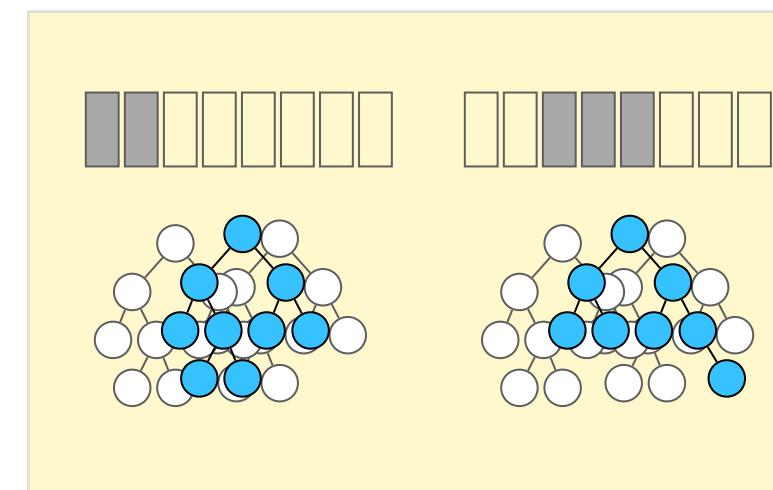
Abstract

The proliferation of networked devices, systems, and applications that we depend on every day makes managing networks more important than ever. The increasing security, availability, and performance demands of these applications suggest that these increasingly difficult network management

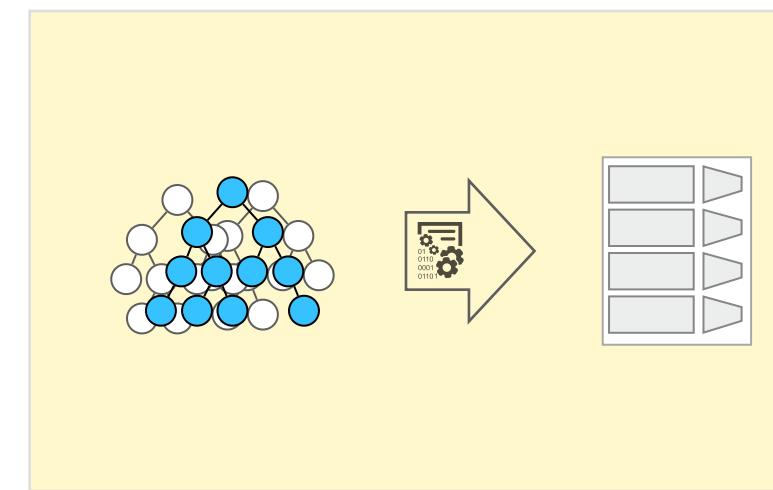
relationships between them and user quality of experience become increasingly complex. Twenty years ago, we had some hope of (and success in) creating clean, closed-form models of individual protocols, applications, and systems [4, 24]; today, many of these are too complicated for closed-form analysis. Prediction problems such as determining how search query response time would vary in response to the placement

Performing automated, real-time inference: The past ten years has demonstrated significant promise in using machine learning to both detect and predict network attacks; we must build on the increasing amount of work in automated infer-

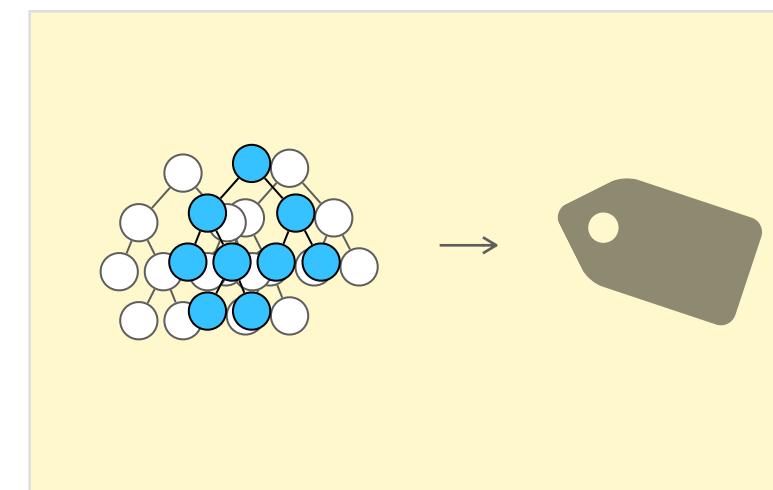
pForest: In-Network Inference with Random Forests



Optimizing random forest models
for programmable network devices



Compiling random forest models
to programmable network devices



Performing runtime classification
at Tbps

Early, accurate & efficient classification of an ongoing flow
as an optimization problem

Given a labeled dataset \mathcal{F} and a threshold τ
find a classifier C such that

- $\text{accuracy}(C) \geq \tau$
- C fits in programmable network devices

while

- minimizing memory usage
- maximizing classification speed

Part II: Taking control

applications

frameworks

methods

some of our challenges
lying ahead

Building and operating truly self-driving networks
will require us to overcome fundamental challenges

data

correctness

interpretability

data

correctness

interpretability



data is network-dependent

poor generalization capabilities

interesting events are rare

few interesting examples (if any)

data can easily be polluted

by anyone with a Internet connection...

data

correctness

interpretability



how do we...

ensure correctness?

network should be reachable

guarantee stability?

agents can clash with each other

reason about optimality?

data-driven > non-data-driven (?)

data

correctness

interpretability

how can operators...

reason about self-driving networks?
especially in partial deployment

debug their self-driving networks?
manual override

Self-Driving Networks

Breaking new ground in network automation



- 1 operator assistance
- 2 partial automation

Part I
assisting operators

- 3 conditional automation
- 4 high automation

Part II
taking control

- 5 full automation

too futuristic? 😊



Huge thanks to my research group!



Roland



Ahmed



Mirja



Edgar



Tobias



Thomas



Rüdiger



Albert



Maria

Merci à tous!

Group

Maria Apostolaki
Rüdiger Birkner
Tobias Bühler
Edgar Costa Molero
Alexander Dietmüller
Ahmed El-Hassany
Albert Gran Alcoz
Thomas Holterbach
Roland Meier

Bernhard Ager
David Gugelmann
Mirja Kühlewind
Brian Trammell

Collaborators

Alberto Dainotti
Dana Drachsler Cohen
Nick Feamster
Arpit Gupta
Markus Happe
Vincent Lenders
Jennifer Rexford
Michael Schapira
Petar Tsankov
Stefano Vissicchio
Martin Vechev
Aviv Zohar

+ many more

Colleagues @ITET

Lothar Thiele
Roger Wattenhofer

Gian-Luca Bona (committee)
Amos Lapidoth
John Lygeros

+ the entire D-ITET faculty & staff

Colleagues @ETH

Schulleitung
Rektorat
Professoren Ruderteam

+ *many* more

Self-Driving Networks

Breaking new ground in network automation



Laurent Vanbever
nsg.ee.ethz.ch

ETH Zurich
May 29 2019